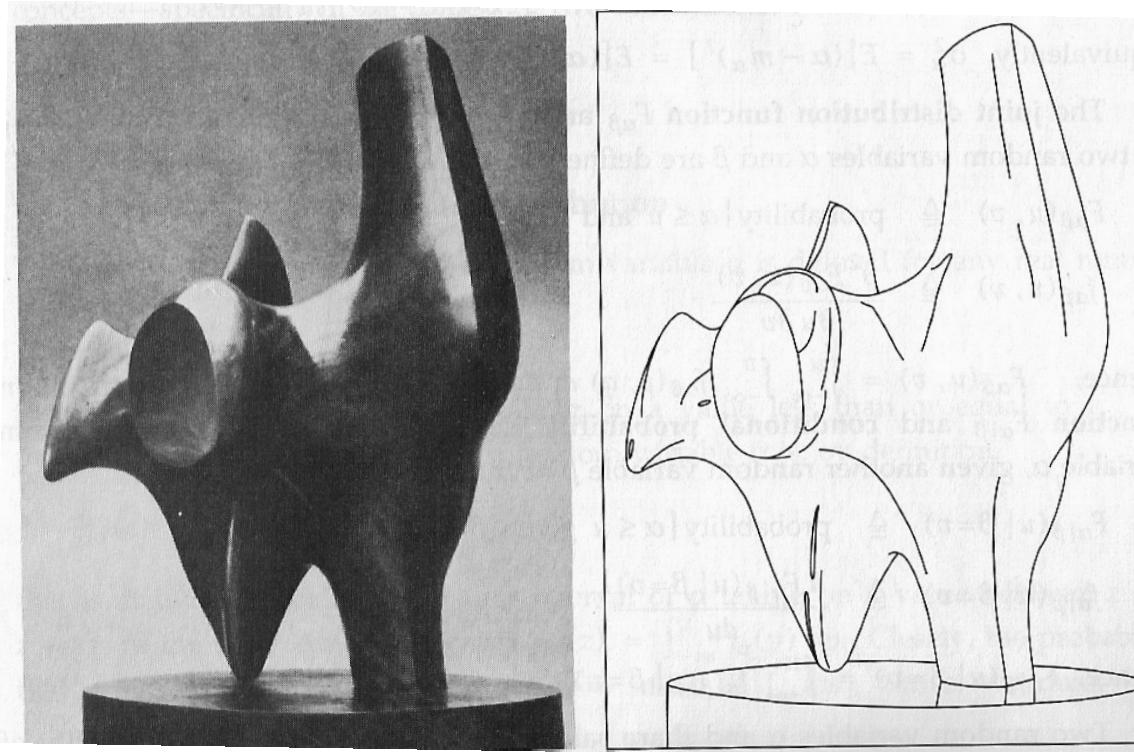
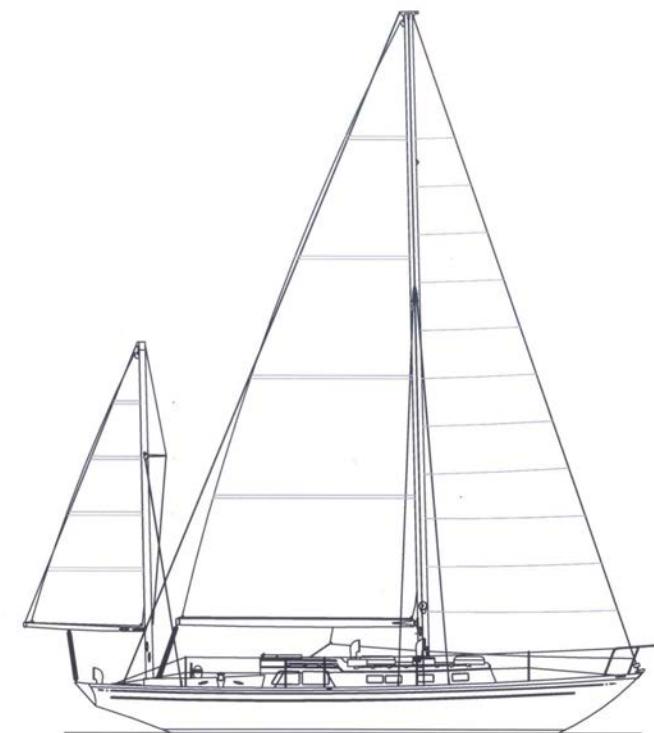
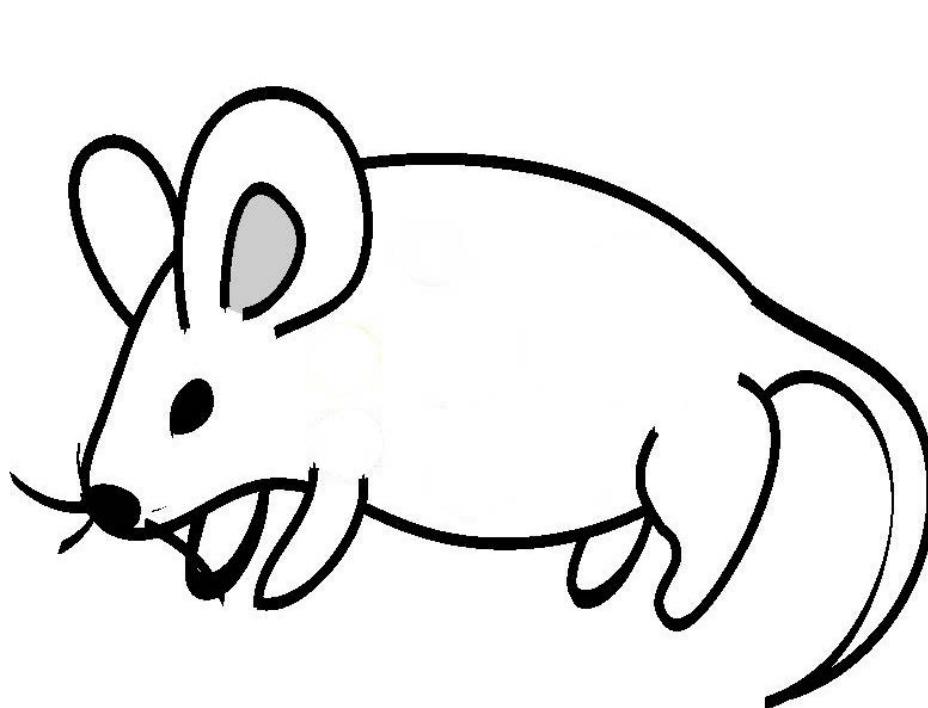


EDGE DETECTION



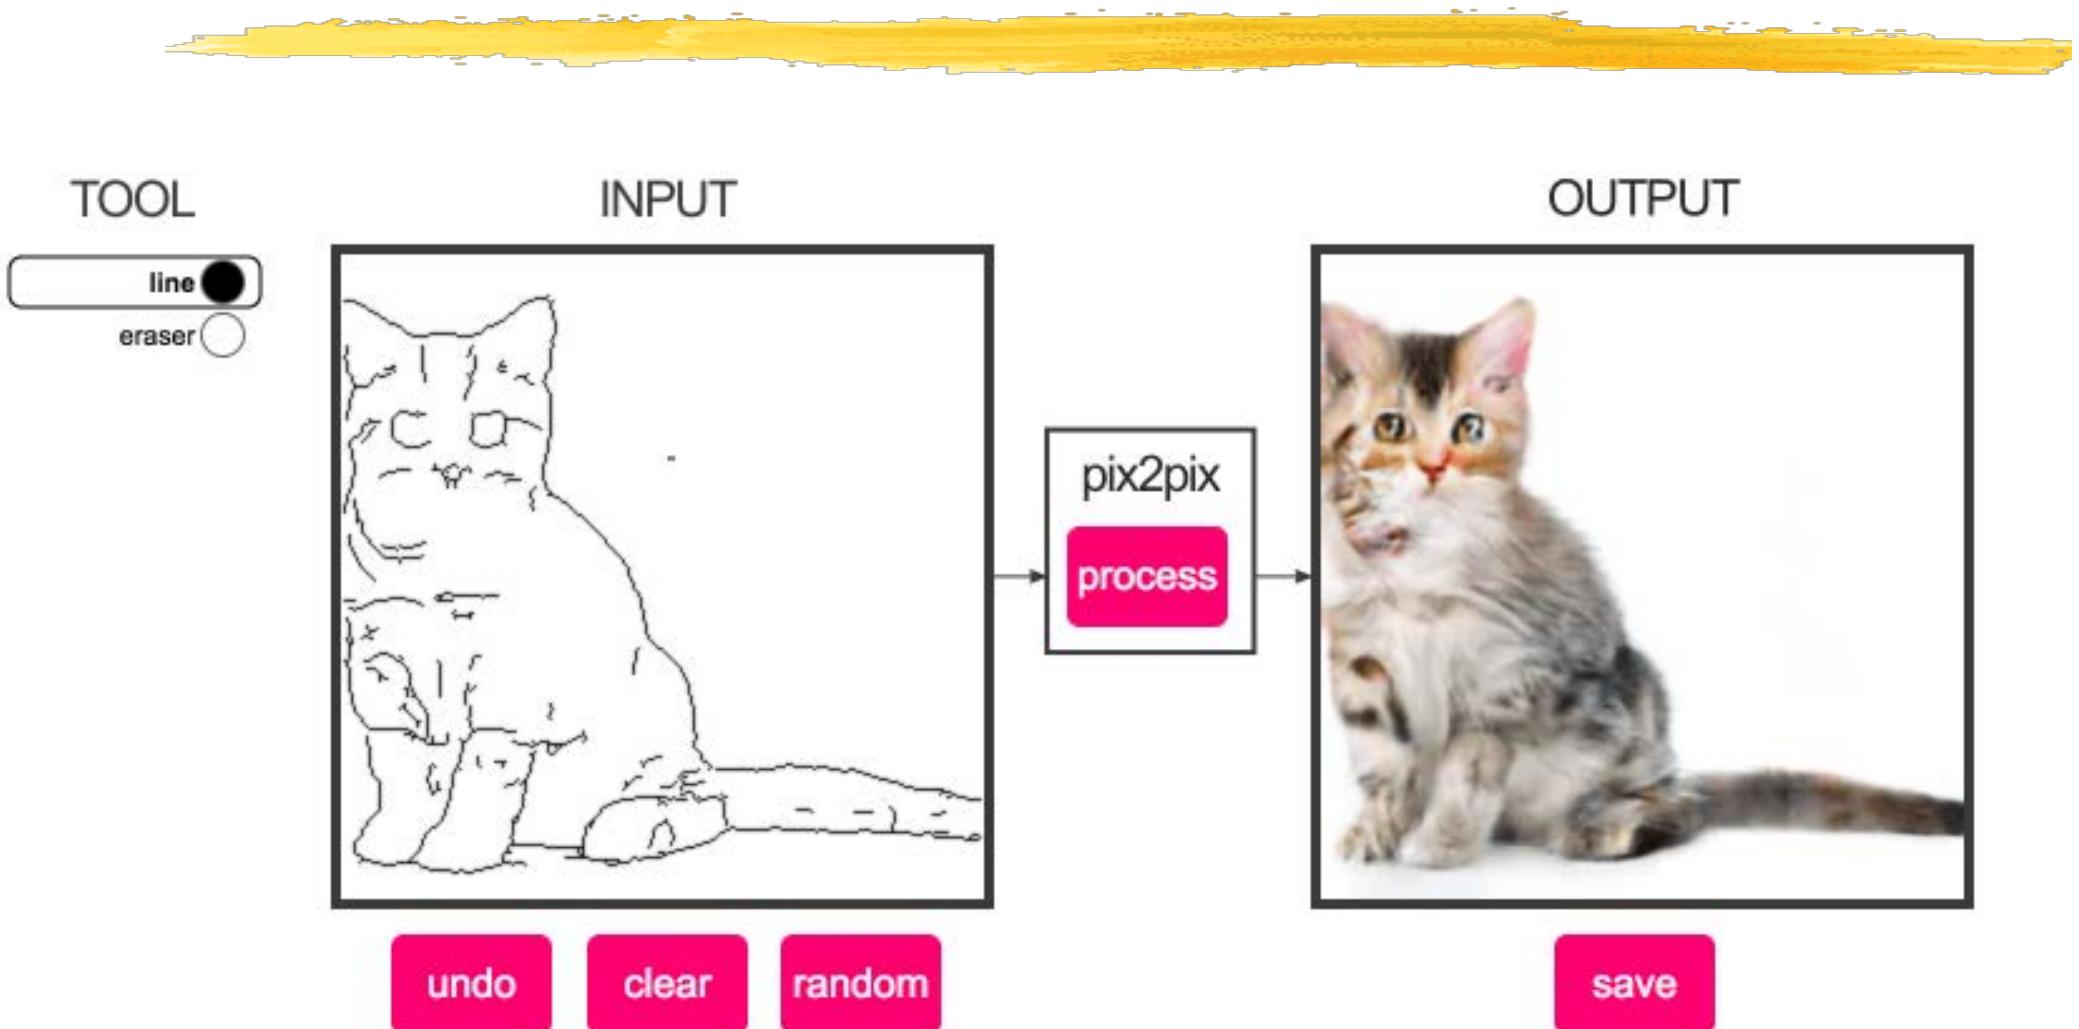
- What's an edge
- Image gradients
- Edge operators

LINE DRAWINGS



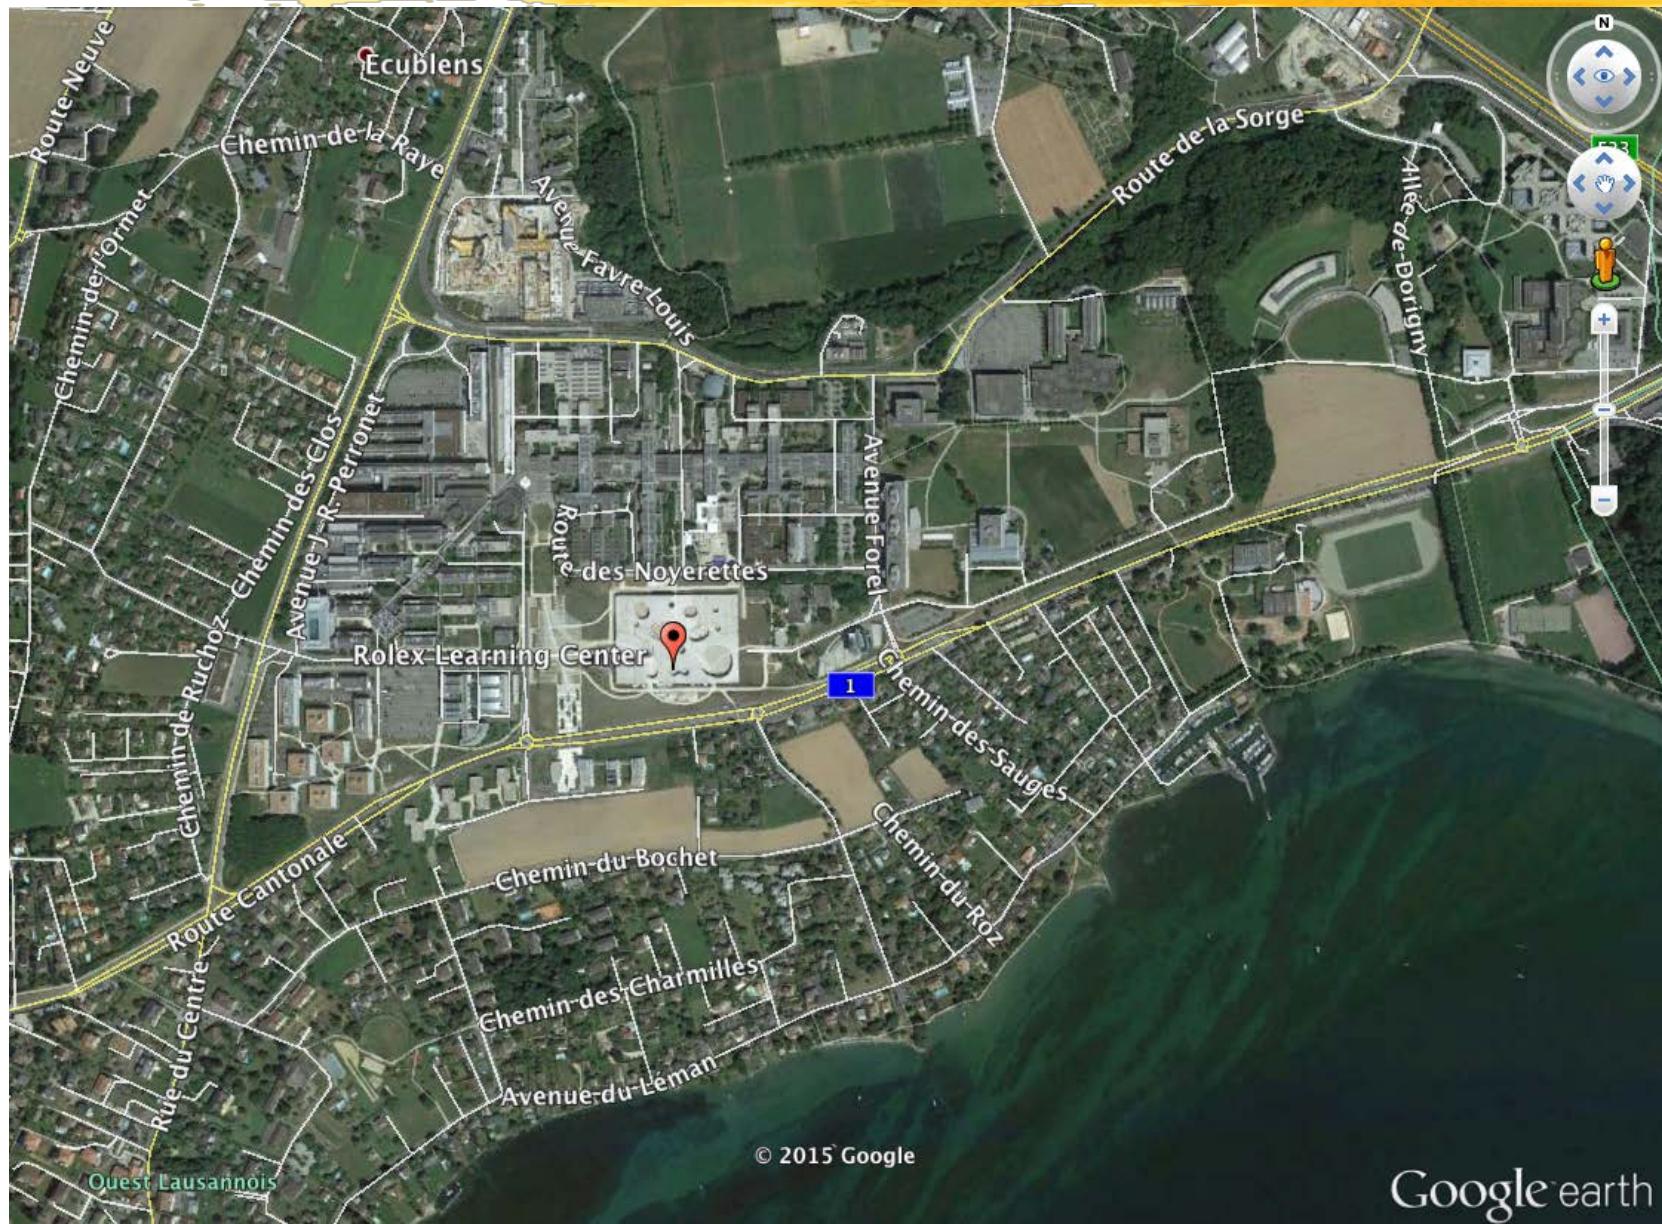
- Edges seem fundamental to human perception.
- They form a compressed version of the image.

FROM EDGES TO CATS



<https://affinelayer.com/pixsrv/>

MAPS AND OVERLAYS



© 2015 Google

Google earth

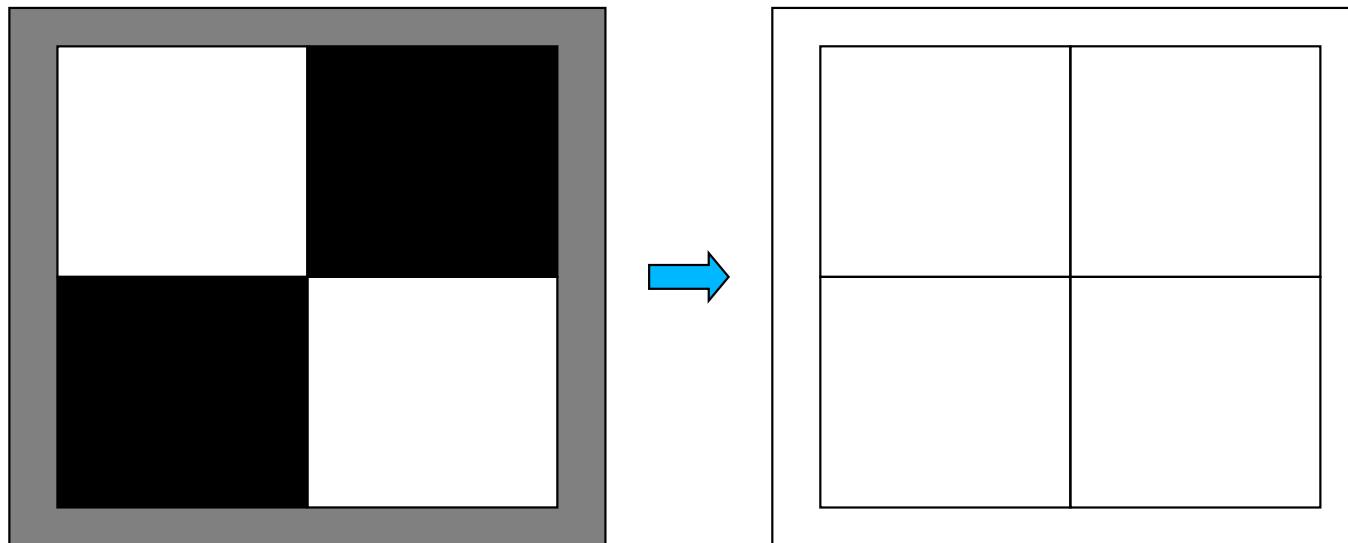
CORRIDOR



CORRIDOR



EDGES AND REGIONS



Edges:

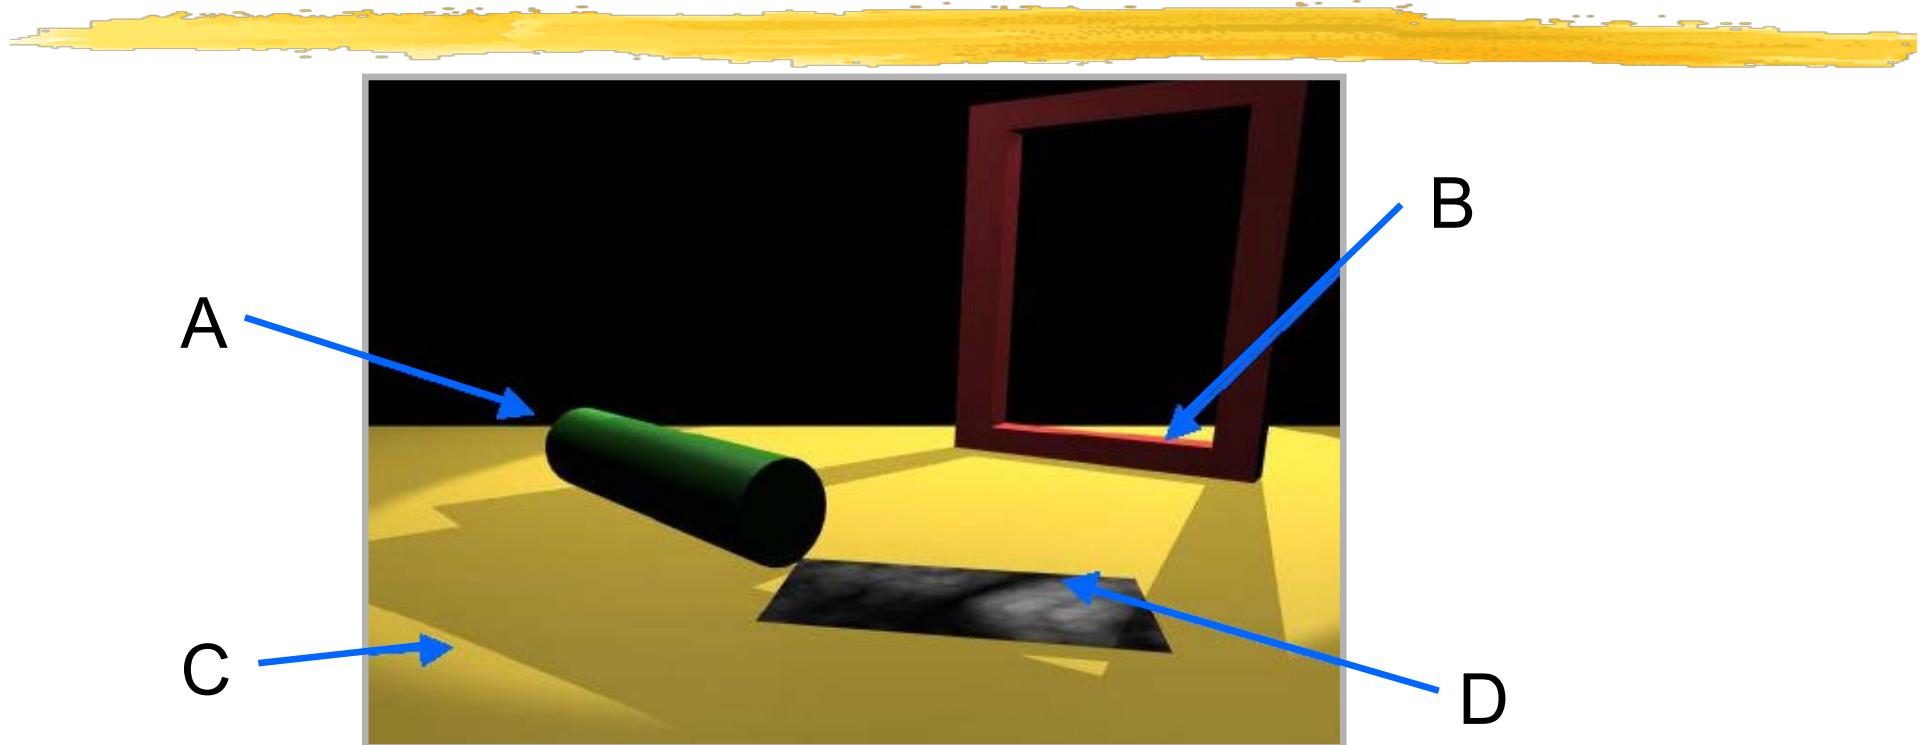
Boundary between bland image regions.

Regions:

Homogenous areas between edges.

→ Duality Edge/Regions.

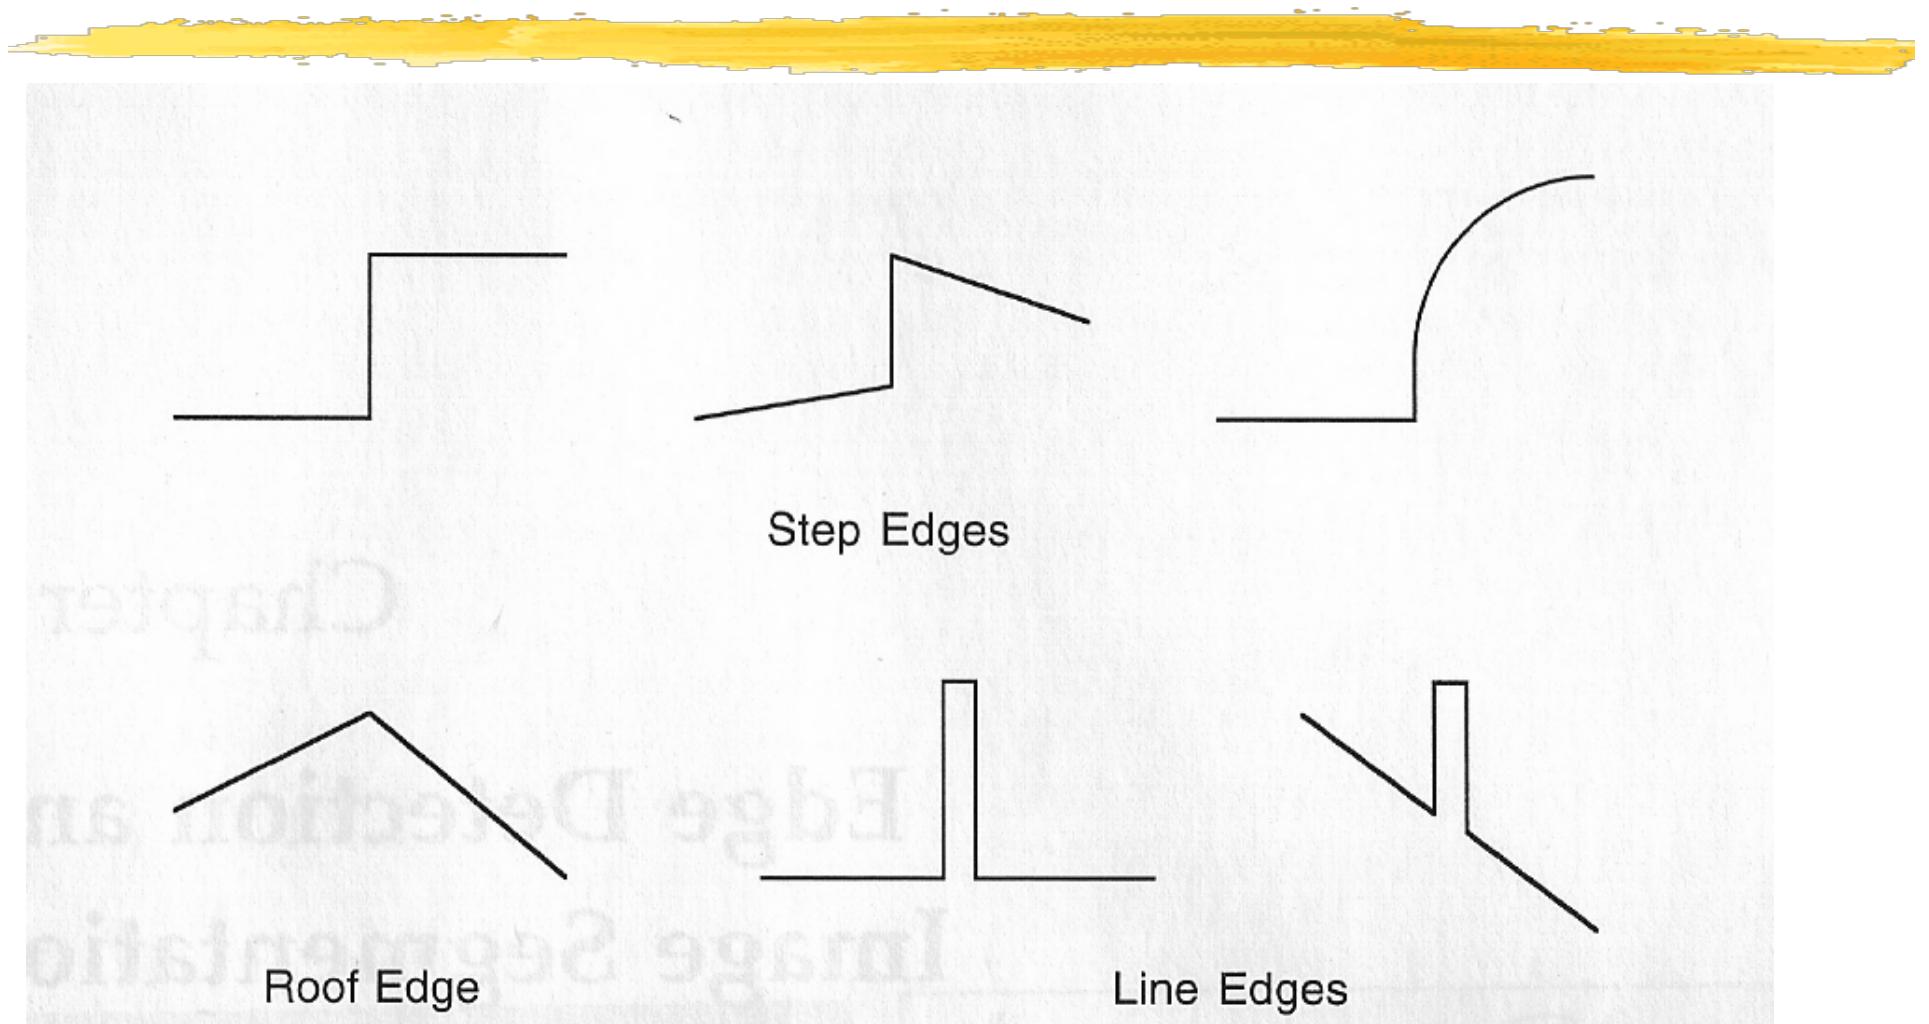
DISCONTINUITIES



- A. Depth discontinuity: Abrupt depth change in the world
- B. Surface normal discontinuity: Change in surface orientation
- C. Illumination discontinuity: Shadows, lighting changes
- D. Reflectance discontinuity: Surface properties, markings

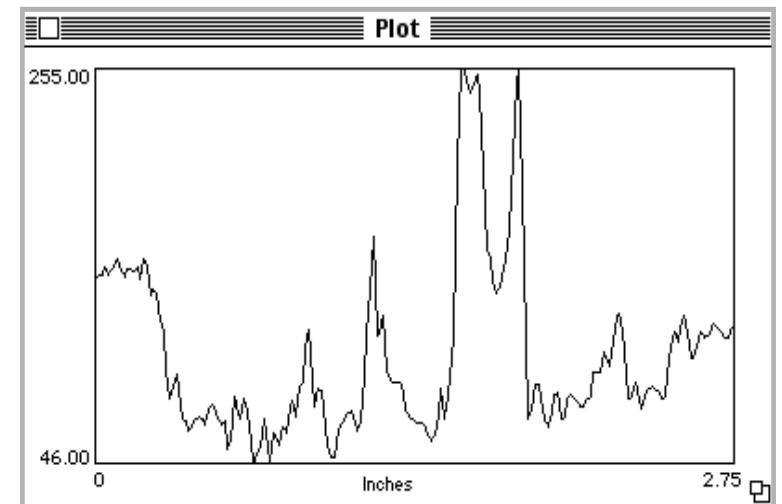
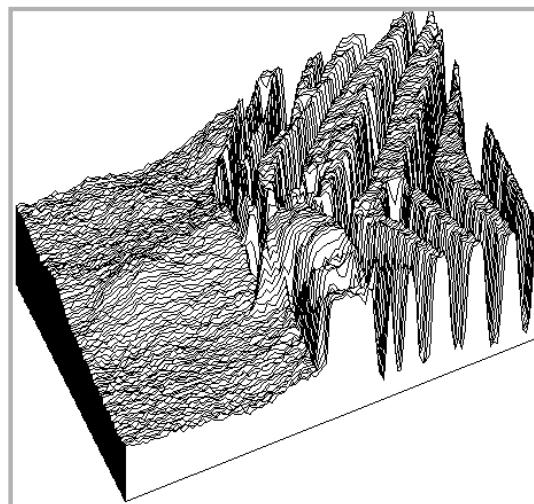
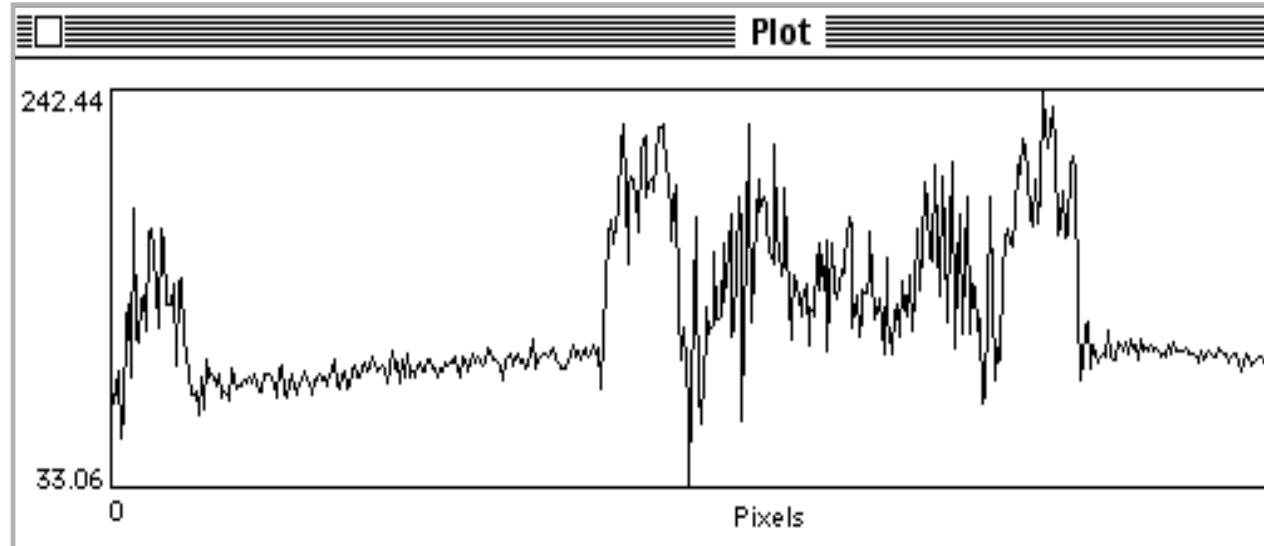
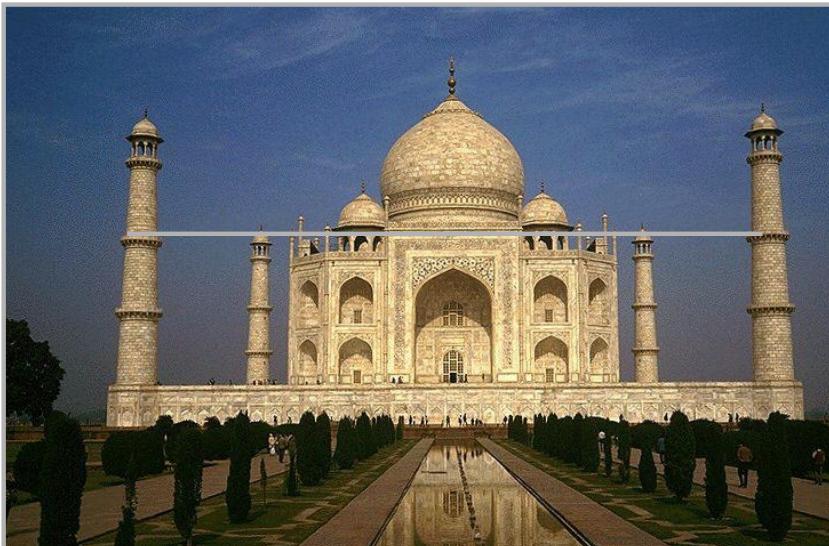
→ **Contrast:** Gray levels different on both sides

EDGE PROFILES

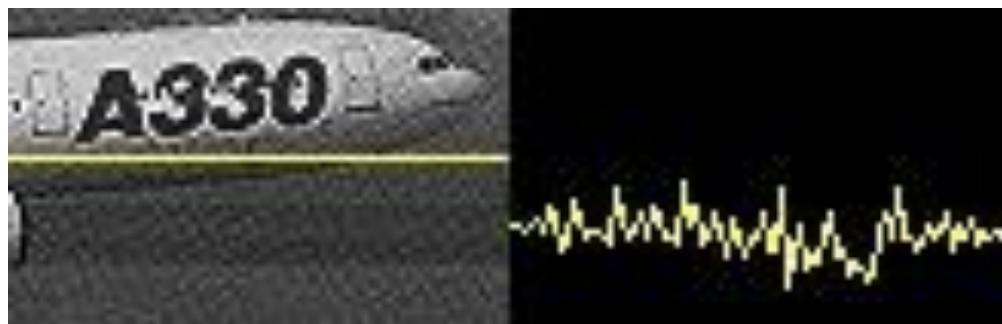
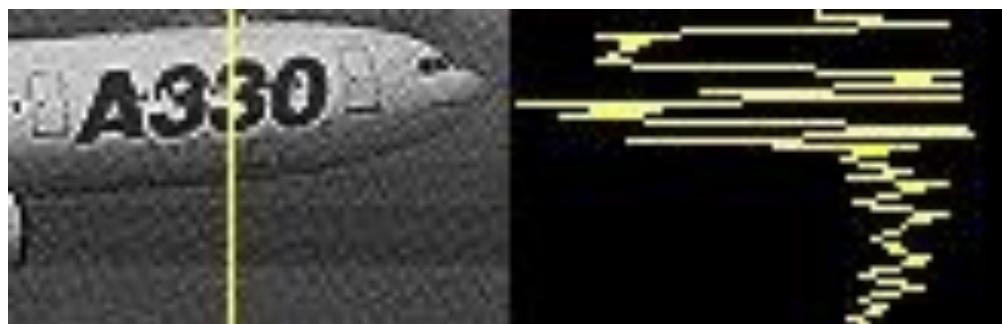


Edges are where a change occurs

REALITY

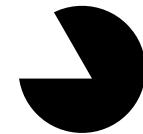


MORE REALITY



Very noisy signals
→ Much knowledge is required!!

ILLUSORY CONTOURS

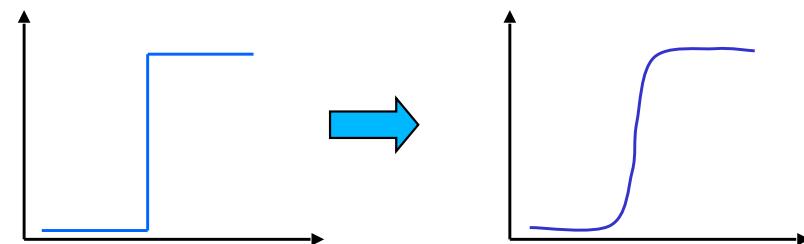


No closed contour, but

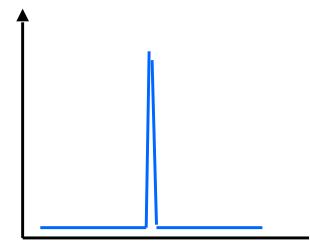
IDEAL STEP EDGE



$f(x) = \text{step edge}$

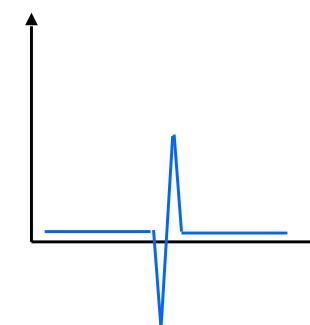


1st Derivative $f'(x)$



maximum

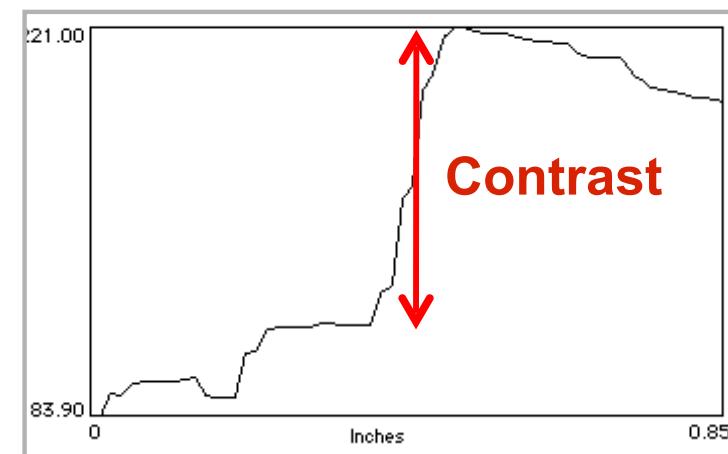
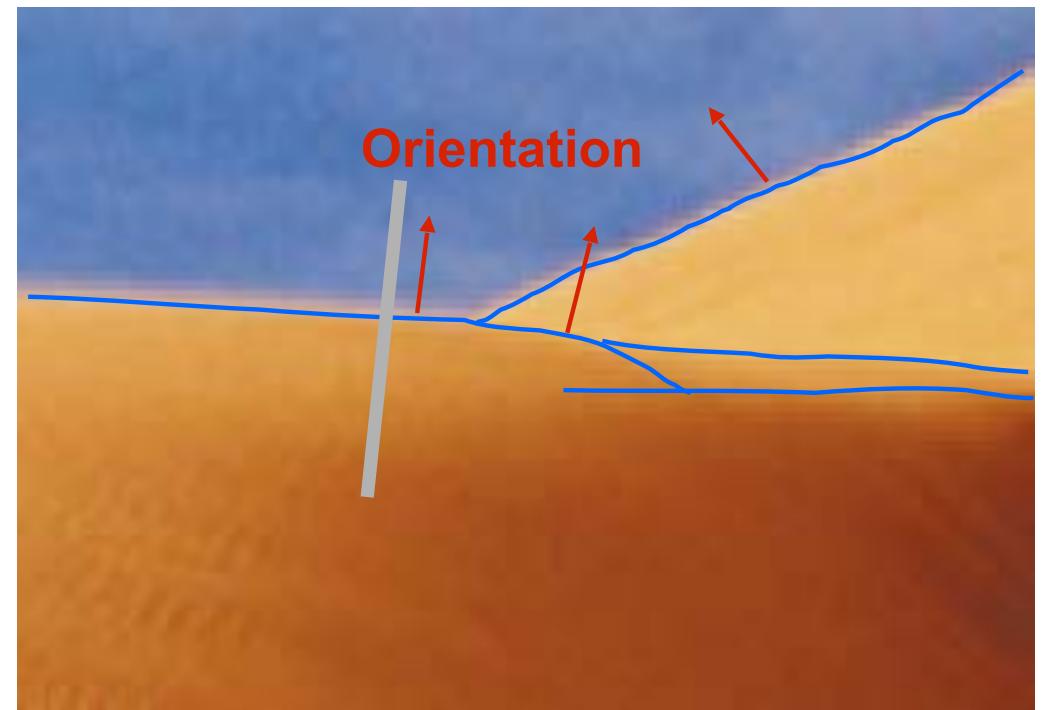
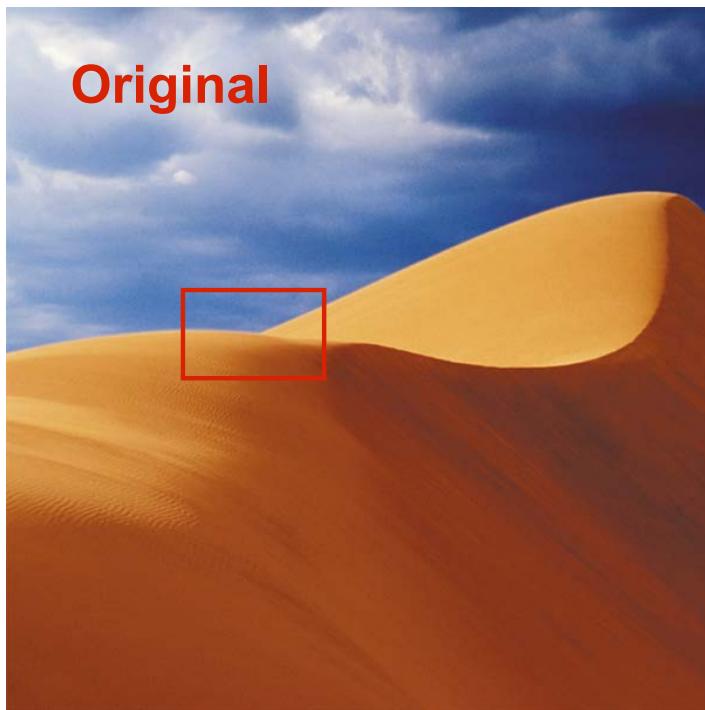
2nd Derivative $f''(x)$



zero crossing

Rapid change in image => High local gradient

EDGE PROPERTIES



EDGE DESCRIPTORS

Edge Normal:

- Unit vector in the direction of maximum intensity change

Edge Direction:

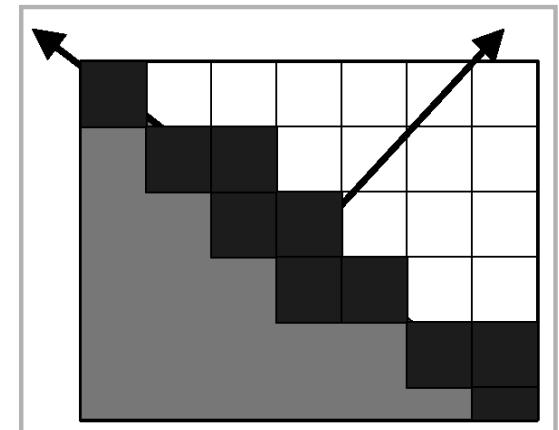
- Unit vector perpendicular to the edge normal

Edge position or center

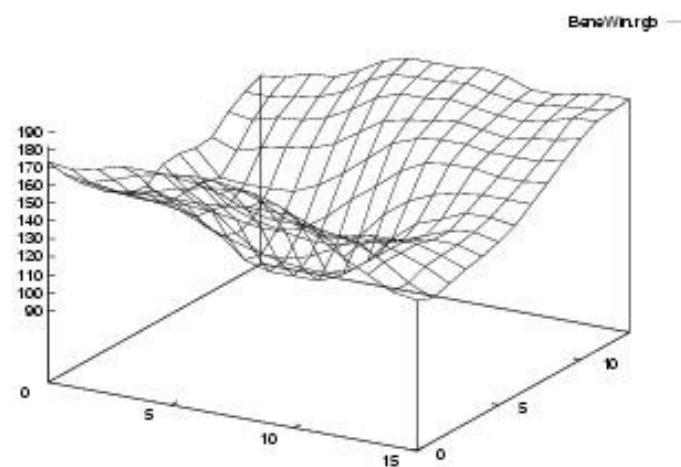
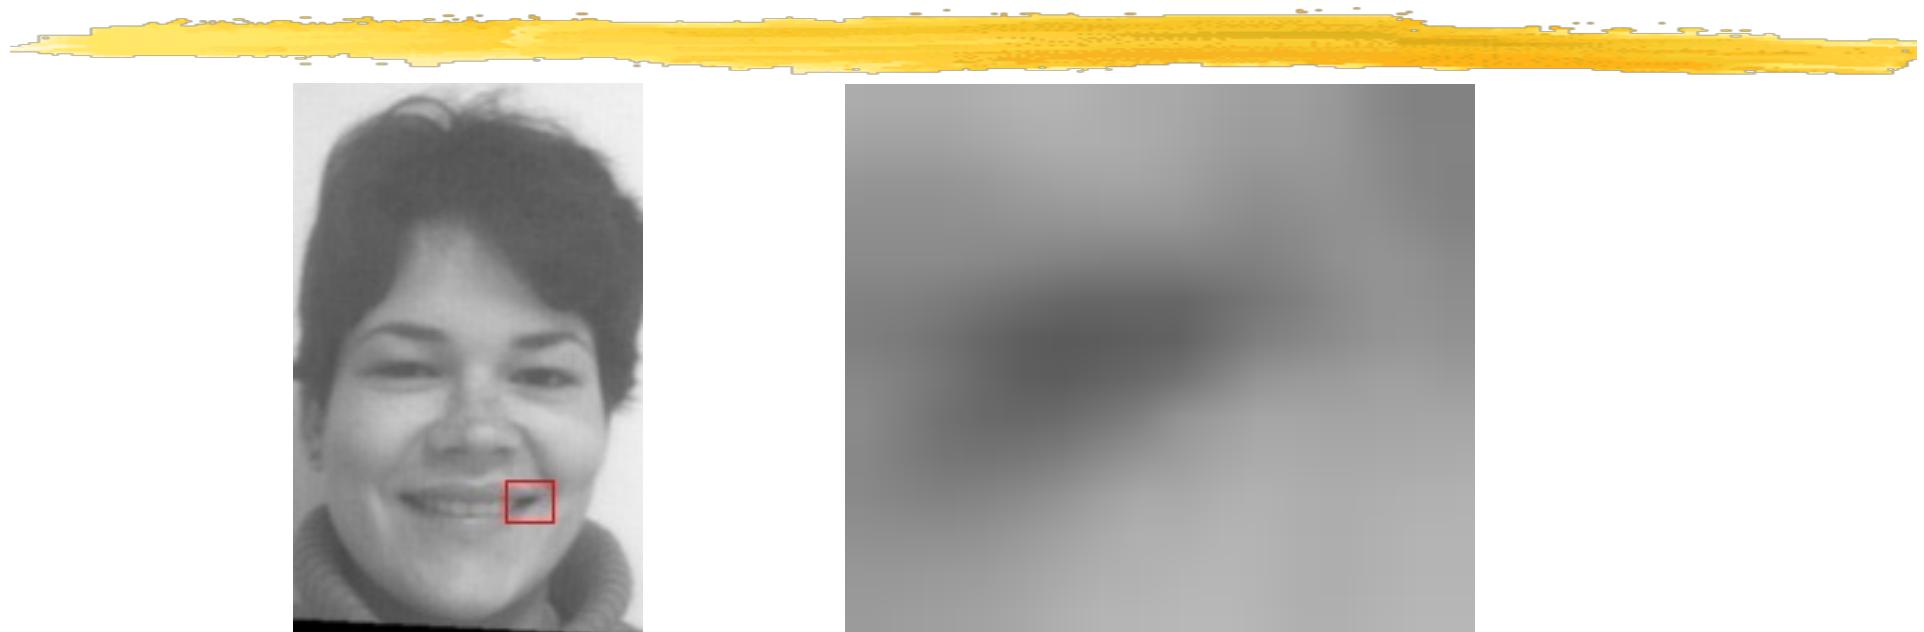
- Image location at which edge is located

Edge Strength

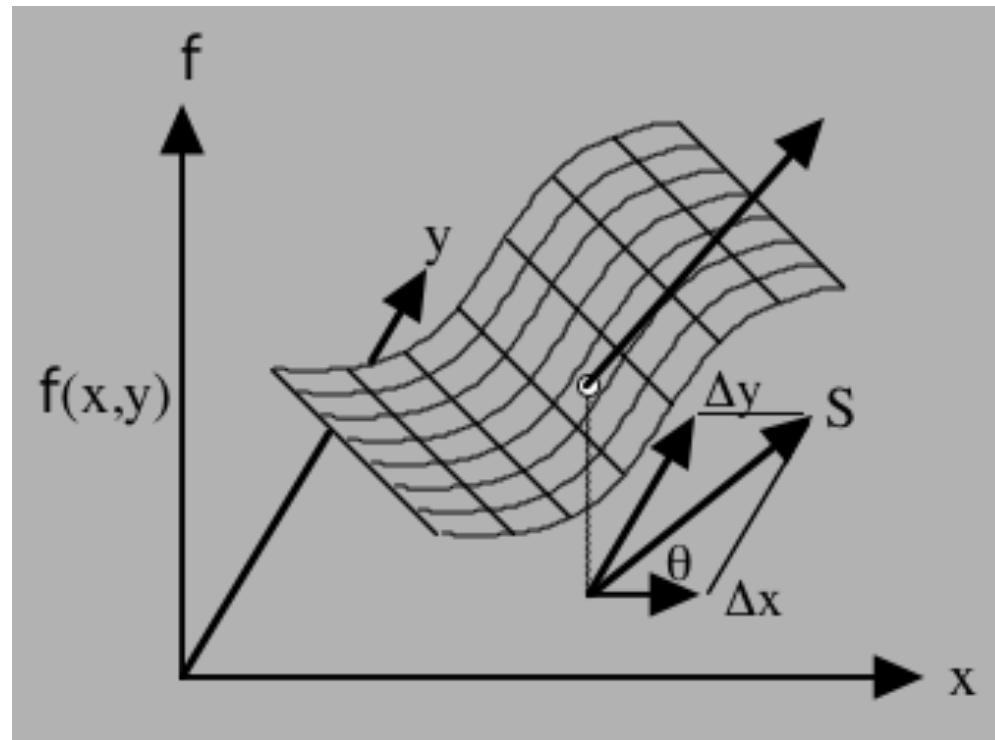
- Speed of intensity variation across the edge.



IMAGES AS 3-D SURFACES



GEOMETRIC INTERPRETATION



Since $I(x,y)$ is not a continuous function:

1. Locally fit a smooth surface.
2. Compute its derivatives.

IMAGE GRADIENT

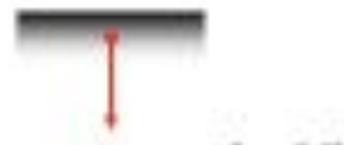
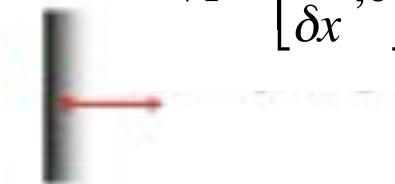


The gradient of an image

$$\nabla I = \left[\frac{\delta I}{\delta x}, \frac{\delta I}{\delta y} \right]$$

points in the direction of most rapid change in intensity.

$$\nabla I = \left[\frac{\delta I}{\delta x}, 0 \right]$$

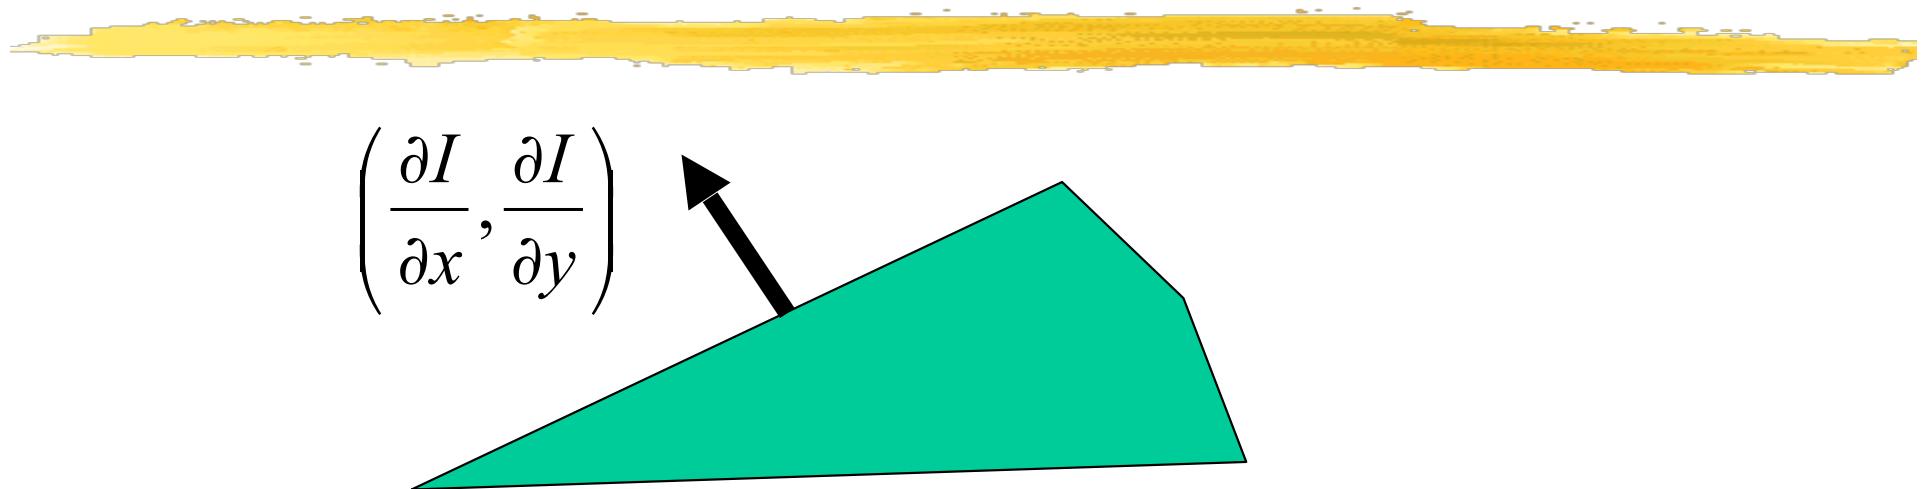


$$\nabla I = \left[\frac{\delta I}{\delta x}, \frac{\delta I}{\delta y} \right]$$

$$\nabla I = \left[0, \frac{\delta I}{\delta y} \right]$$



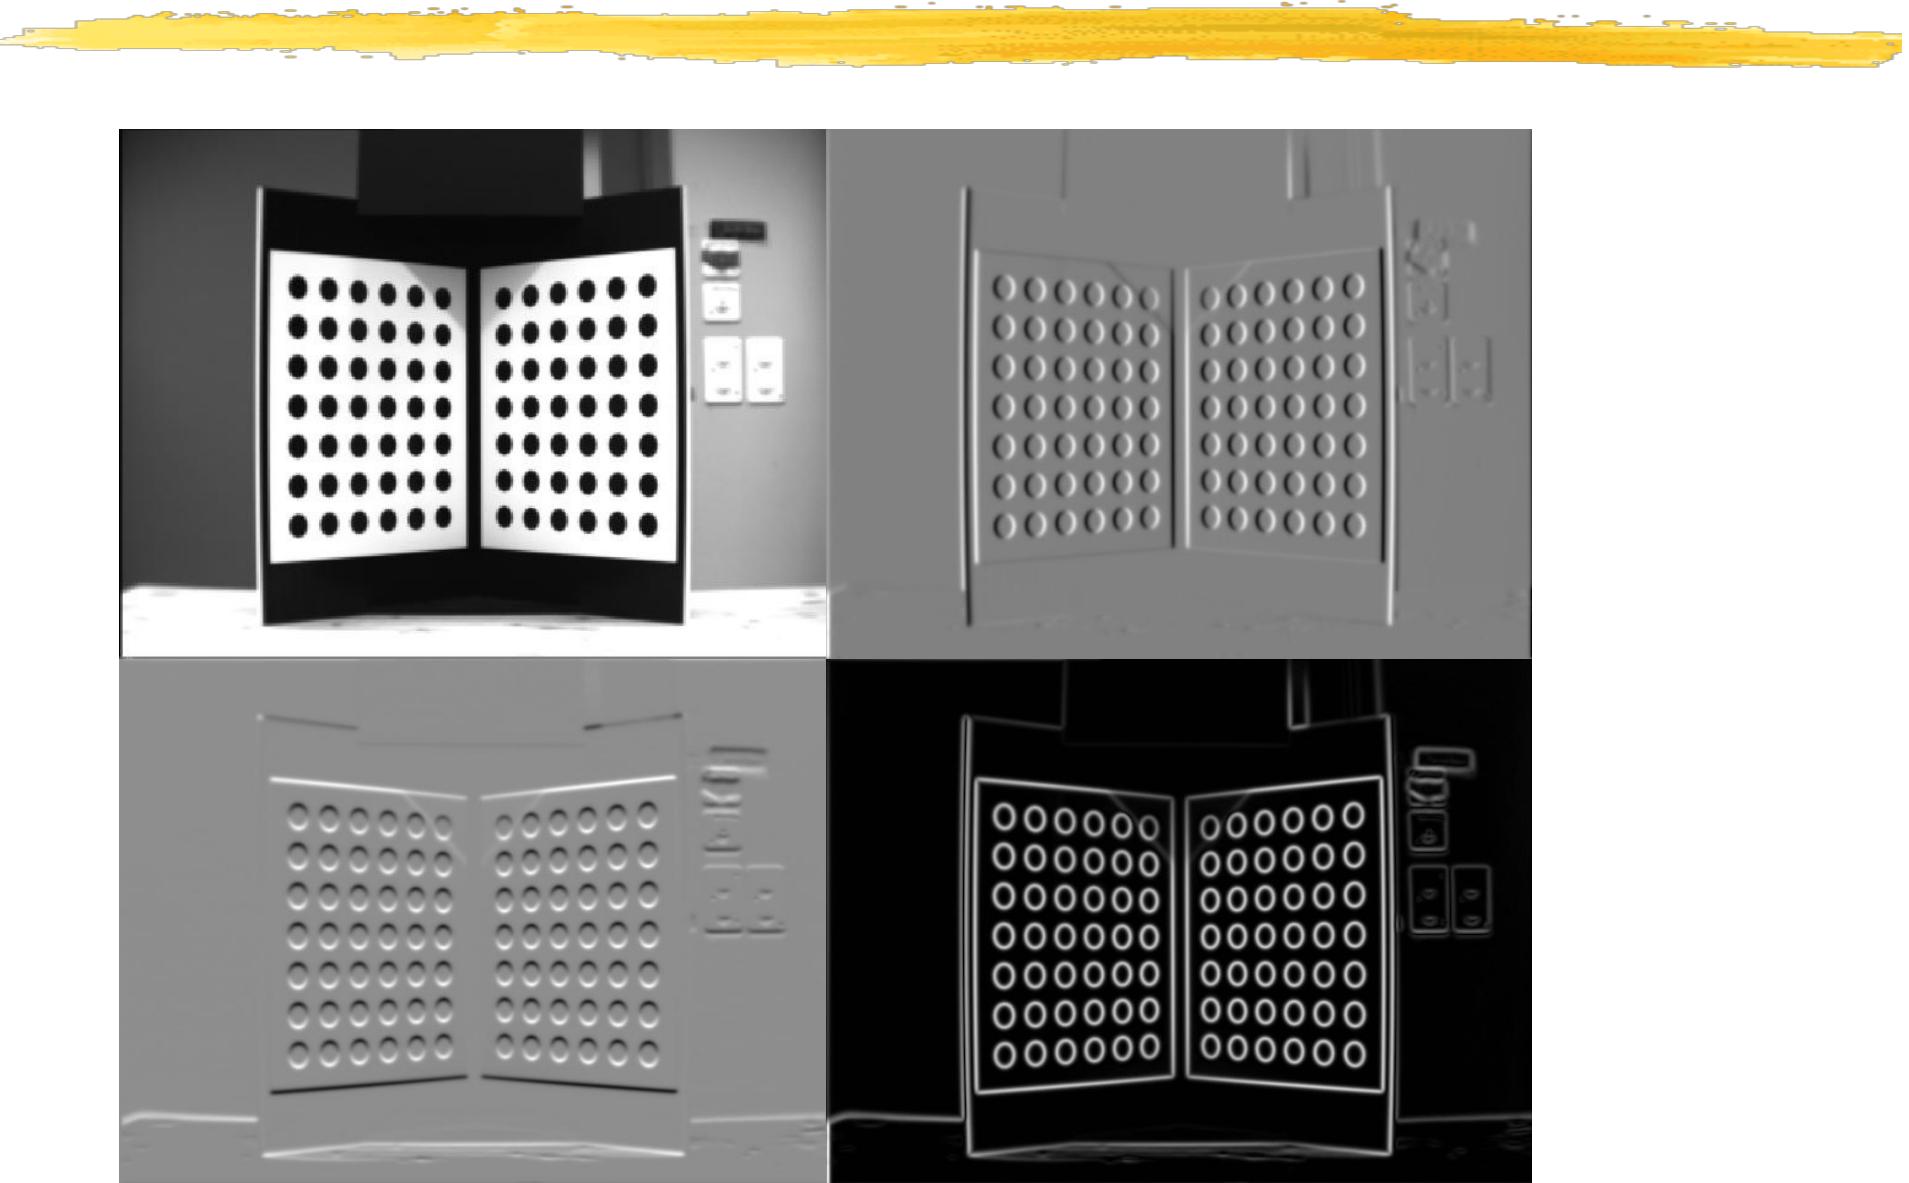
MAGNITUDE AND ORIENTATION



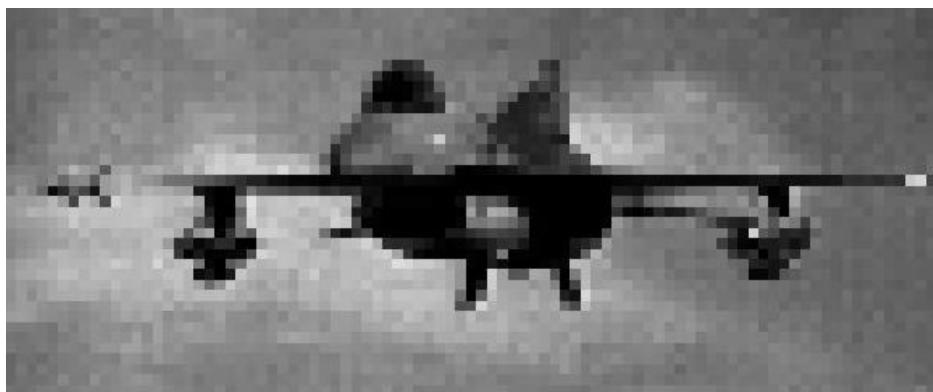
$$\text{Measure of contrast : } G = \sqrt{\frac{\partial I^2}{\partial x} + \frac{\partial I^2}{\partial y}}$$

$$\text{Edge orientation : } \theta = \arctan\left(\frac{\partial I}{\partial y}, \frac{\partial I}{\partial x}\right)$$

GRADIENT IMAGES



REAL IMAGES

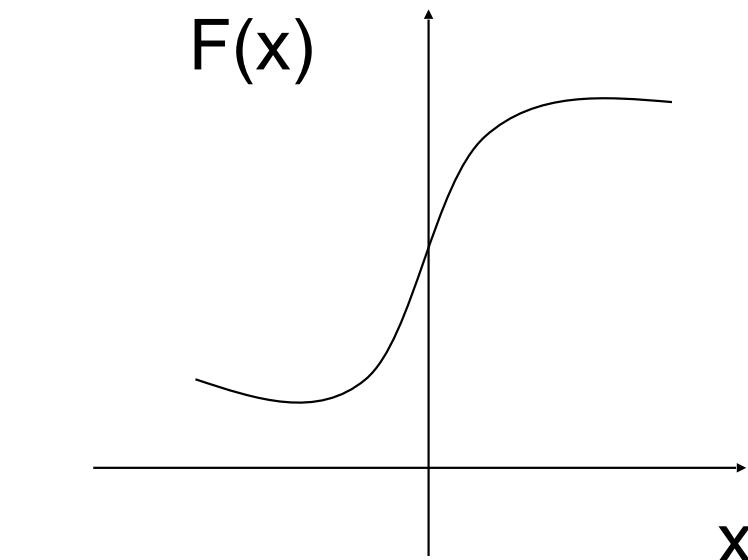


EDGE OPERATORS

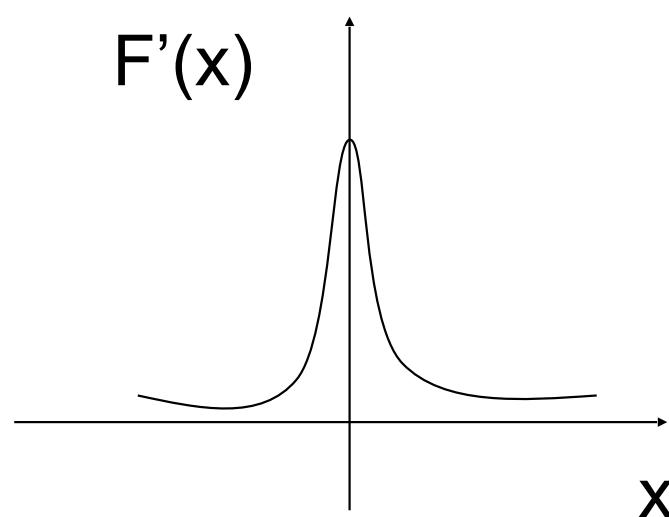


- Difference Operators
- Convolution Operators
- Parametric Matchers
- Trained Detectors
- Deep Nets

GRADIENT METHODS



Edge = Sharp variation

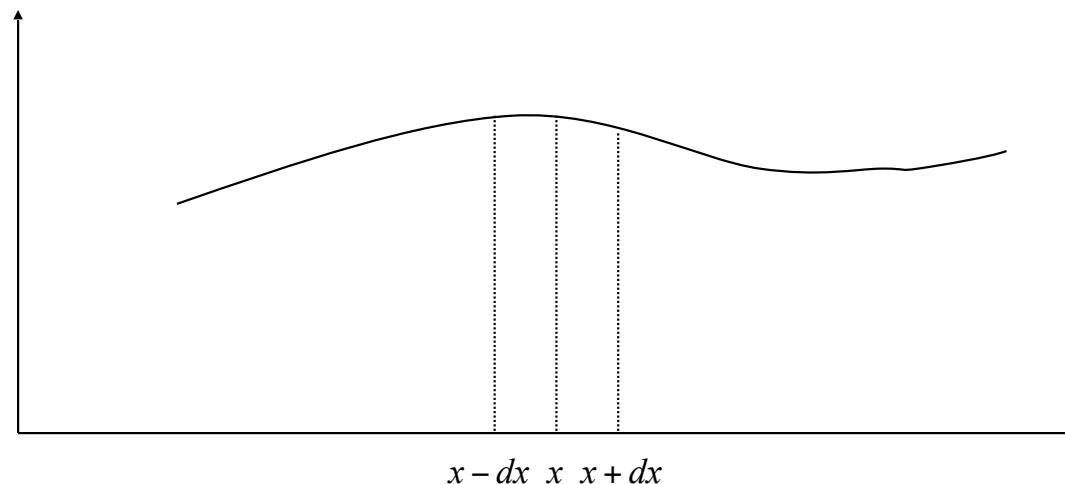


Large first derivative



1D FINITE DIFFERENCES

In one dimension:



$$\frac{df}{dx} \approx \frac{f(x + dx) - f(x)}{dx} \approx \frac{f(x + dx) - f(x - dx)}{2dx}$$

$$\frac{d^2f}{dx^2} \approx \frac{f(x + dx) - 2f(x) + f(x - dx)}{dx^2}$$

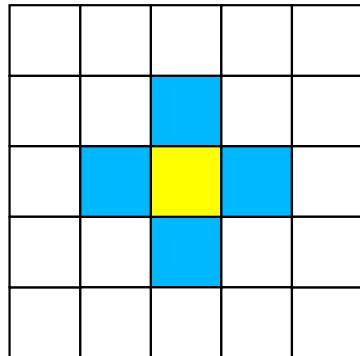
1D FINITE DIFFERENCES IN C



Line stored as an array:

```
{  
    int i;  
    for(i=0;i<n;i++){  
        q[i]=p[i+1]-p[i];  
    }  
}
```

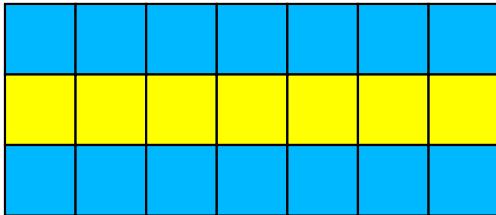
2D FINITE DIFFERENCES



$$\frac{\partial f}{\partial x} \approx \frac{f(x + dx, y) - f(x, y)}{dx} \approx \frac{f(x + dx, y) - f(x - dx, y)}{2dx}$$

$$\frac{\partial f}{\partial y} \approx \frac{f(x, y + dy) - f(x, y)}{dy} \approx \frac{f(x, y + dy) - f(x, y - dy)}{2dy}$$

2D FINITE DIFFERENCES IN C



→ p →



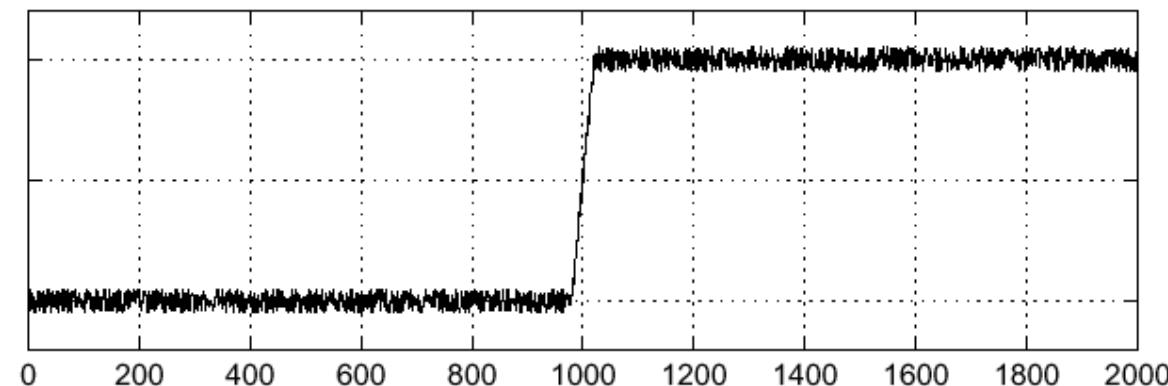
Image stored in raster format:

```
{  
    int i;  
    for(i=0;i<xdim;i++){  
        dx[i] = p[i+1]-p[i];  
        dy[i] = p[i+xdim]-p[i];  
    }  
}
```

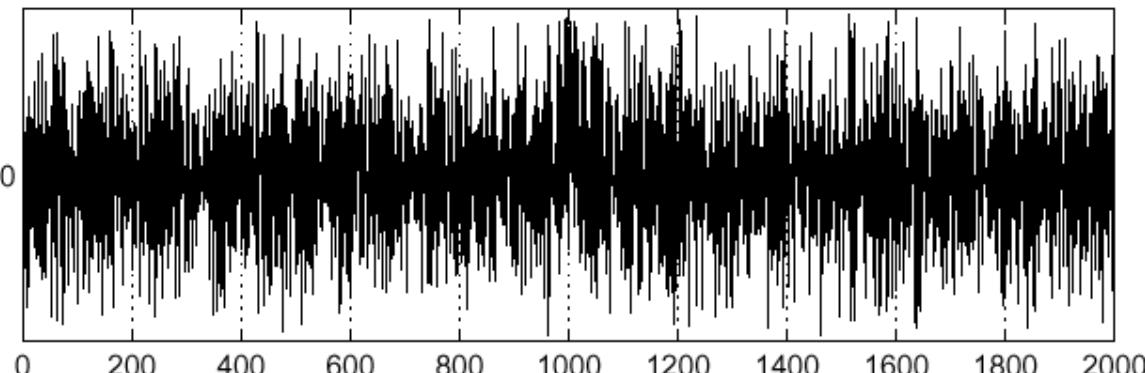
NOISE IN 1D



Consider a single row or column of the image:



$$\frac{d}{dx} f(x)_0$$



FOURIER INTERPRETATION



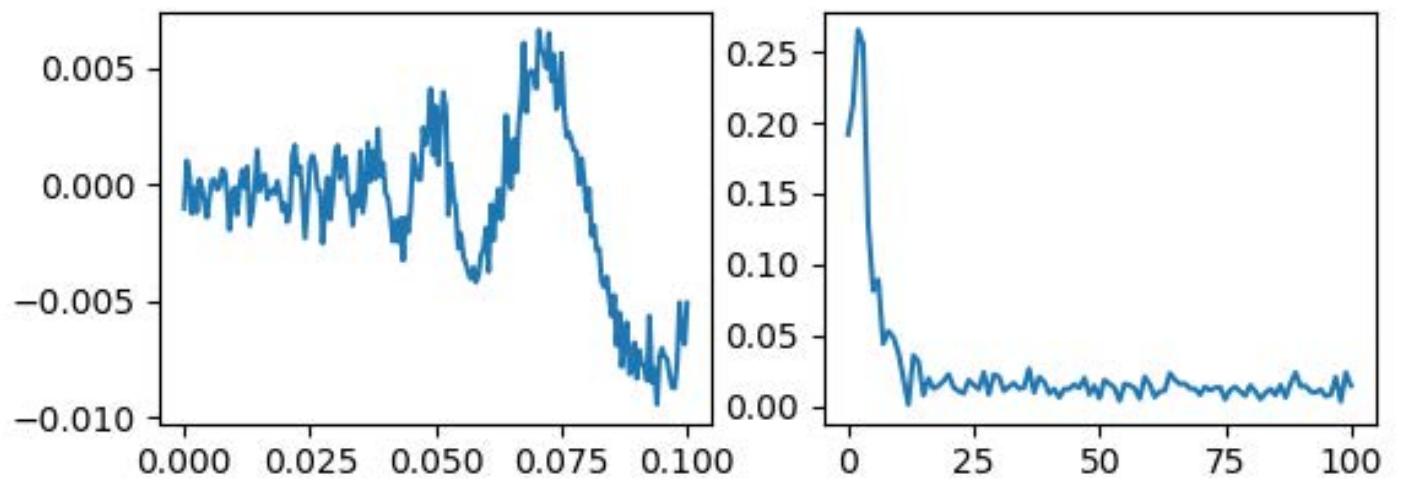
Function	Fourier Transform
$\frac{df}{dx}(x)$	$uF(u)$
$\frac{\delta f}{\delta x}(x, y)$	$uF(u, v)$
$\frac{\delta f}{\delta y}(x, y)$	$vF(u, v)$

→ Differentiating emphasizes high frequencies
and therefore noise!

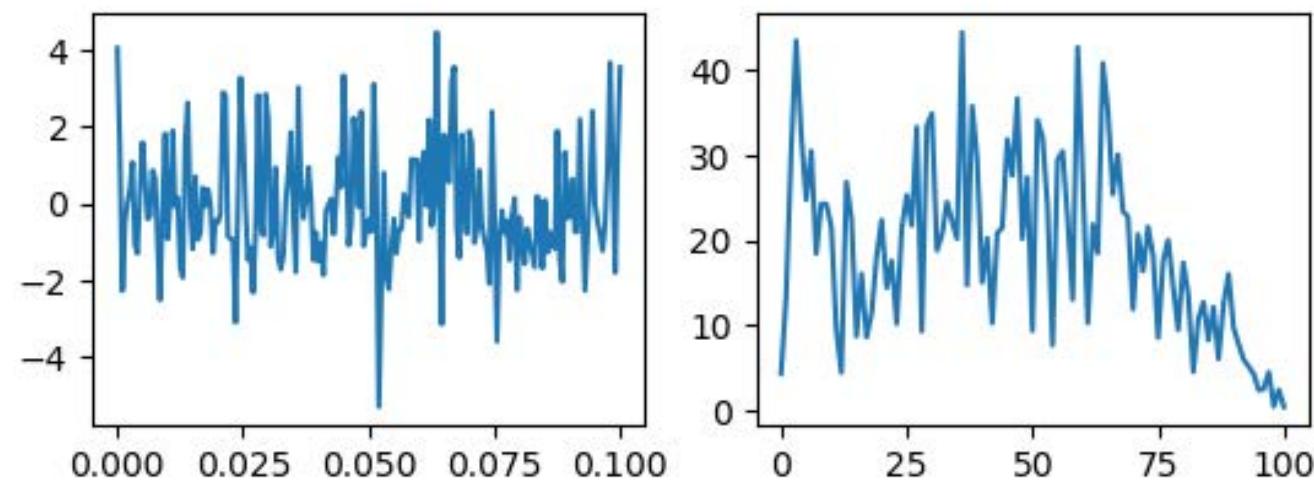
$$f(x) = x^2 \sin(1/x)$$



f



$\frac{df}{dx}$



F

uF

REMOVING NOISE



Problem:

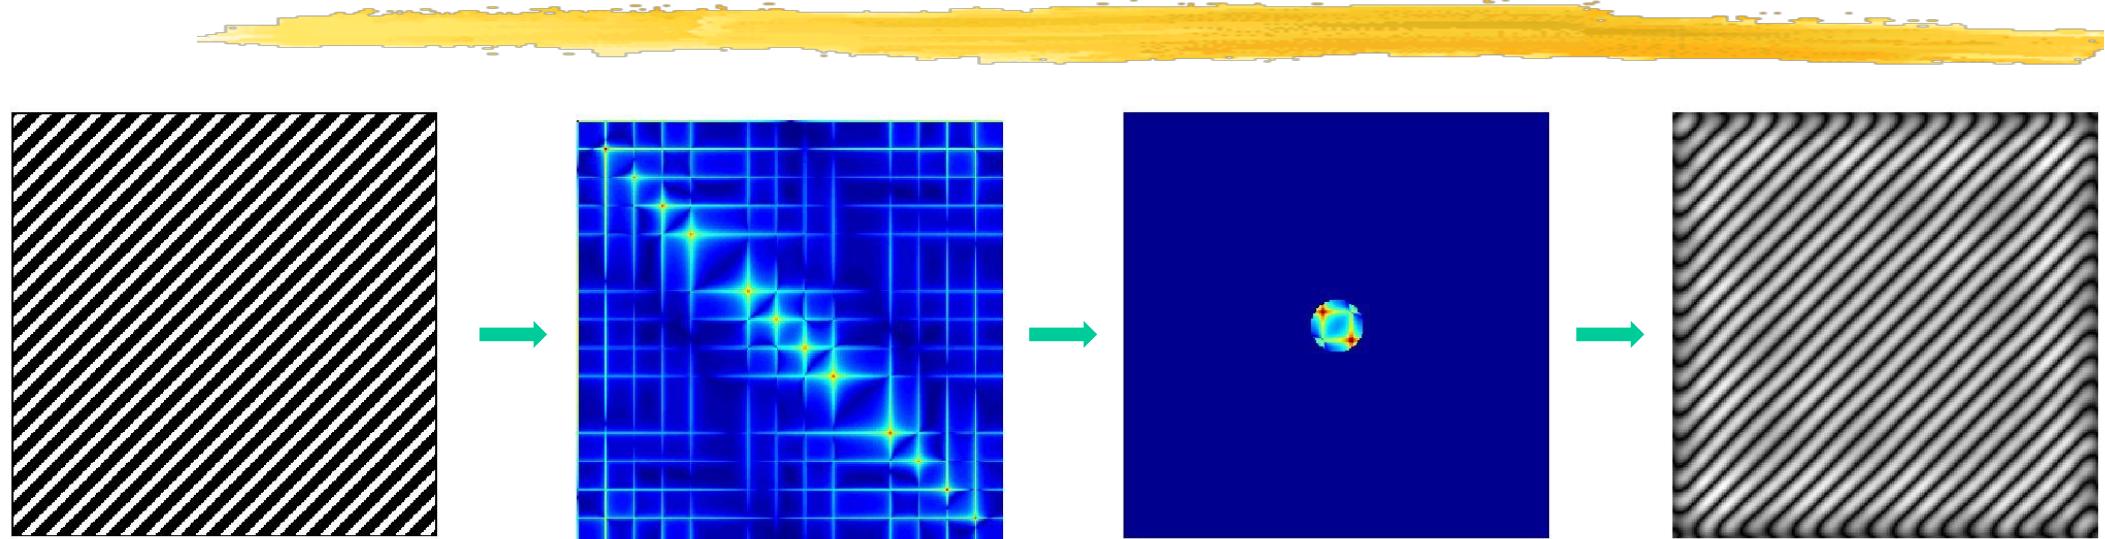
High frequencies lead to trouble with derivation.

Solution:

Suppress high frequencies by

- multiplying DFT of the signal with something that suppresses high frequencies.

DIAGONAL STRUCTURES



Rotated stripes:

- Dominant diagonal structures
- Discretization produces additional harmonics

Removing higher frequencies and reconstructing:

- Smoothed image

REMOVING NOISE



Problem:

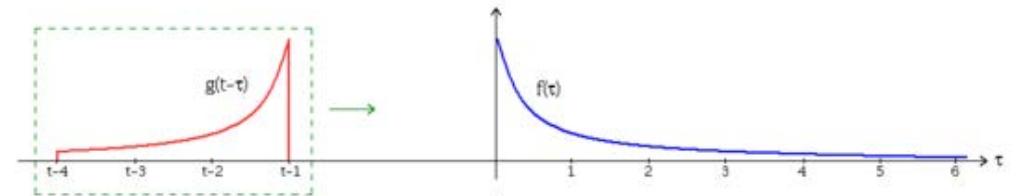
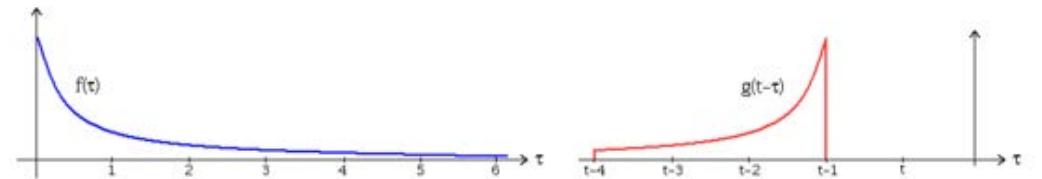
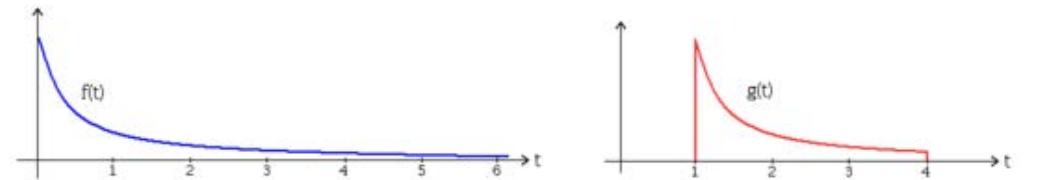
High frequencies lead to trouble with derivation.

Solution:

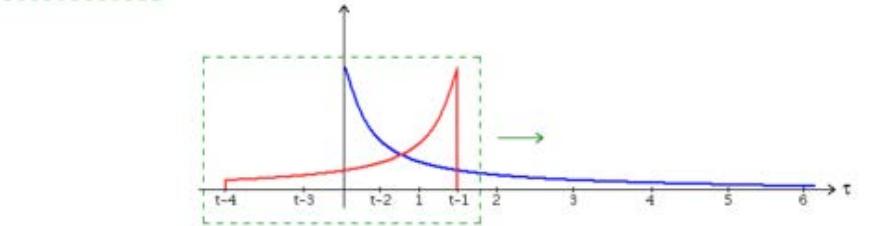
Suppress high frequencies by

- multiplying DFT of the signal with something that suppresses high frequencies;
- convolving with a low-pass filter.

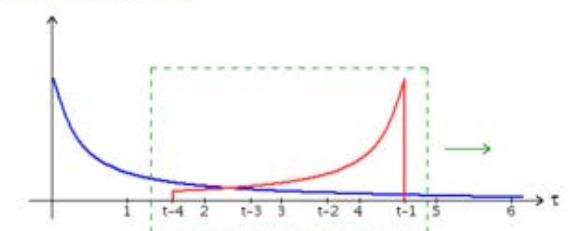
1D CONVOLUTION



$$(f \star g)(t) = \int_{\tau} f(\tau)g(t - \tau)d\tau$$



$$(f \star g)(m) = \sum_n f(n)g(m - n)$$

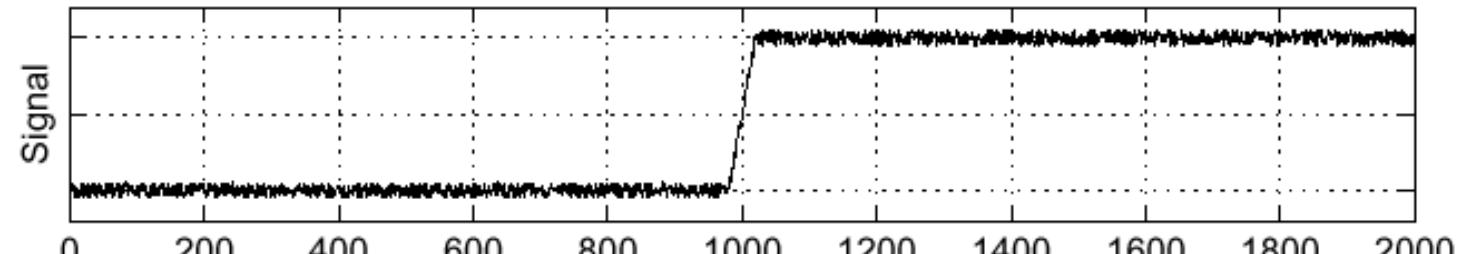


SOLUTION: SMOOTH FIRST

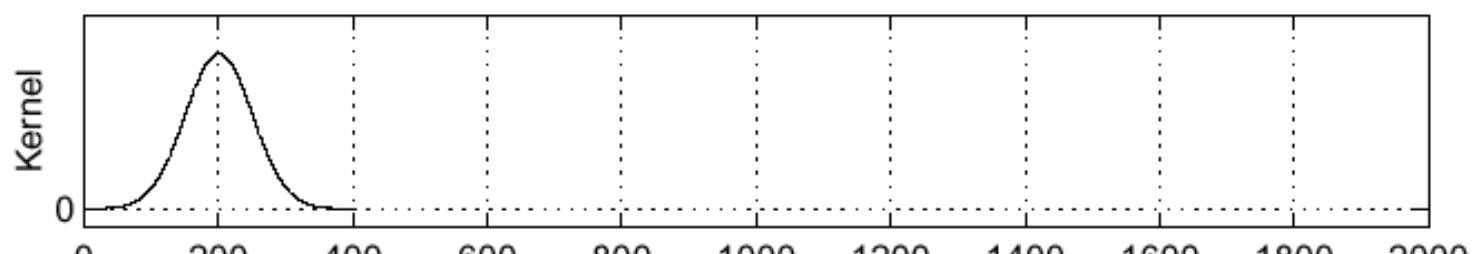


Sigma = 50

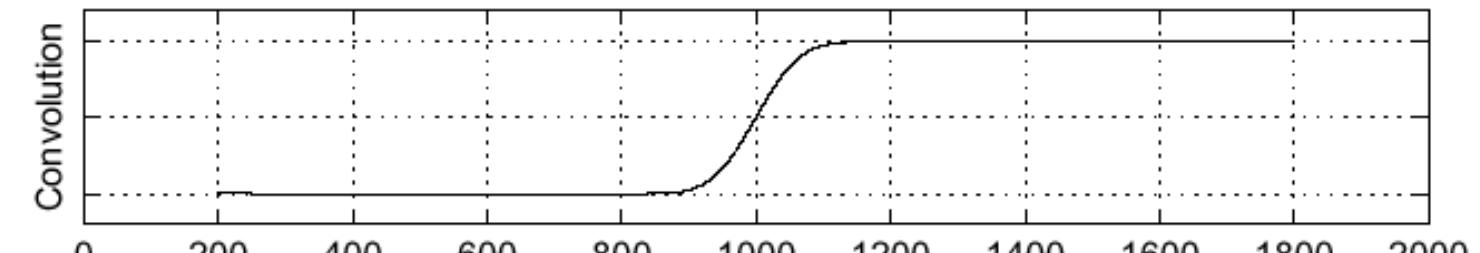
f



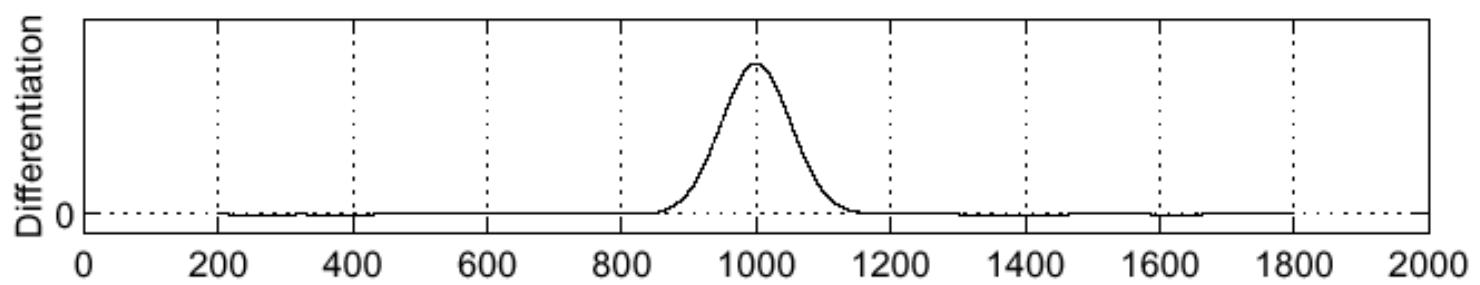
g



$g * f$



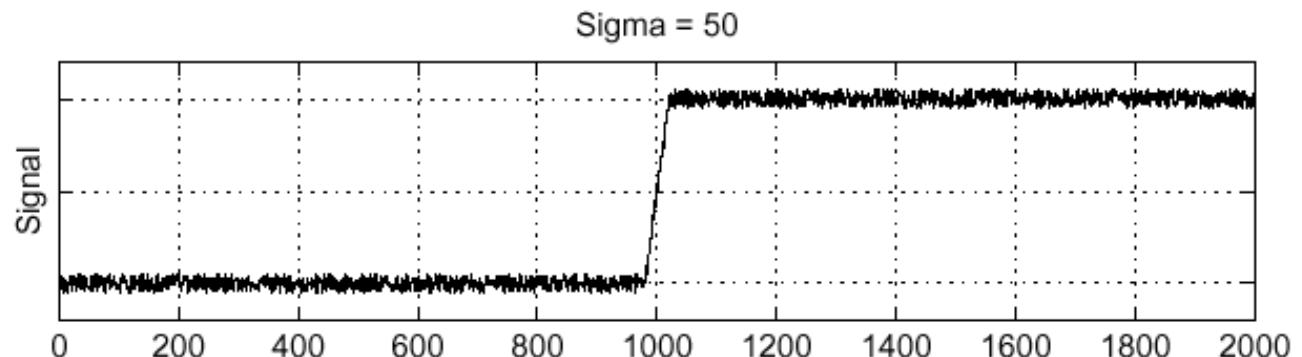
$\frac{\partial}{\partial x}(g * f)$



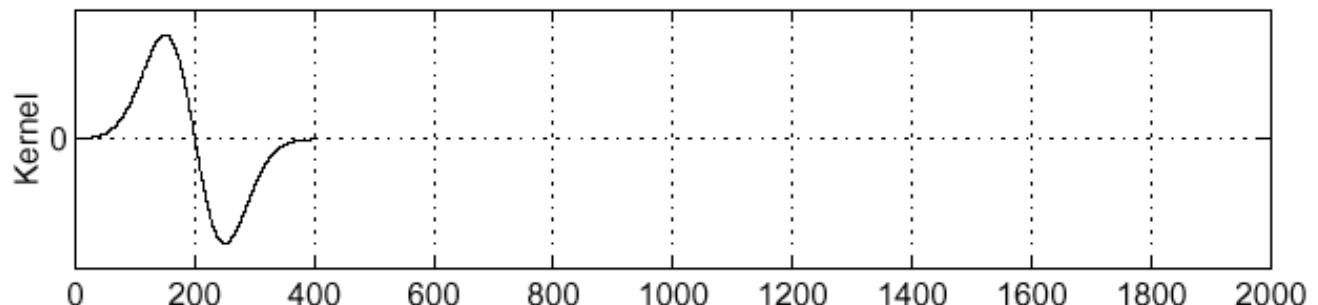
DERIVATIVE THEOREM OF CONVOLUTION



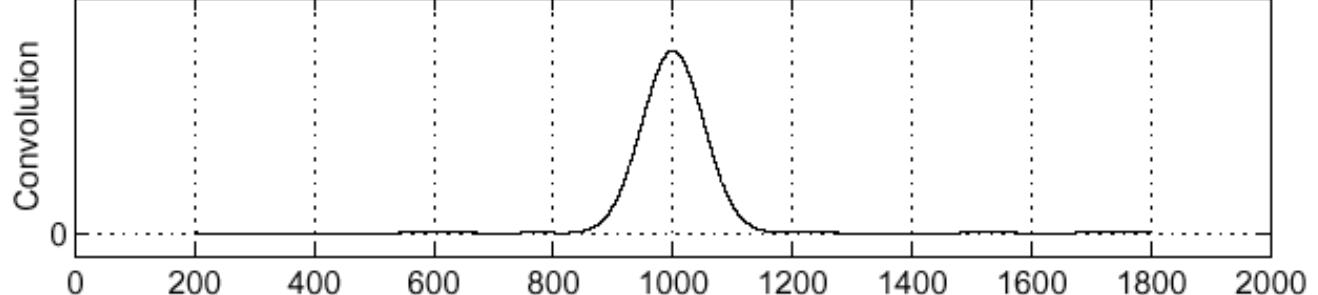
f



$\frac{\partial g}{\partial x}$



$$\frac{\partial}{\partial x} (g * f) = \frac{\partial g}{\partial x} * f$$



--> Faster because dg/dx can be precomputed.

1D AND 2D CONVOLUTION

Continuous case:

$$m \bullet f(x) = \int_u m(u) f(x - u) du$$

$$m \bullet f(x, y) = \iint_{u \ v} m(u, v) f(x - u, y - v) du dv$$

Discrete case:

$$m \bullet f(x) = \sum_{i=-w}^w m(i) f(x - i)$$

$$m \bullet f(x, y) = \sum_{i=-w}^w \sum_{j=-h}^h m(i, j) f(x - i, y - j)$$

CONVOLUTION IN C

Naive implementation:

```
static double g[][]={{{-1.0,-2.0,-1.0},{0.0,0.0,0.0},{1.0,2.0,1.0}};  
{  
    for(i=i0;i<N;i++)  
        for(j=j0;j<N;j++){  
            q[i][j]=0;  
            for(a=a0;a<W;a++)  
                for(b=b0;b<W,b++)  
                    q[i][j]+=g[a][b]*p[i-a][j-b];  
        }  
}
```

Computational complexity:

- N^2W^2 multiplications for a $N \times N$ image and a $W \times W$ mask.
- Lots of memory access

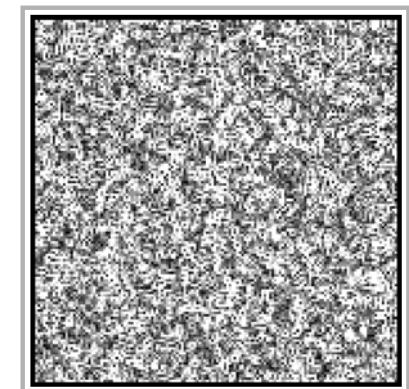
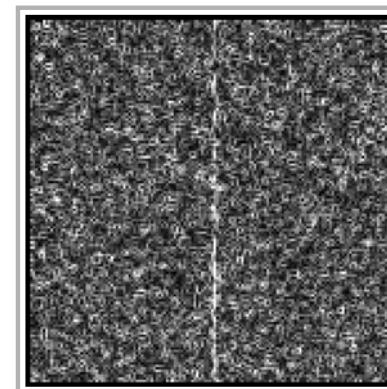
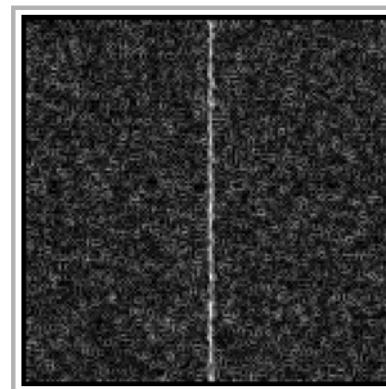
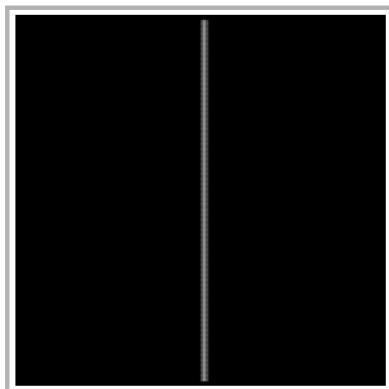
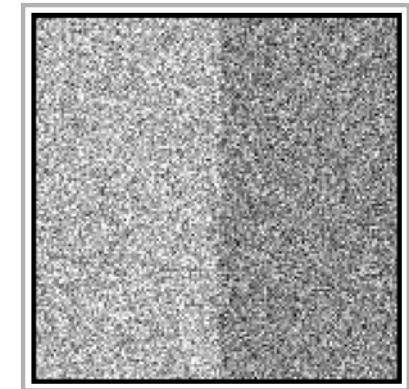
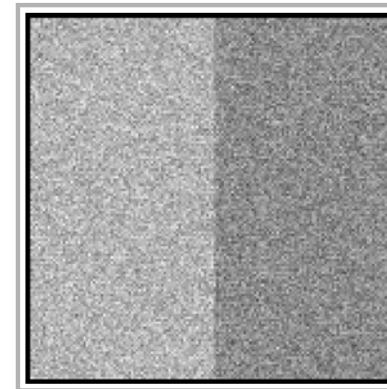
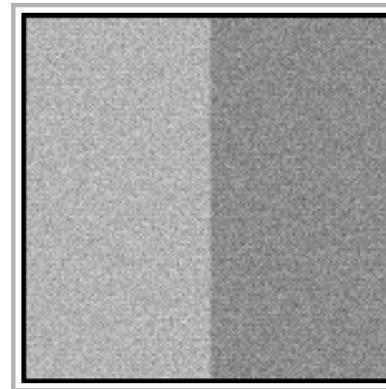
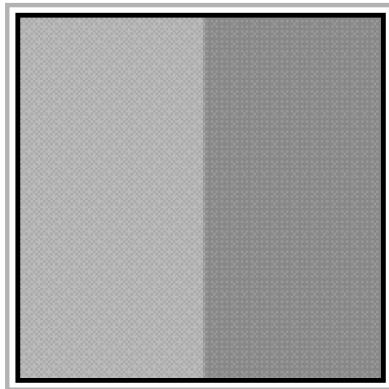
→ Slow, but can be sped up when the filters are separable.

NOISE IN 2D



Ideal step edge

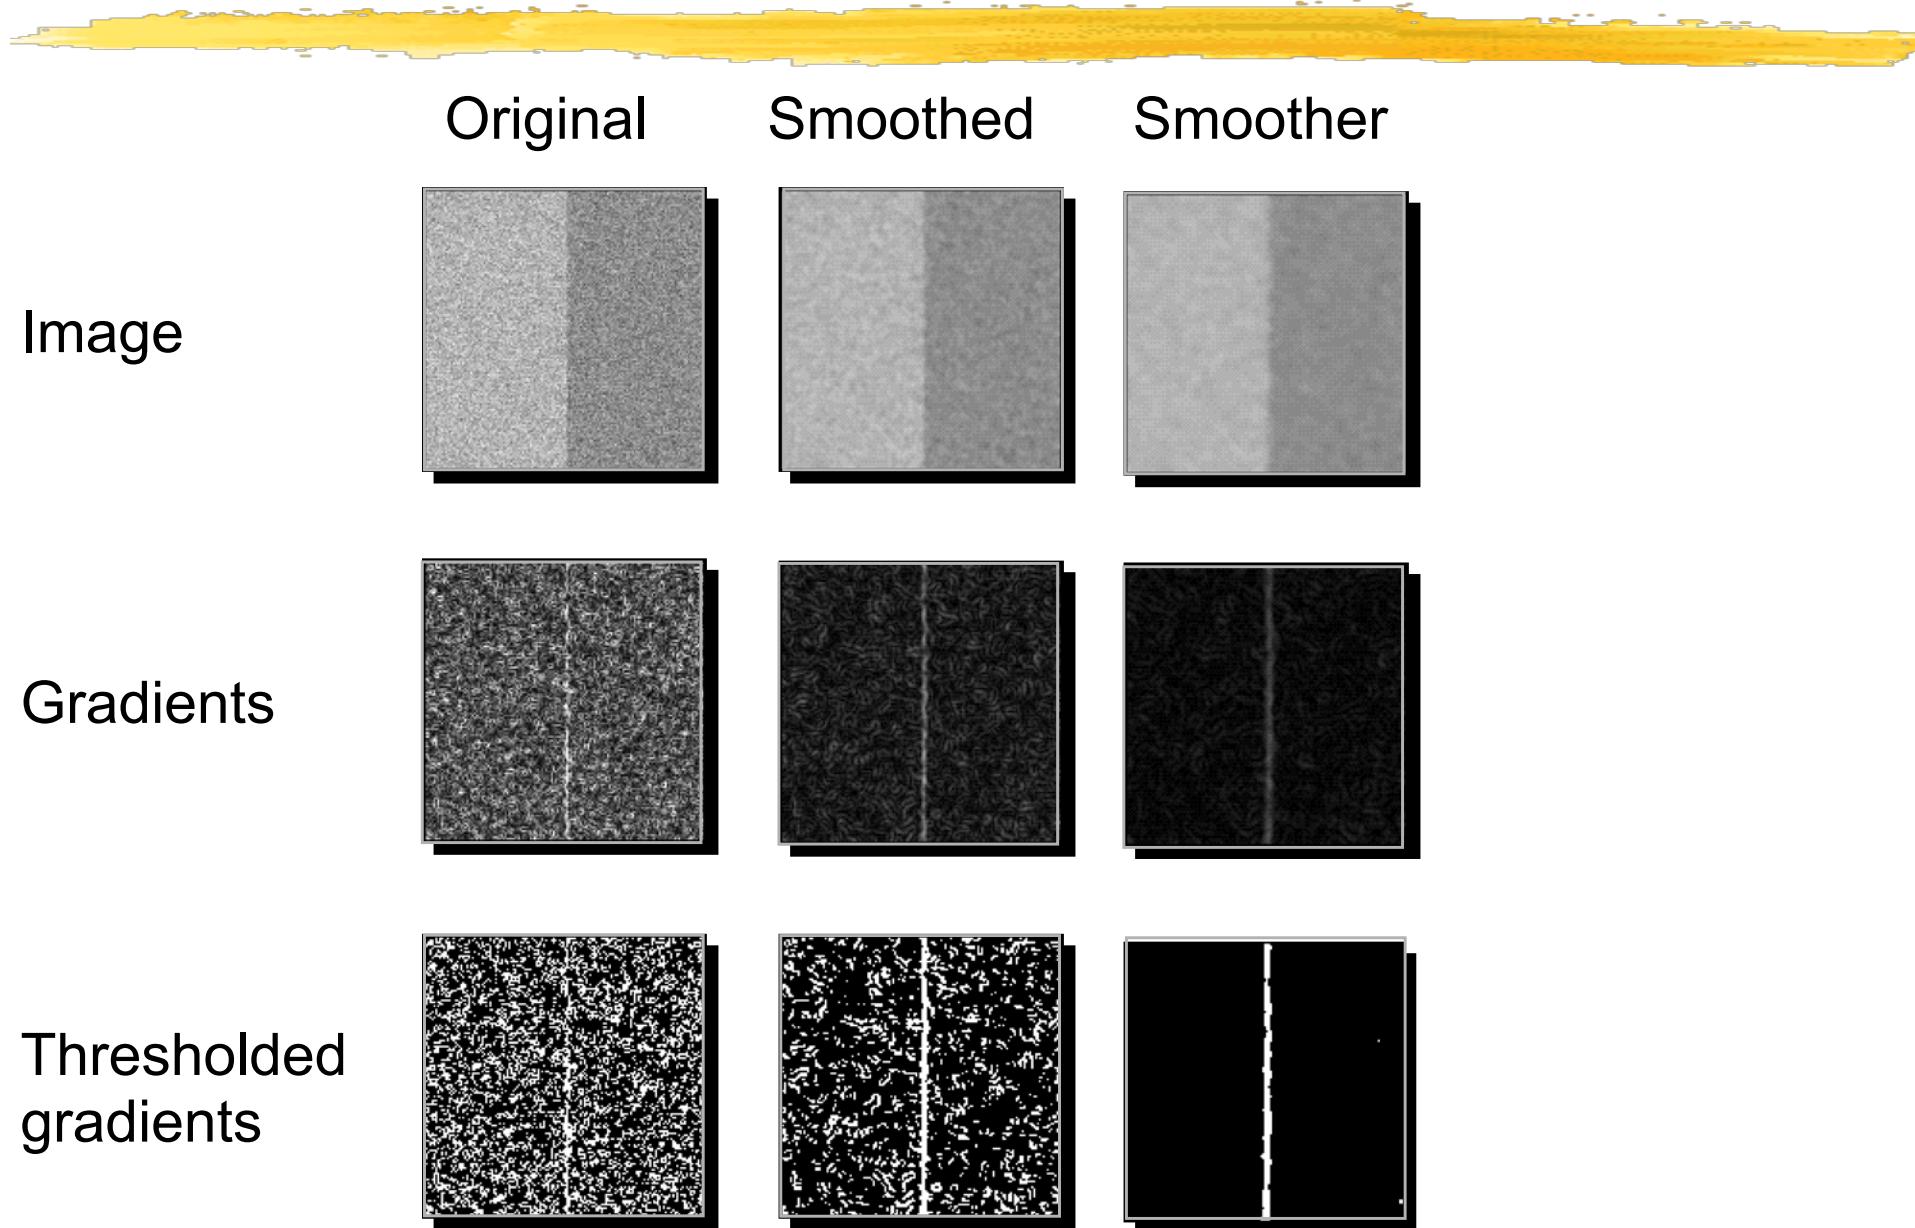
Step edge + noise



Increasing noise



BENEFITS OF SMOOTHING



DIFFERENTIATION AS CONVOLUTION

$$[-1,1] \rightarrow \frac{\partial f}{\partial x} \approx \frac{f(x+dx, y) - f(x, y)}{dx}$$

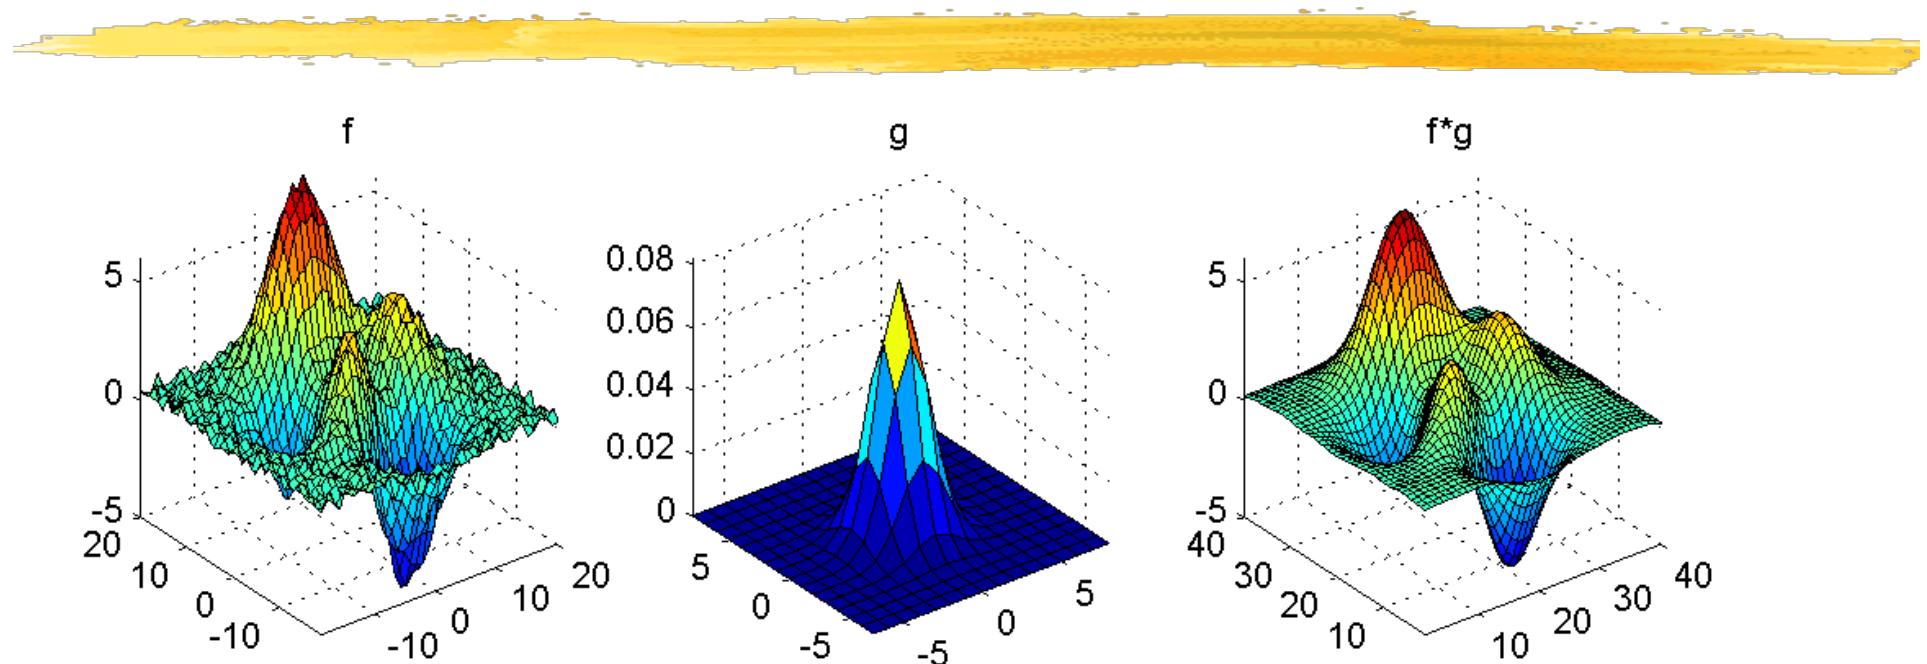
$$[-0.5, 0, 0.5] \rightarrow \frac{\partial f}{\partial x} \approx \frac{f(x+dx, y) - f(x-dx, y)}{2dx}$$

$$\begin{bmatrix} -1 \\ 1 \end{bmatrix} \rightarrow \frac{\partial f}{\partial y} \approx \frac{f(x, y+dy) - f(x, y)}{dy}$$

$$\begin{bmatrix} -0.5 \\ 0 \\ 0.5 \end{bmatrix} \rightarrow \frac{\partial f}{\partial y} \approx \frac{f(x, y+dy) - f(x, y-dy)}{2dy}$$

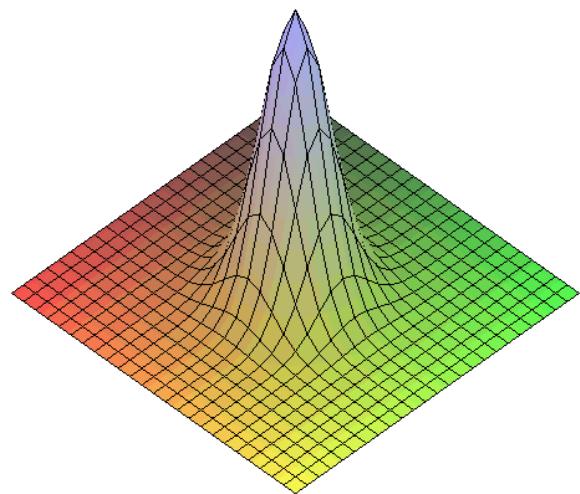
→ Use wider masks to add an element of smoothing

GAUSSIAN SMOOTHING

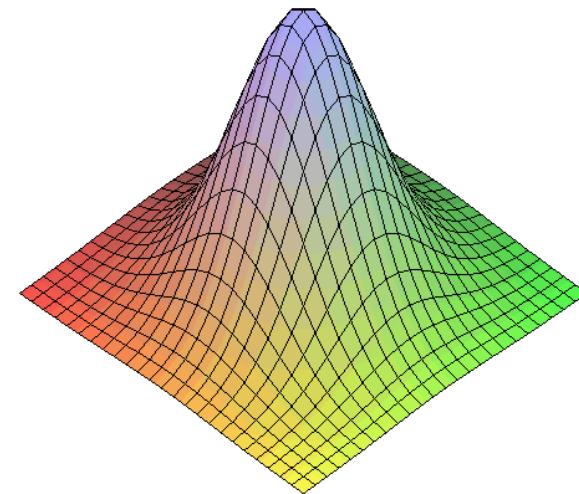


- Eliminates high frequency noise.
- Is fast because the kernel is
 - small,
 - separable.

GAUSSIAN MASKS



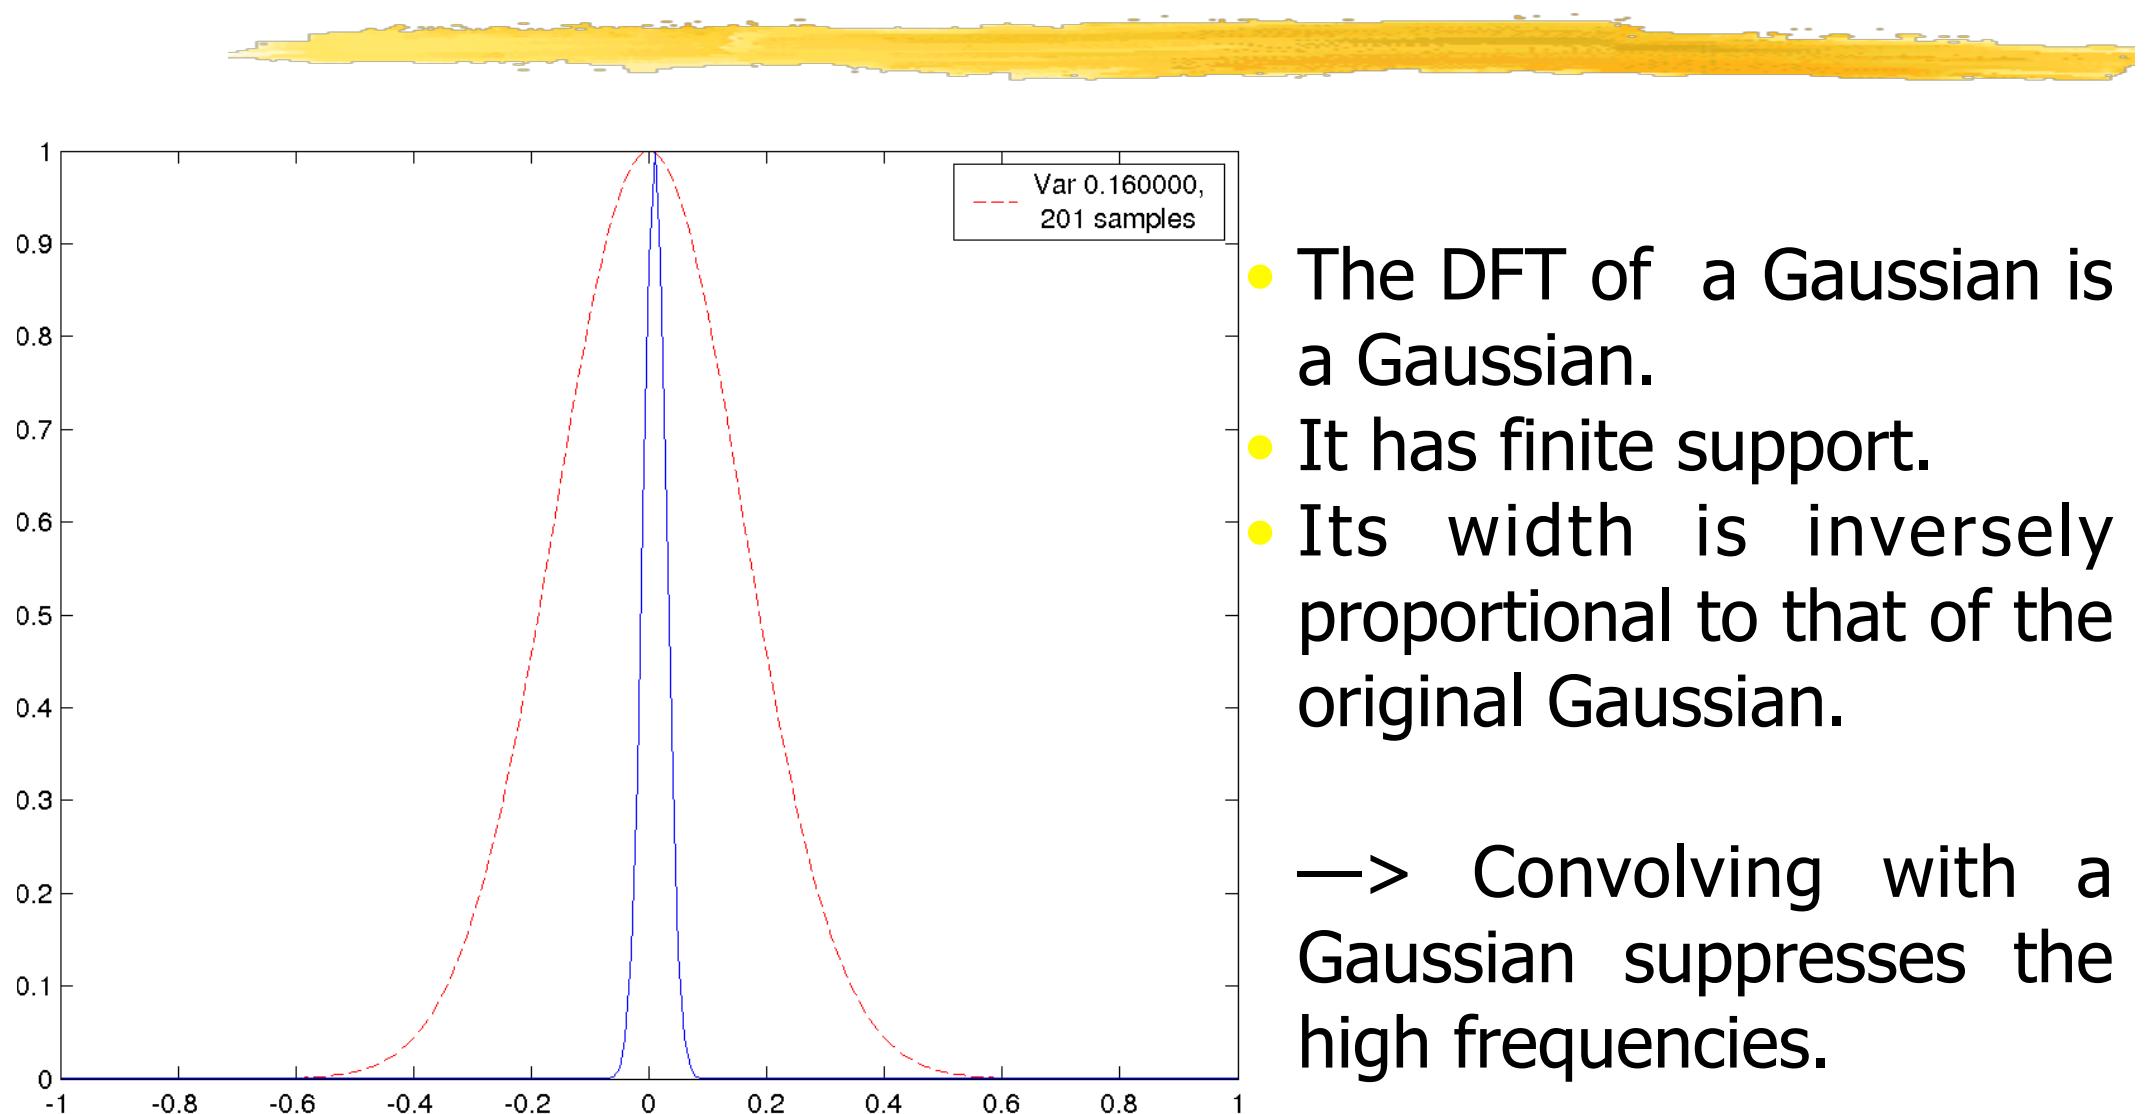
$$\sigma = 1$$



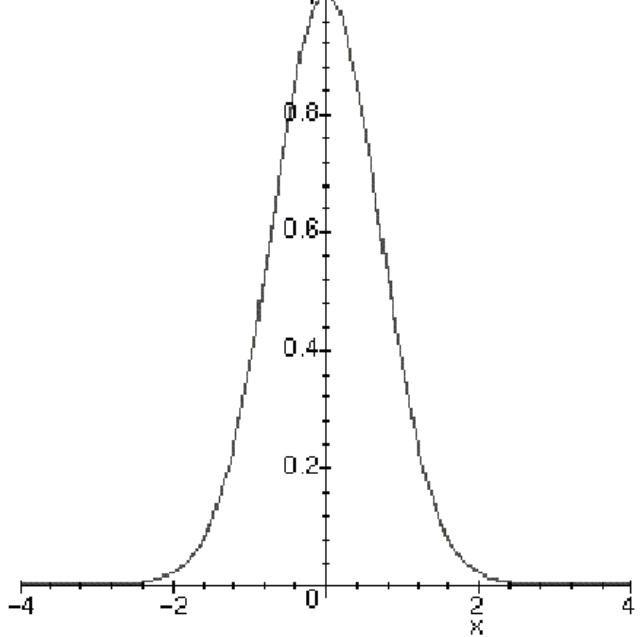
$$\sigma = 2$$

$$g_2(x, y) = \frac{1}{2\pi\sigma^2} \exp(-(x^2 + y^2)/2\sigma^2)$$

FOURIER TRANSFORM



SEPARABILITY

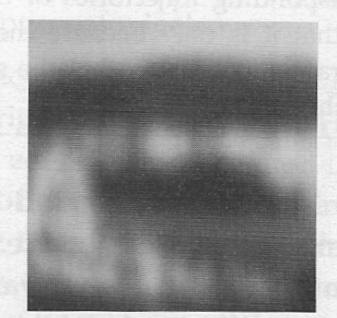
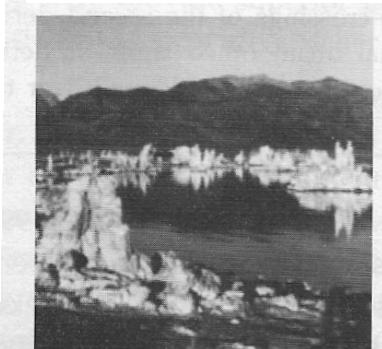
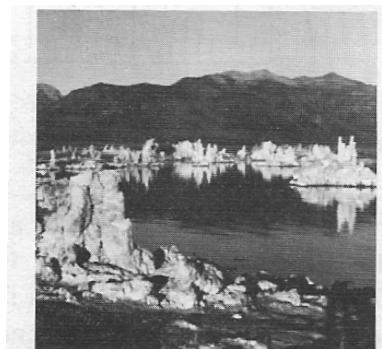
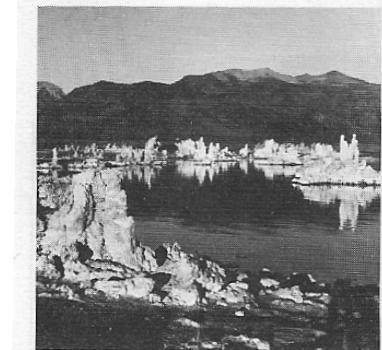


$$g_1(x) = \frac{1}{\sqrt{\pi}\sigma} \exp(-x^2/\sigma^2)$$

$$g_2(x, y) = g_1(x)g_1(y)$$

$$\begin{aligned} \iint_{uv} g_2(u, v) f(x-u, y-v) du dv &= \int_u g_1(u) \left(\int_v g_1(v) f(x-u, y-v) dv \right) du \\ &= \int_v g_1(v) \left(\int_u g_1(u) f(x-u, y-v) du \right) dv \end{aligned}$$

SMOOTHED IMAGES



GAUSSIAN DERIVATIVES

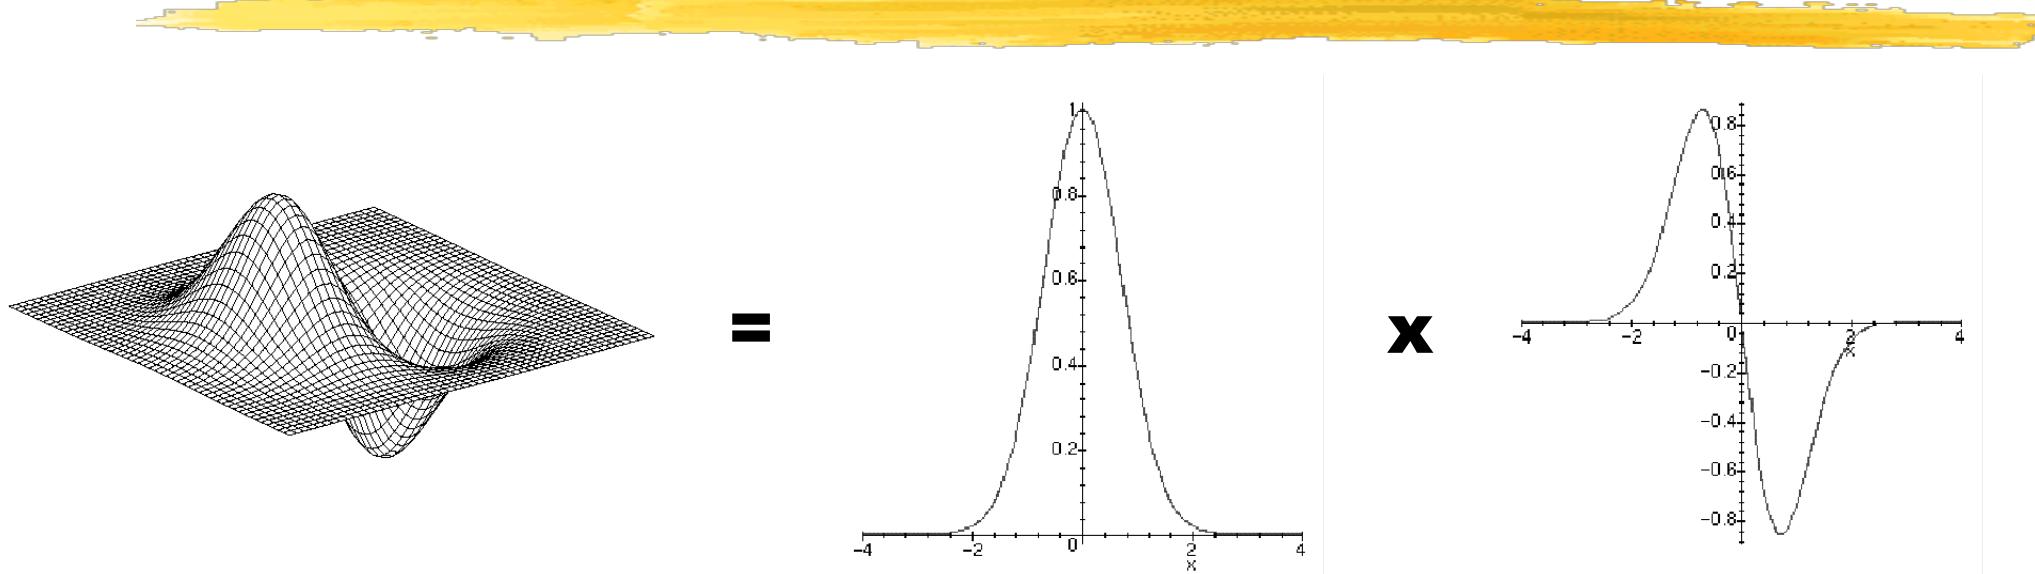


Image derivatives computed by convolving
with the derivative of a Gaussian:

$$\frac{\partial}{\partial x} \iint_{u,v} g_2(u,v) f(x-u, y-v) du dv = \int_u g_1'(u) \left(\int_v g_1(v) f(x-u, y-v) dv \right) du$$

$$\frac{\partial}{\partial y} \iint_{u,v} g_2(u,v) f(x-u, y-v) du dv = \int_v g_1'(v) \left(\int_u g_1(u) f(x-u, y-v) du \right) dv$$

GAUSSIAN MASKS



Sigma=1:

g : 0.000070 0.010332 0.207532 0.564131 0.207532 0.010332 0.000070

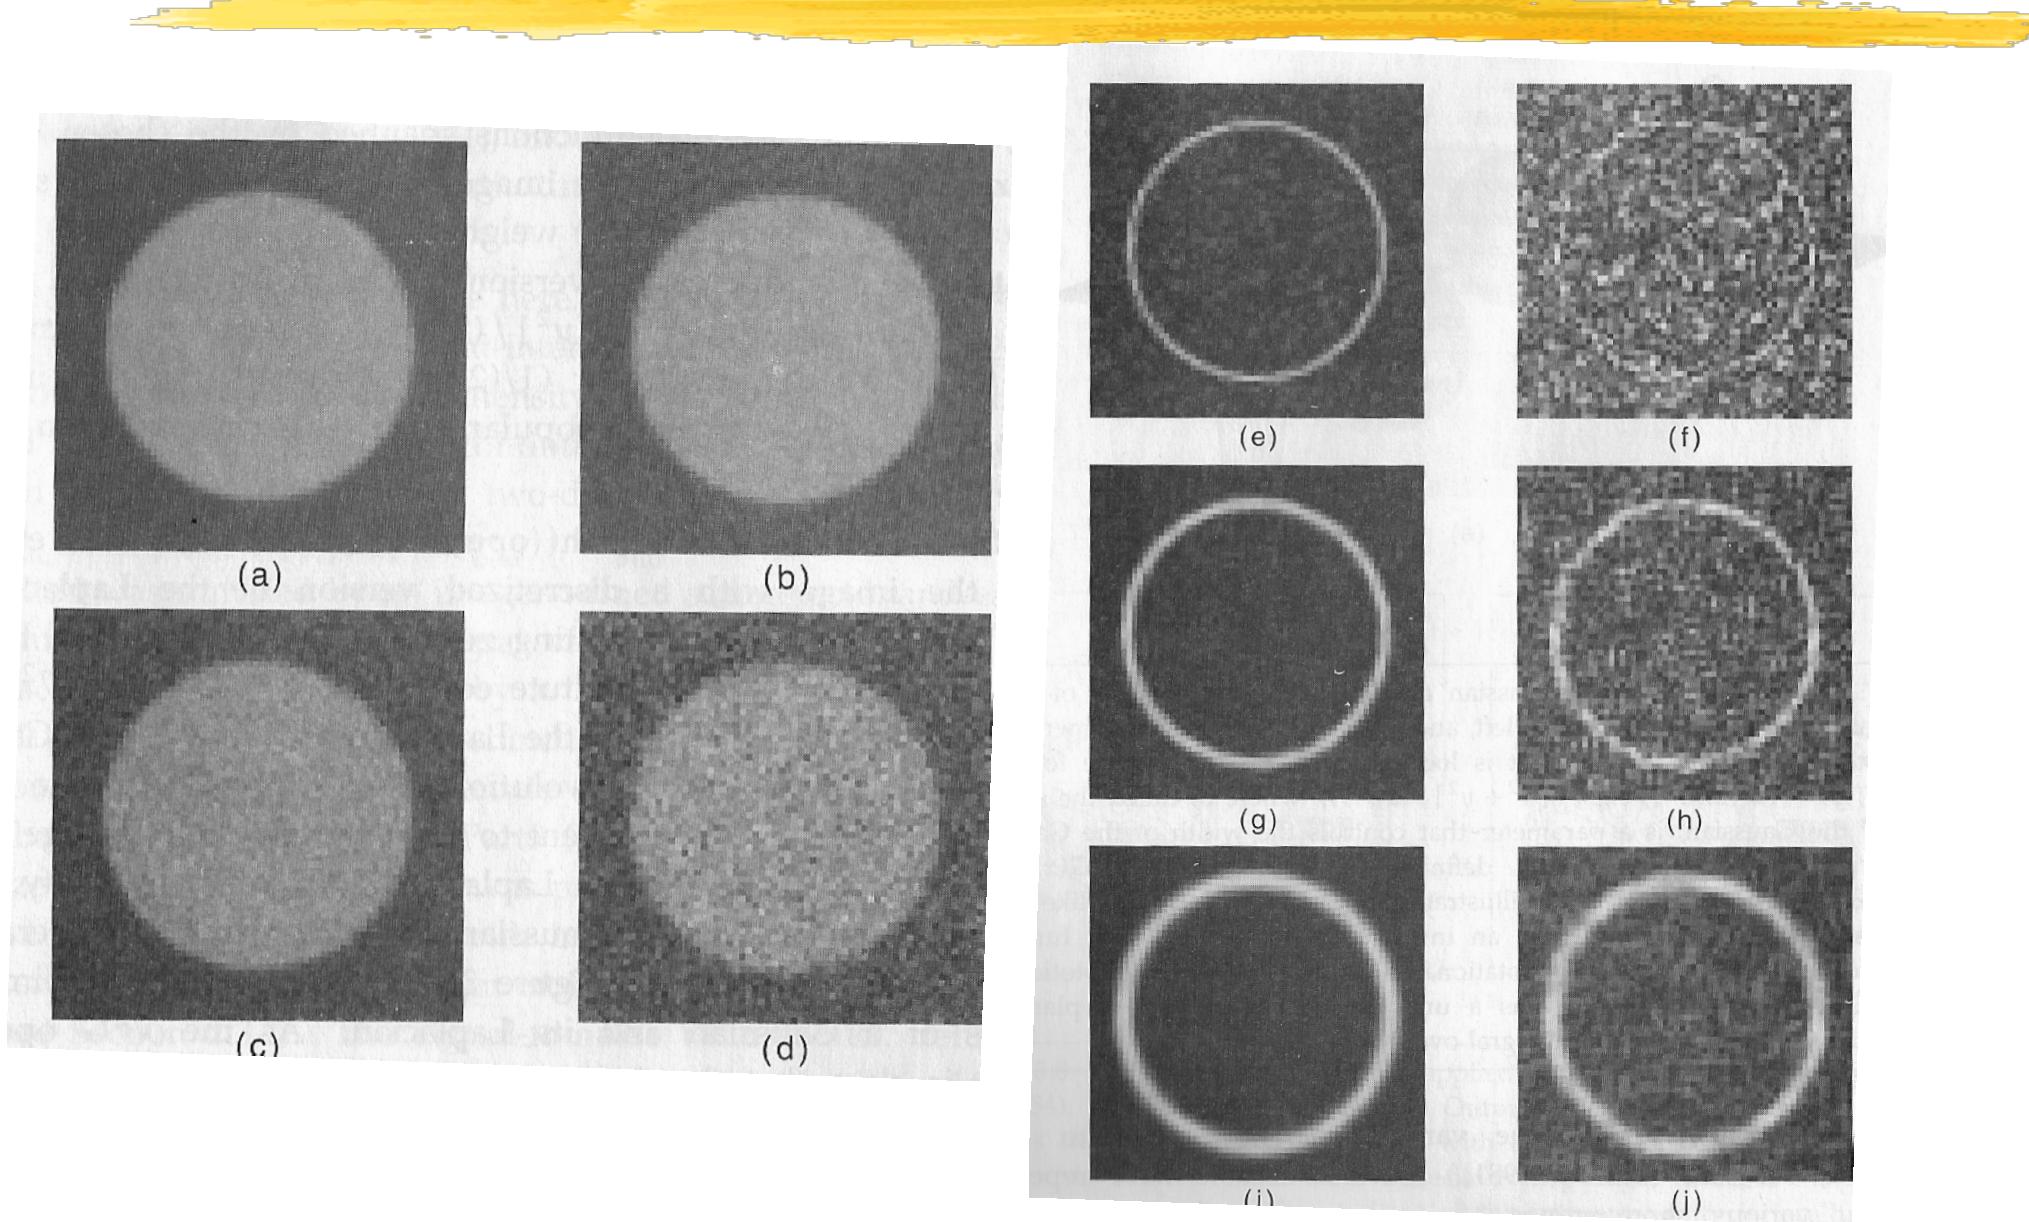
g' : 0.000418 0.041330 0.415065 0.000000 -0.415065 -0.041330 -0.000418

Sigma=2:

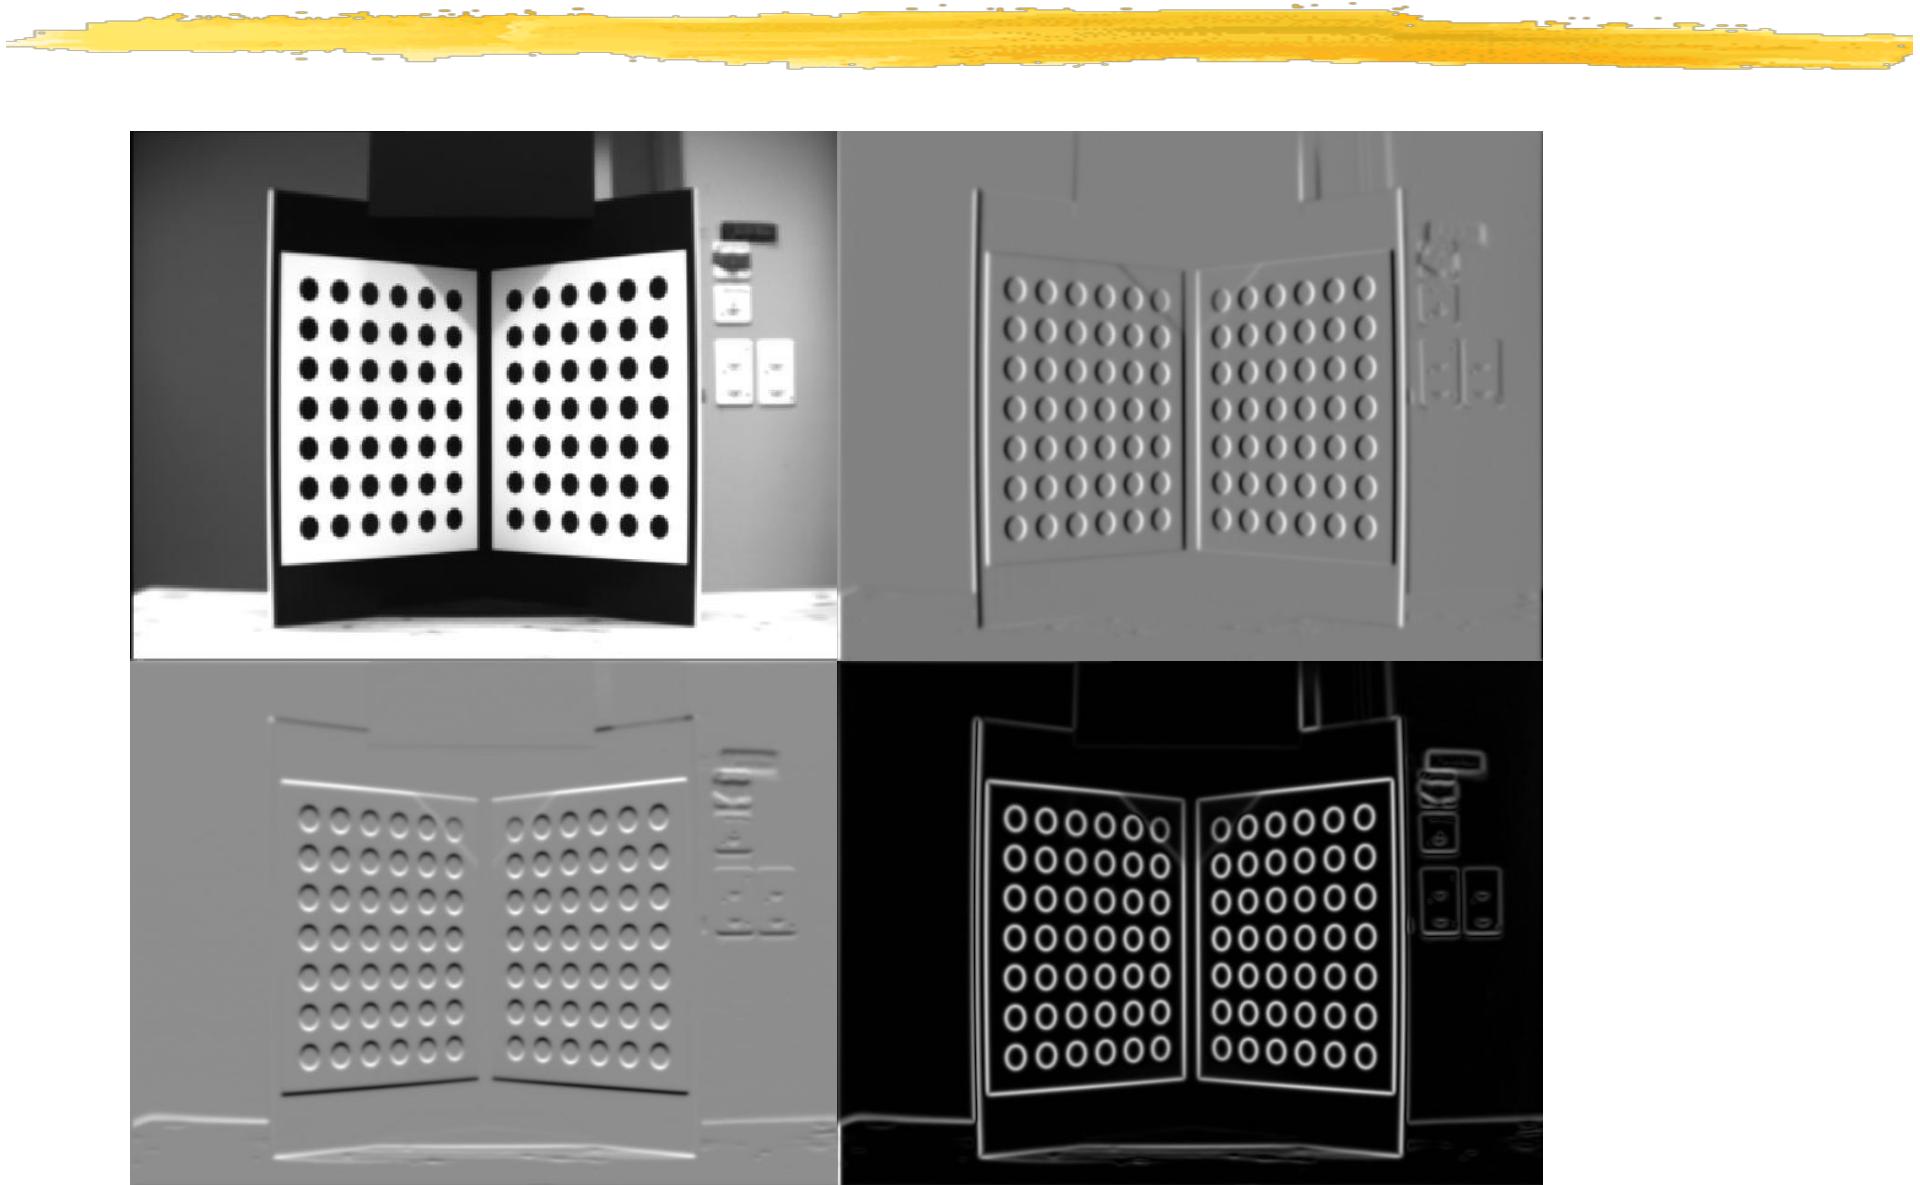
g : 0.005167 0.029735 0.103784 0.219712 0.282115 0.219712 0.103784 0.029735 0.005167

g' : 0.010334 0.044602 0.103784 0.109856 0.000000 -0.109856 -0.103784 -0.044602 -0.010334

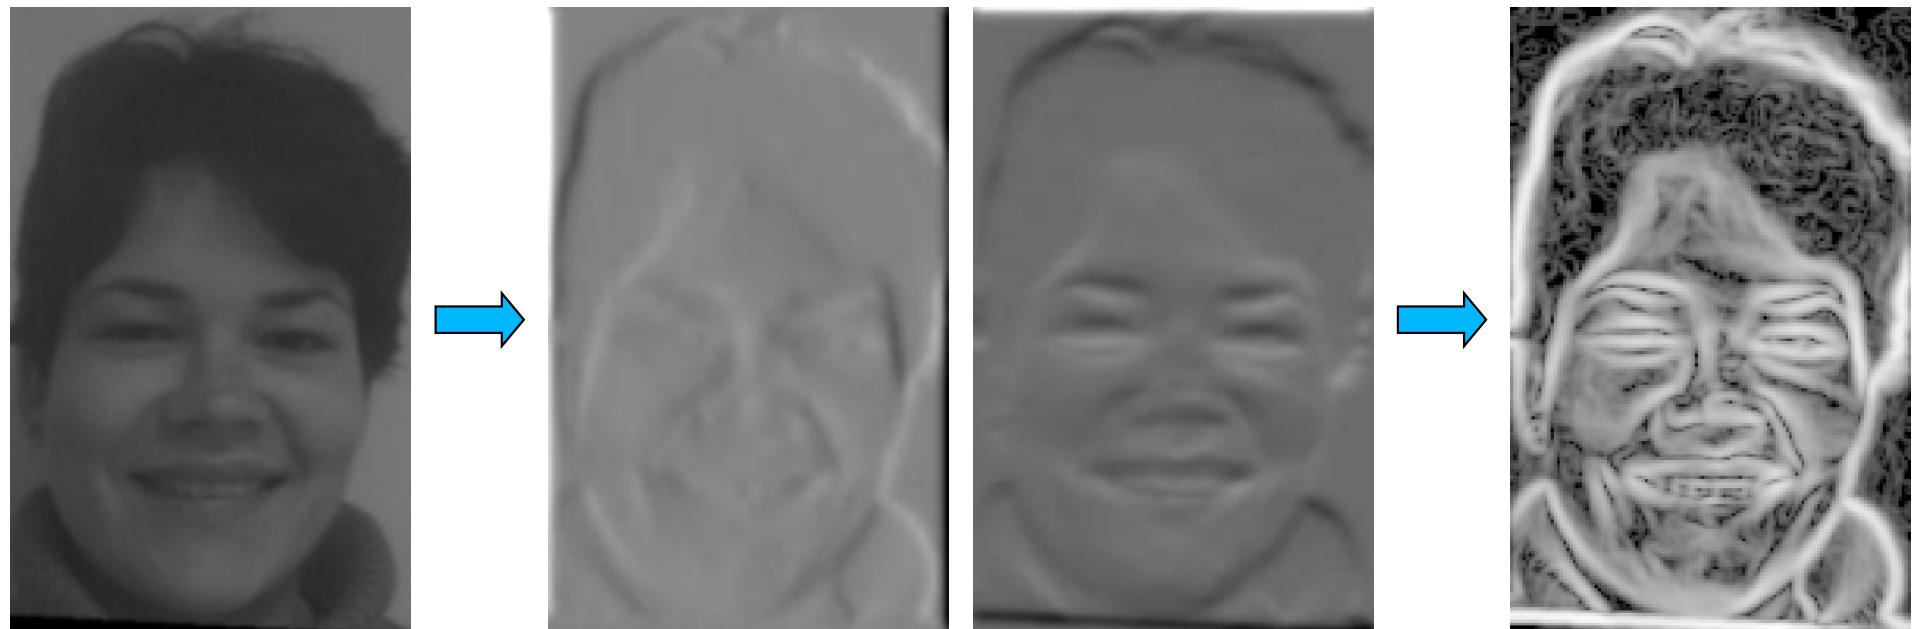
MASK SIZE



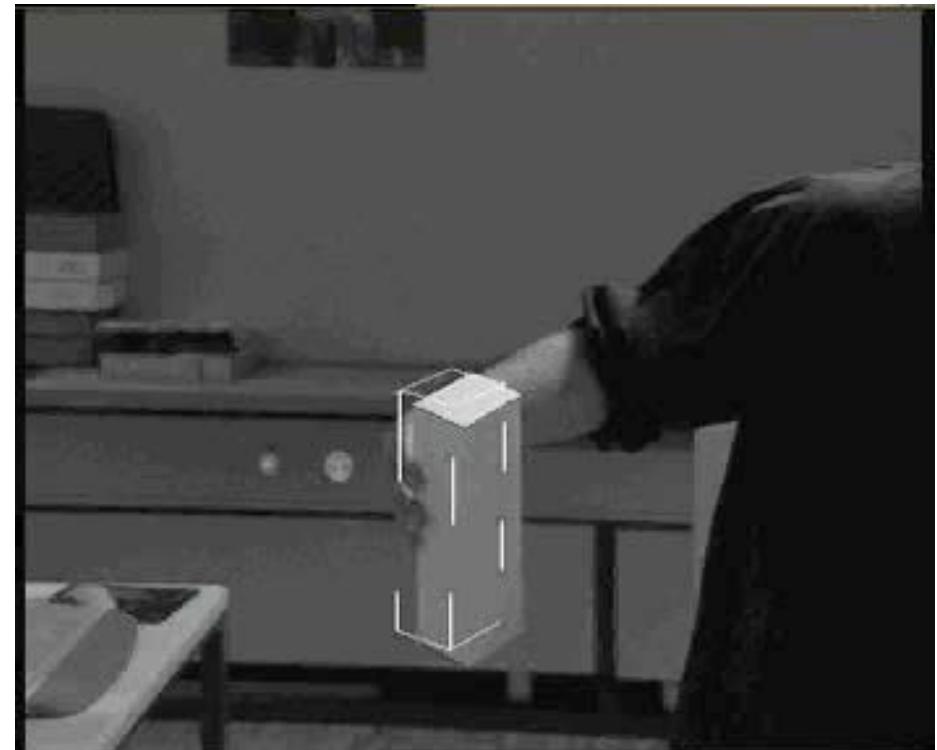
DERIVATIVE IMAGES



DERIVATIVE IMAGES

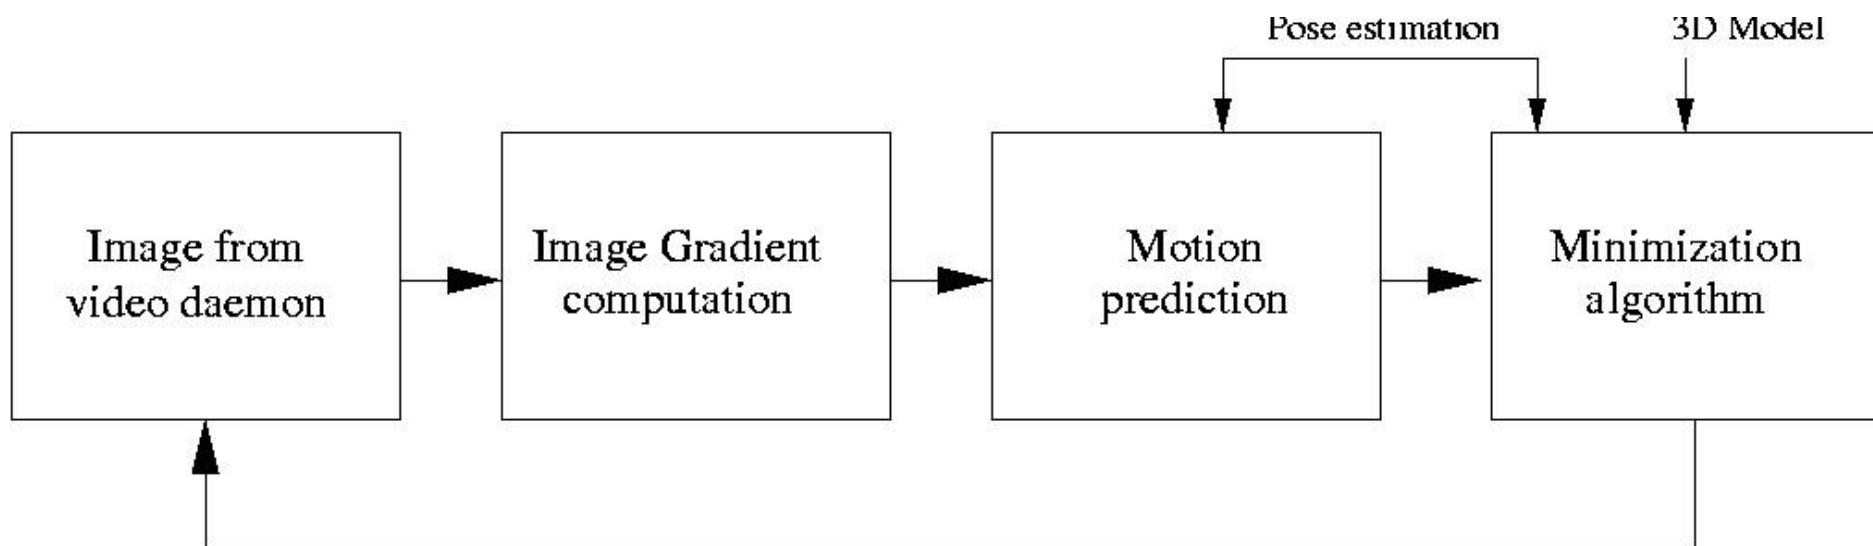
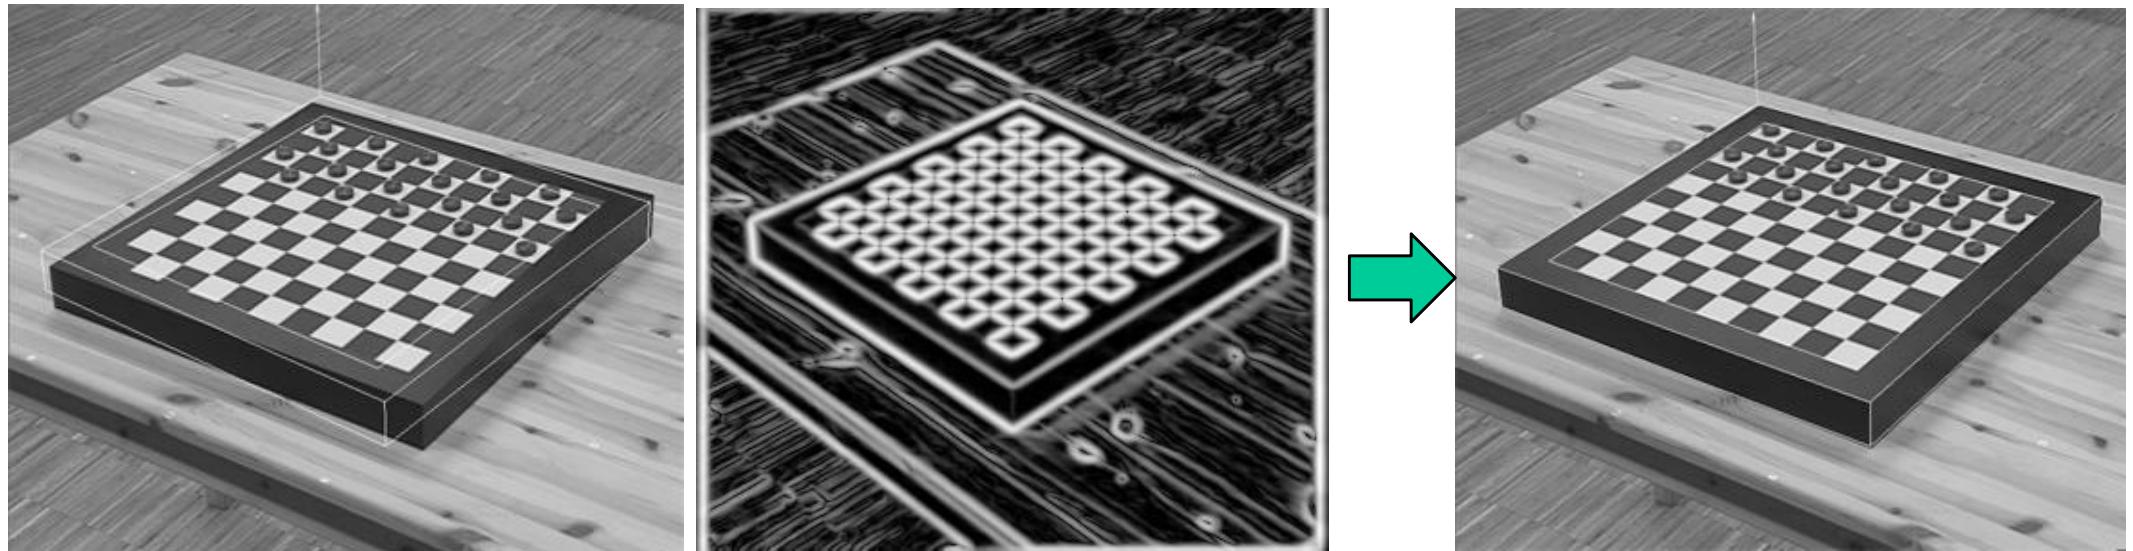


GRADIENT-BASED TRACKING



Maximize edge-strength along projection of the 3—D wireframe.

GRADIENT MAXIMIZATION

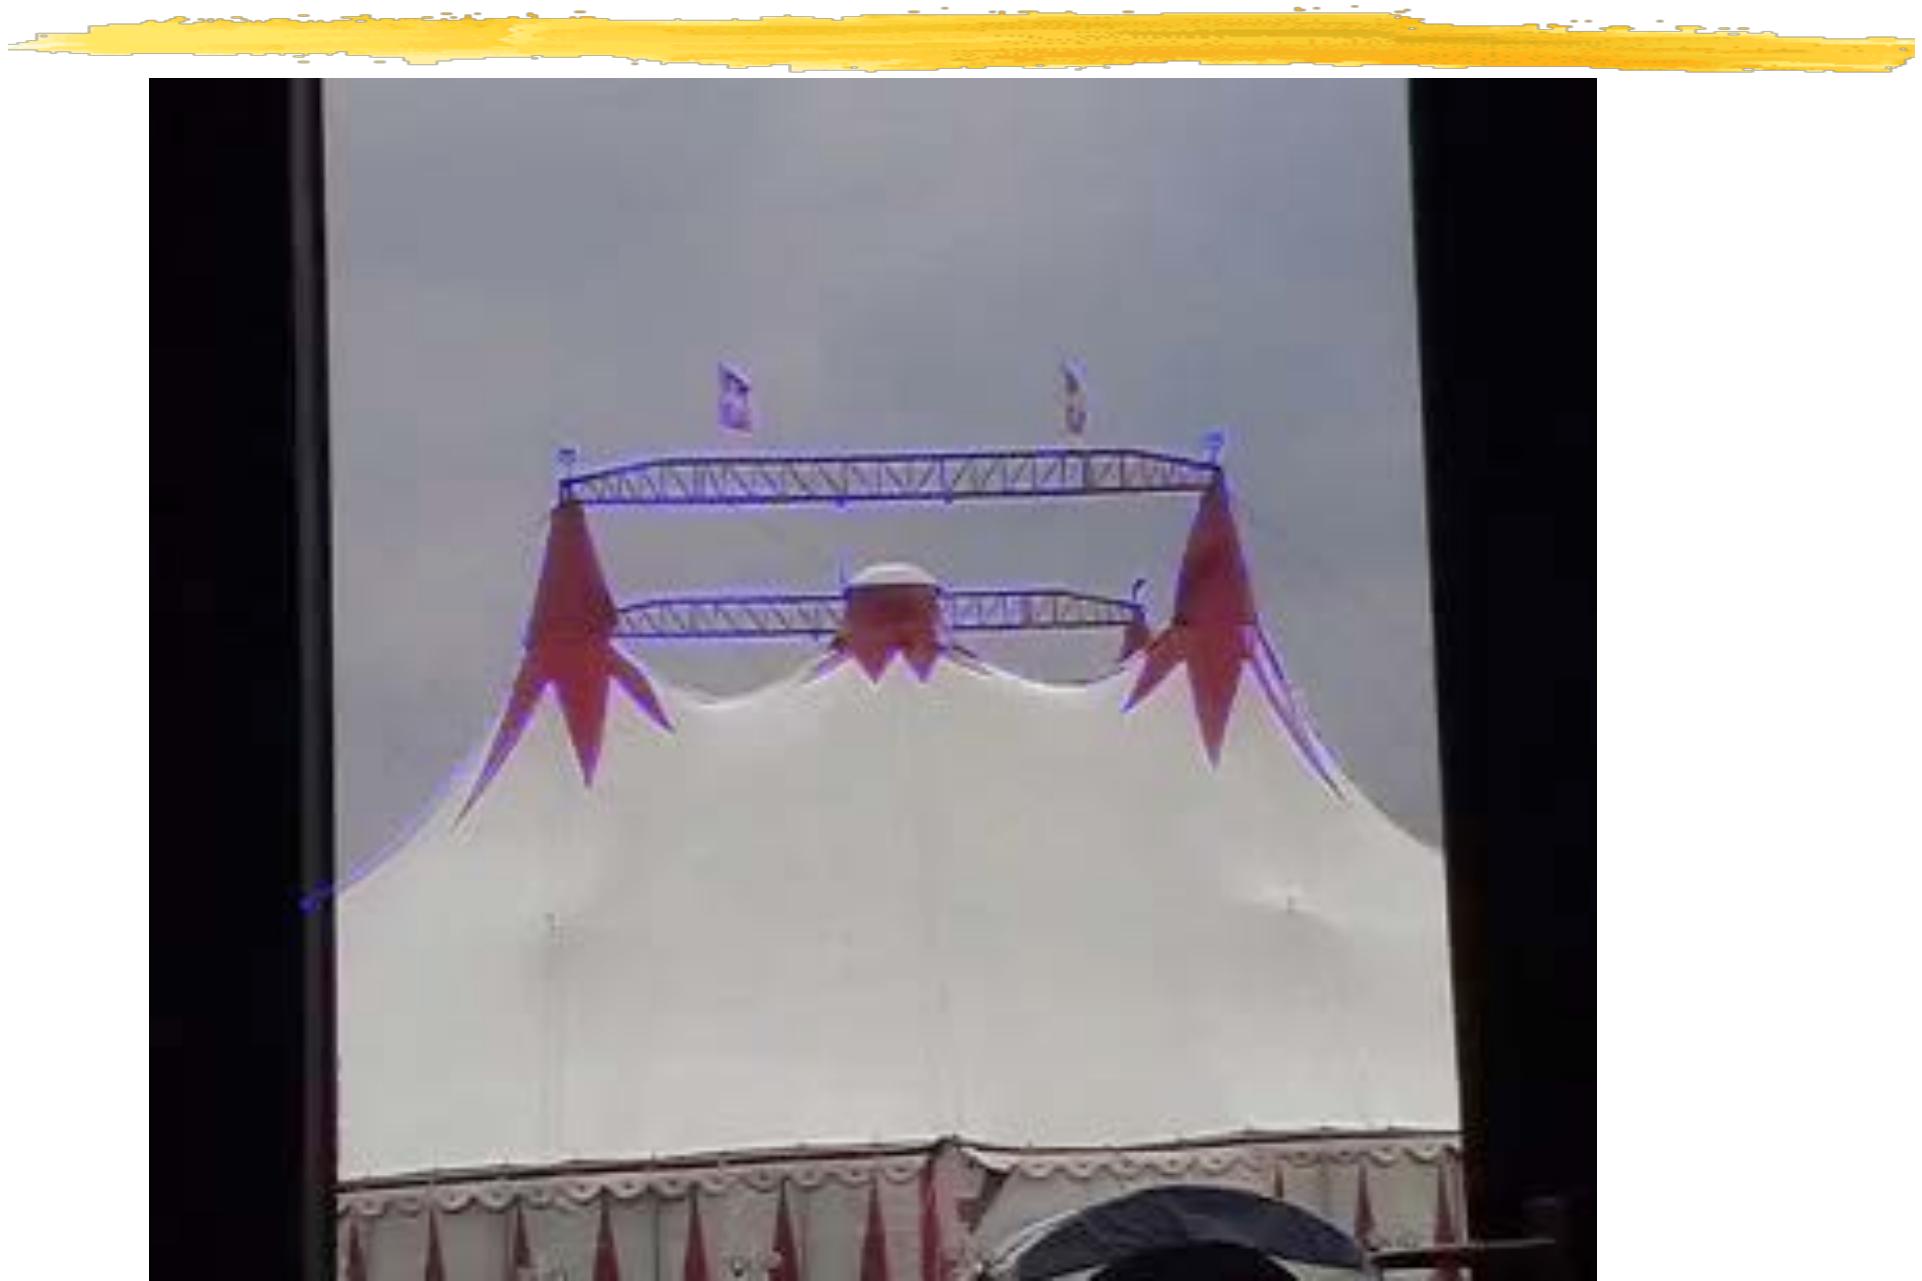


WING DEFORMATION

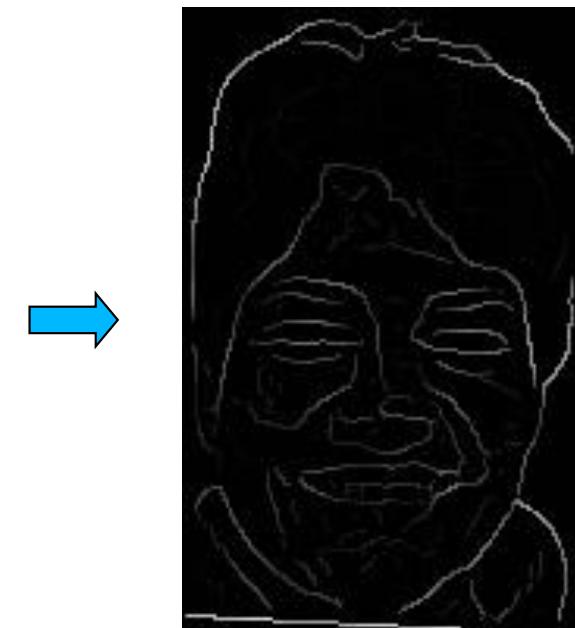
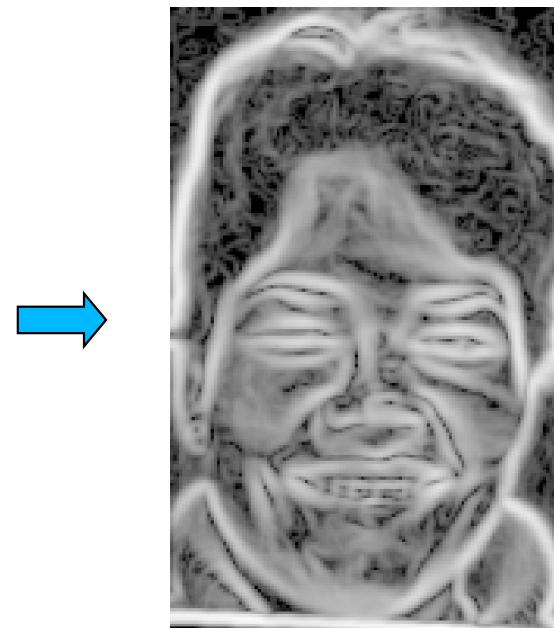
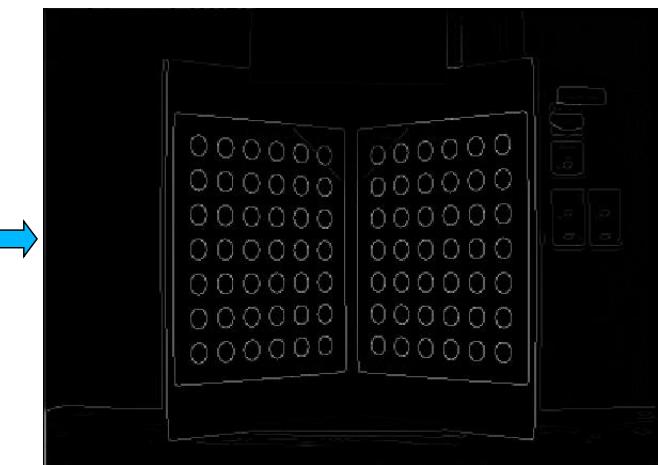
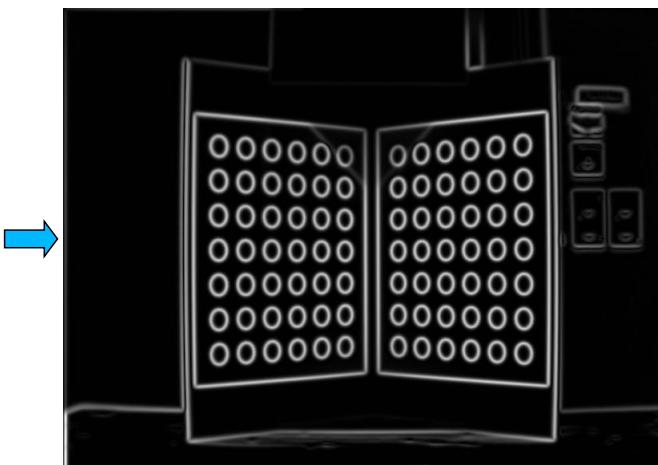
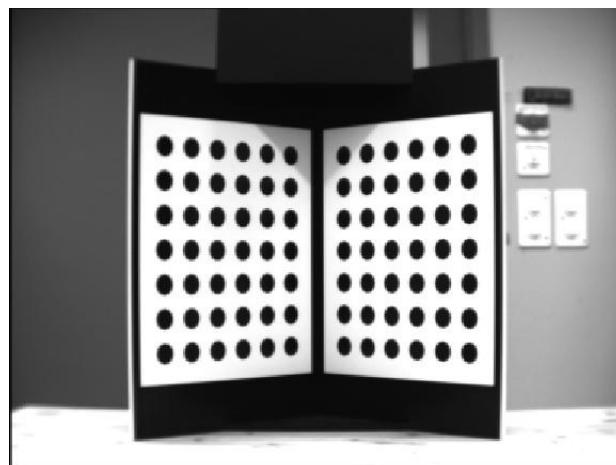


- Measure true behavior in flight.
- Validate computational models.

REAL-TIME TRACKING



CANNY EDGE DETECTOR



CANNY EDGE DETECTOR



Convolution

- Gradient strength
- Gradient direction

Thresholding

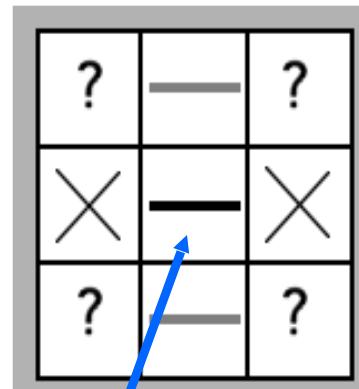
Non Maxima Suppression

Hysteresis Thresholding

NON-MAXIMA SUPPRESSION

In parallel, at each pixel in edge image, use window to select as a function of edge orientation:

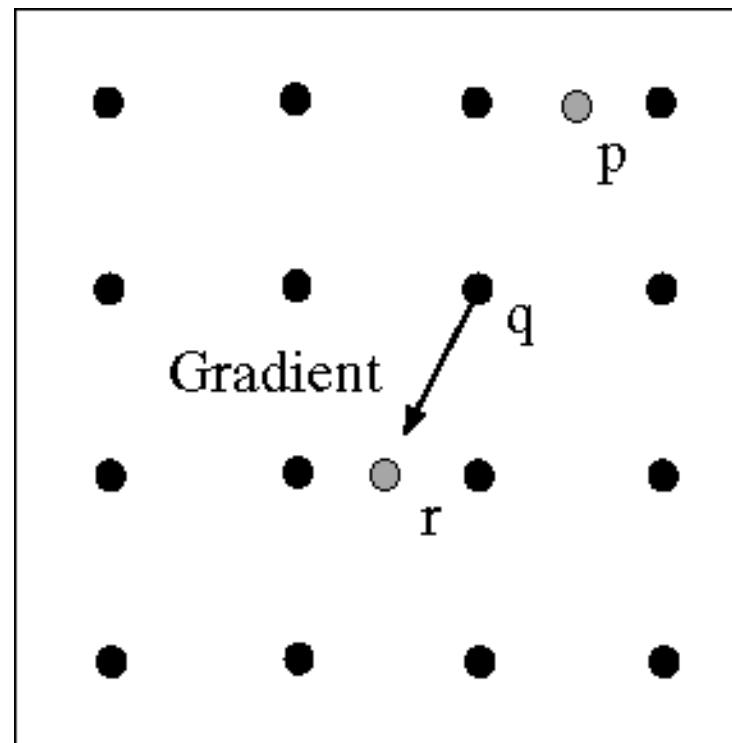
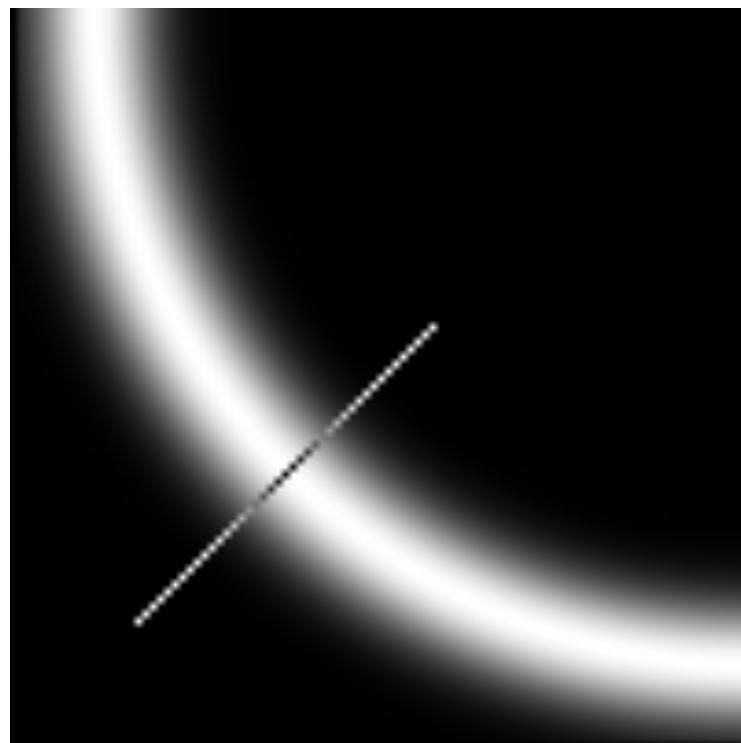
Window W



Central Edge

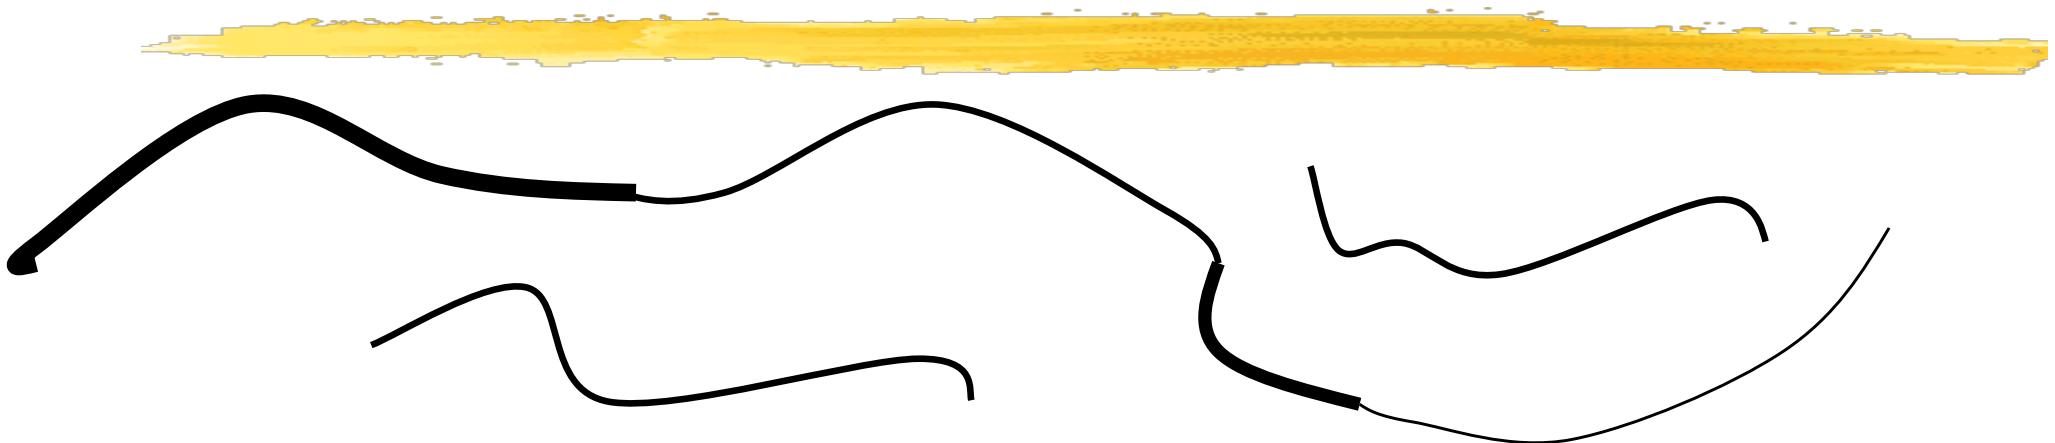
- Definitely consider these edges
- X Do not consider these edges
- ? Maybe consider them, depending on algorithm

NON-MAXIMA SUPPRESSION



Check if pixel is local maximum along gradient direction,
which requires checking interpolated pixels p and r.

HYSTERESIS THRESHOLDING



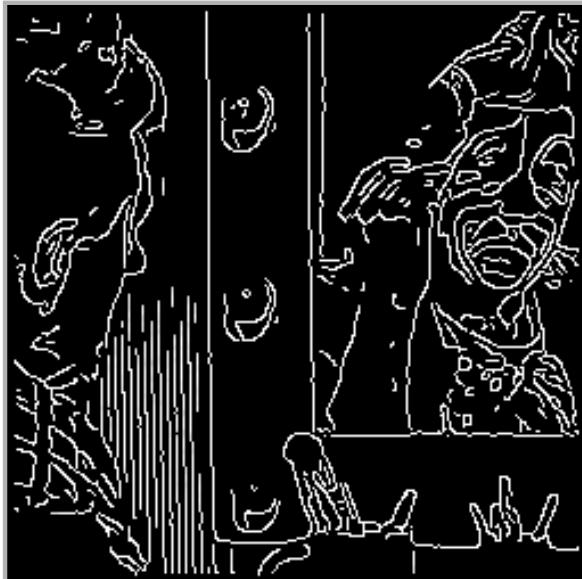
Algorithm takes two thresholds: high & low

- A pixel with edge strength above high threshold is an edge.
- Any pixel with edge strength below low threshold is not.
- Any pixel above the low threshold and next to an edge is an edge.

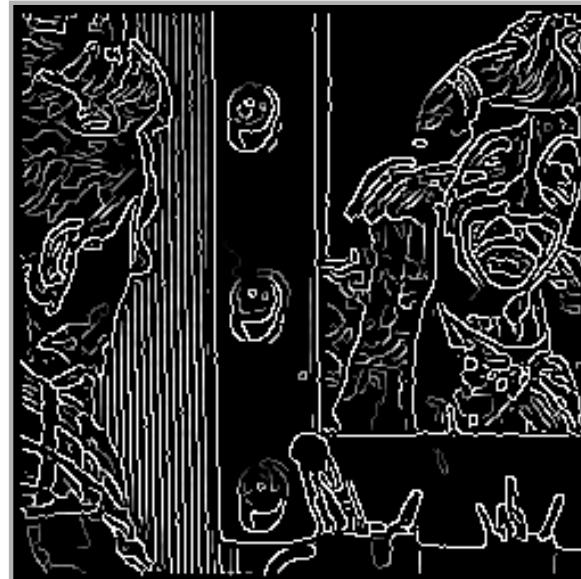
Iteratively label edges

- Edges grow out from 'strong edges'
- Iterate until no change in image.

CANNY RESULTS



$\sigma=1$, $T_2=255$, $T_1=220$



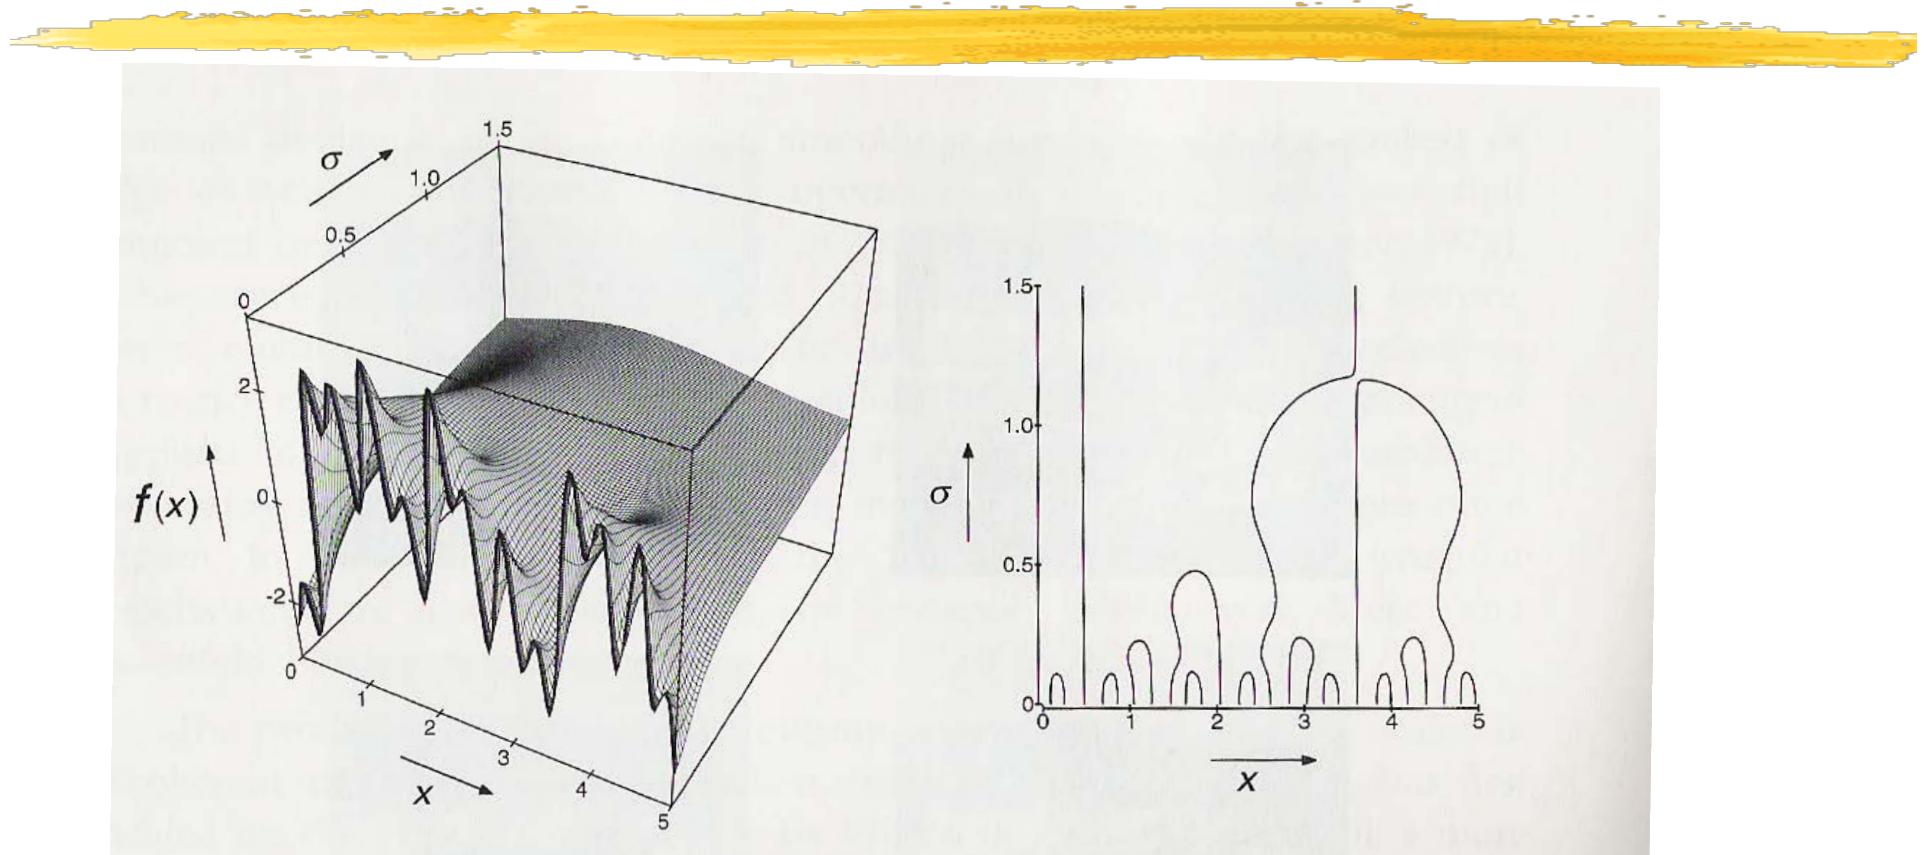
$\sigma=1$, $T_2=128$, $T_1=1$



$\sigma=2$, $T_2=128$, $T_1=1$

M. Heath, S. Sarkar, T. Sanocki, and K.W. Bowyer, "A Robust Visual Method for Assessing the Relative Performance of Edge-Detection Algorithms" IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. 19, No. 12, December 1997, pp. 1338-1359.

SCALE SPACE



Increasing scale (σ) removes details but never adds new ones:

- Edge position may shift.
- Two edges may merge.
- An edge may **not** split into two.

MULTIPLE SCALES



$\sigma = 1$



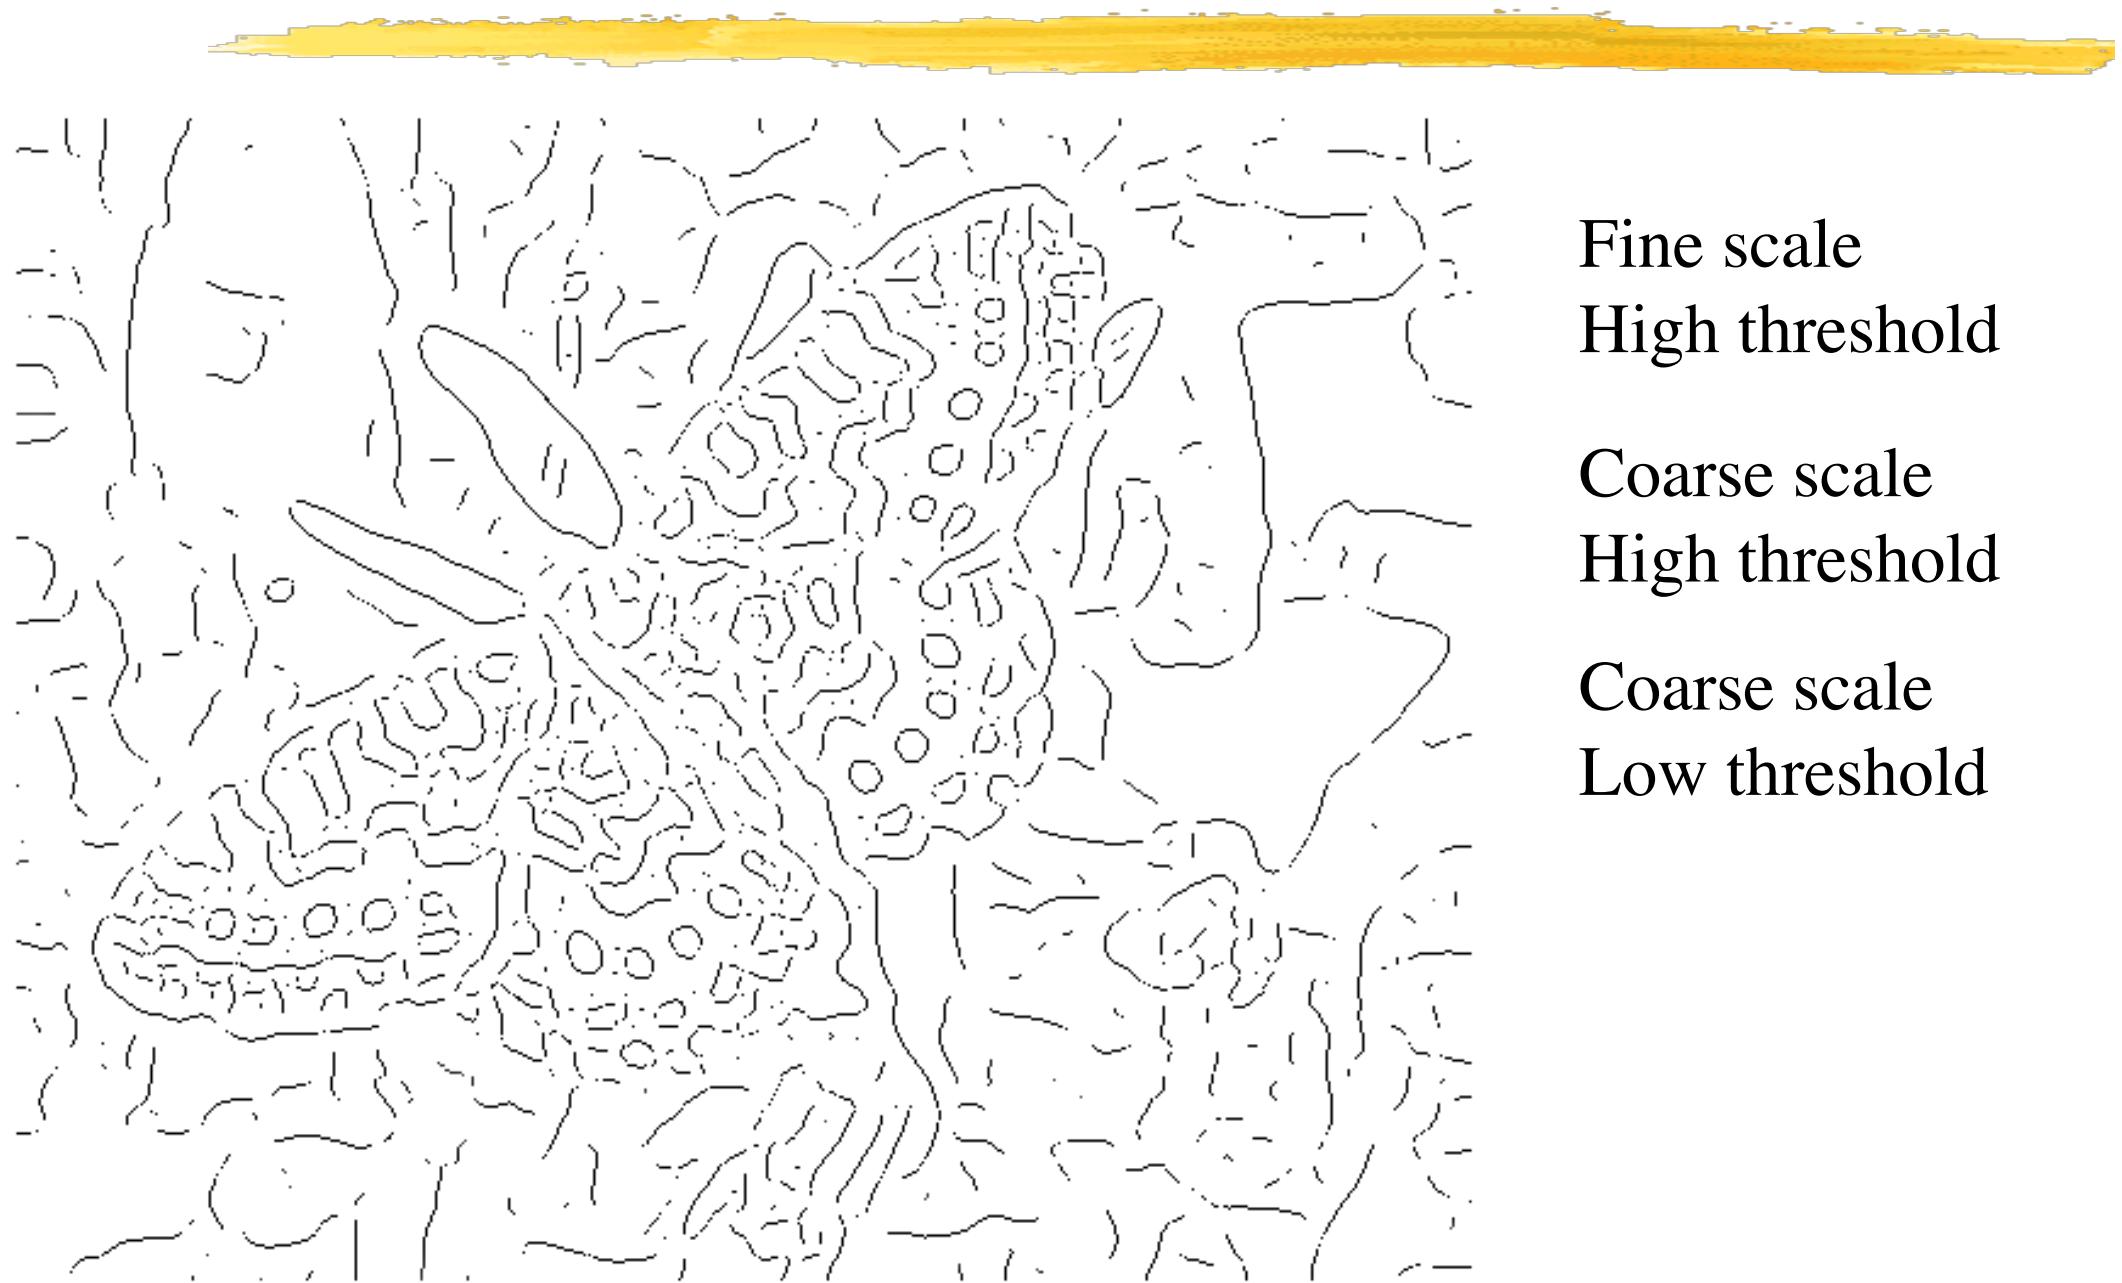
$\sigma = 2$



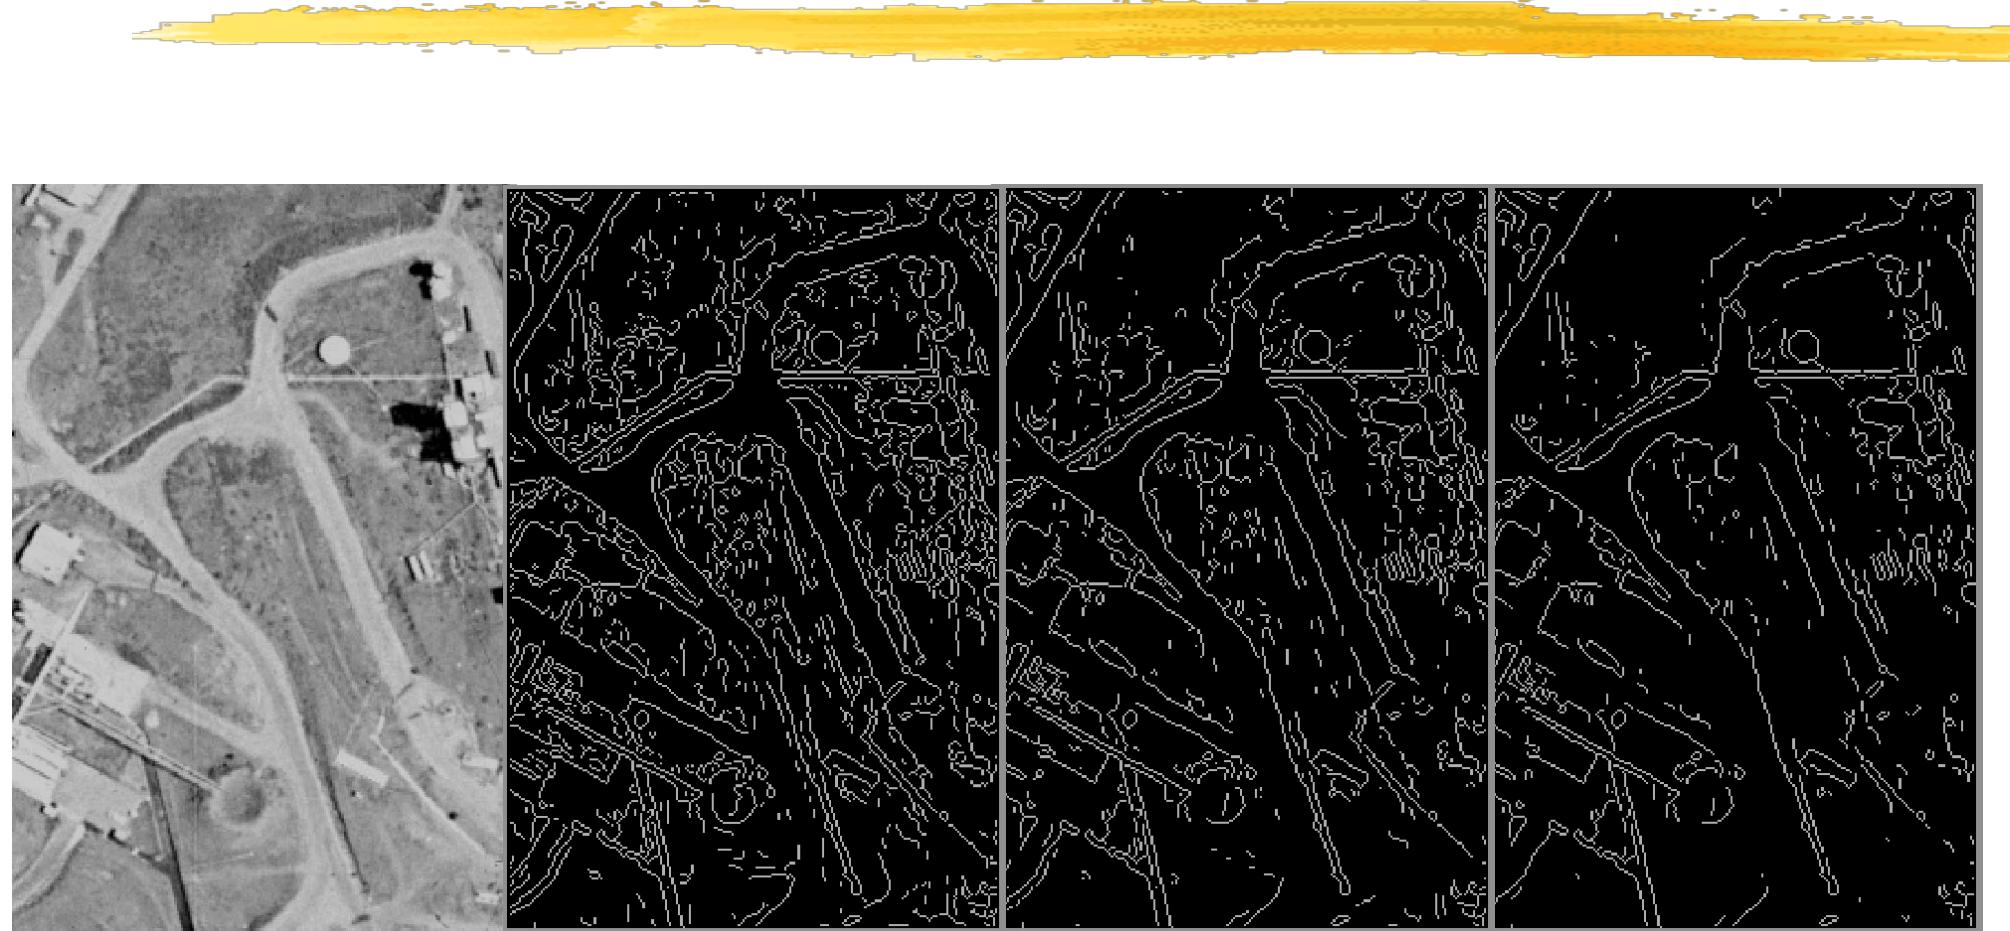
$\sigma = 4$

→ Choosing the right scale is a difficult semantic problem.

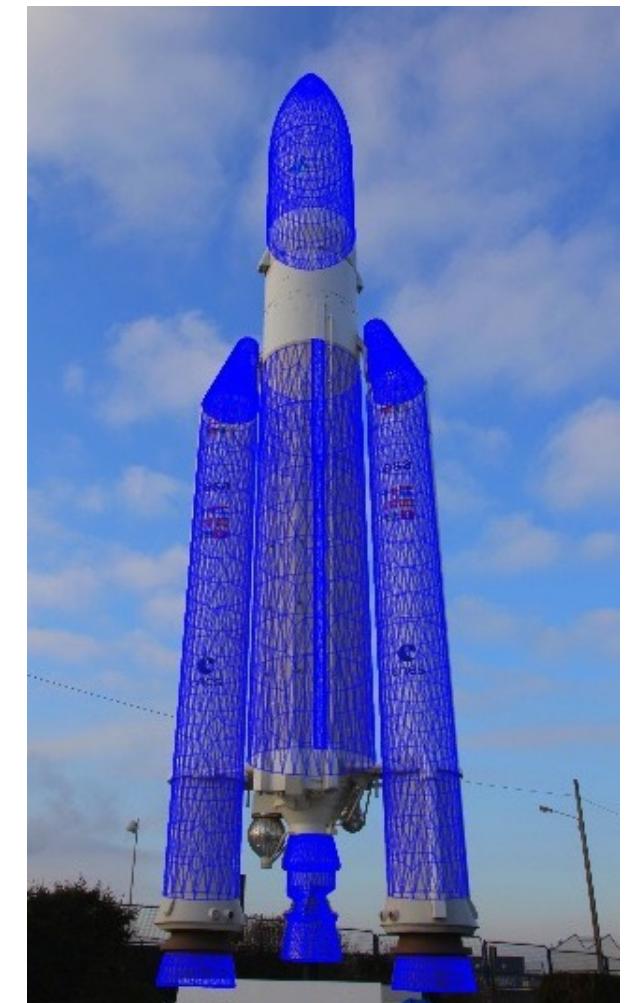
SCALE vs THRESHOLD



ROAD IMAGE



TRACKING ARIANE



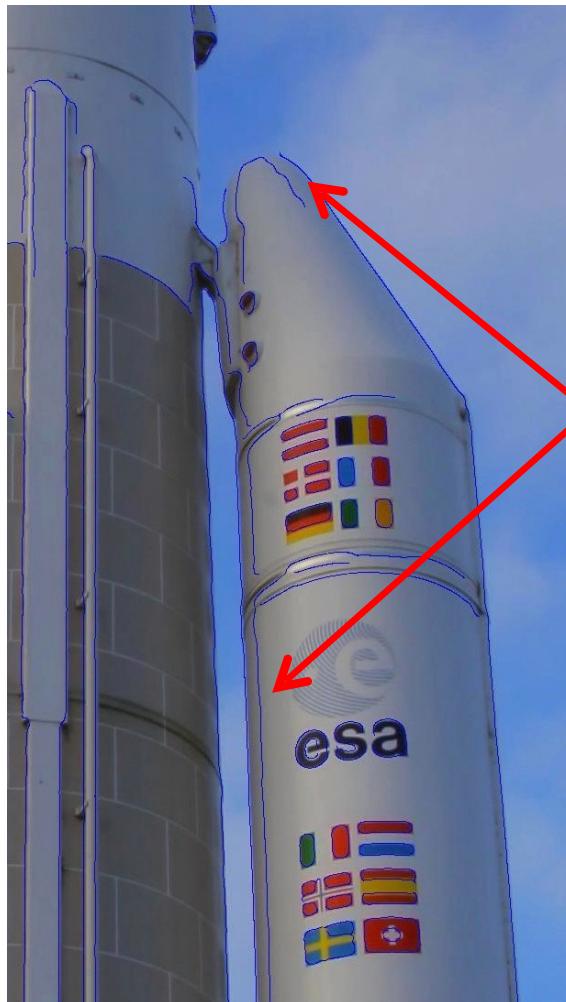
RAPID TRACKER



Given an initial pose estimate:

- Estimate the occluding contours.
- Find closest edge points in the normal direction.
- Re-estimate pose to minimize distances in the least squares sense.
- Iterate until convergence.

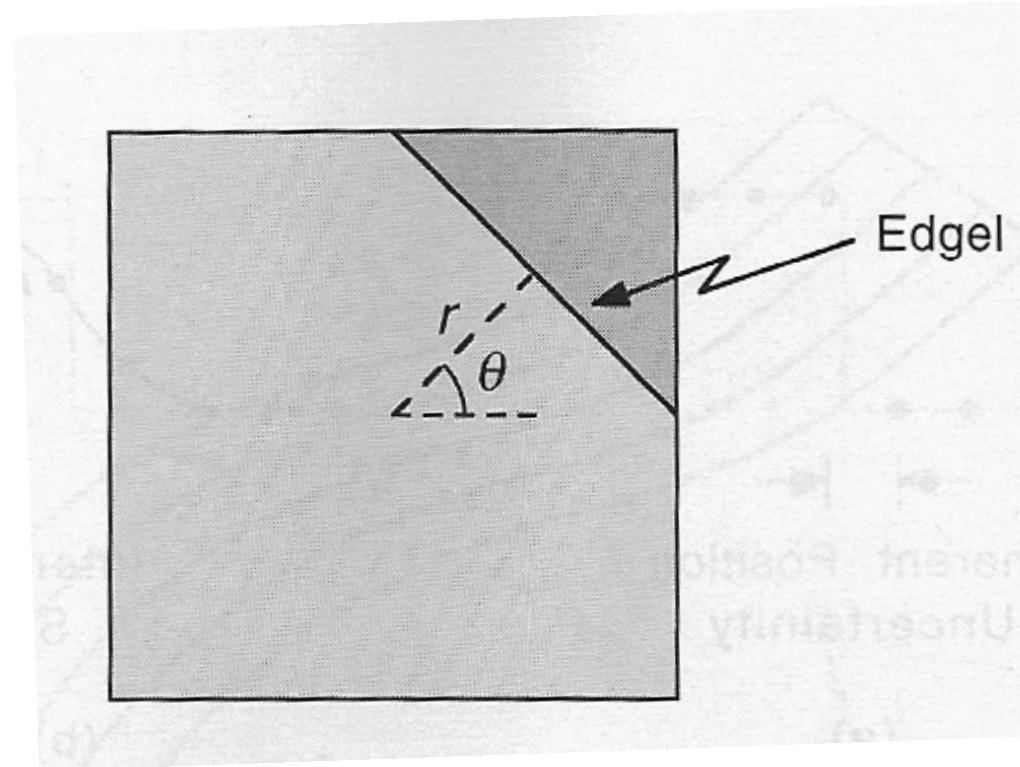
STEEP SMOOTH SHADING



- Rapidly varying gray levels.
- Large gradients.

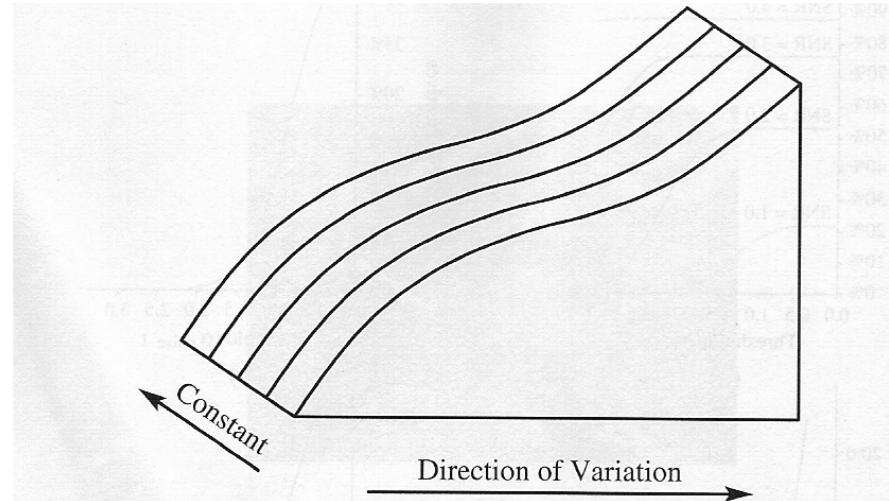
→ Shading can produce spurious edges.

PARAMETRIC MODEL MATCHERS



→ 4 parameters model to be fit in the least squares sense.

SURFACE FITTING



1. Estimate the edgel direction by fitting a cubic surface.
2. Fit a 1-D surface in the direction of the edgel
 - Step shaped surface,
 - Quadratic polynomial.
3. Declare an edge if step shape better than quadratic.

TEXTURE BOUNDARIES

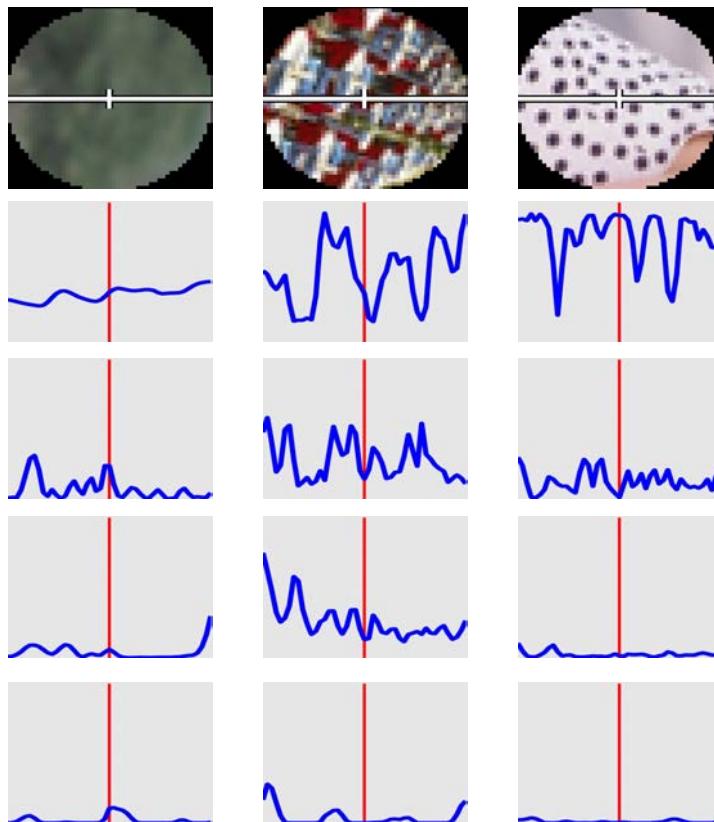


- Not all image contours are characterized by strong contrast.
- Sometimes, textural changes are just as significant.

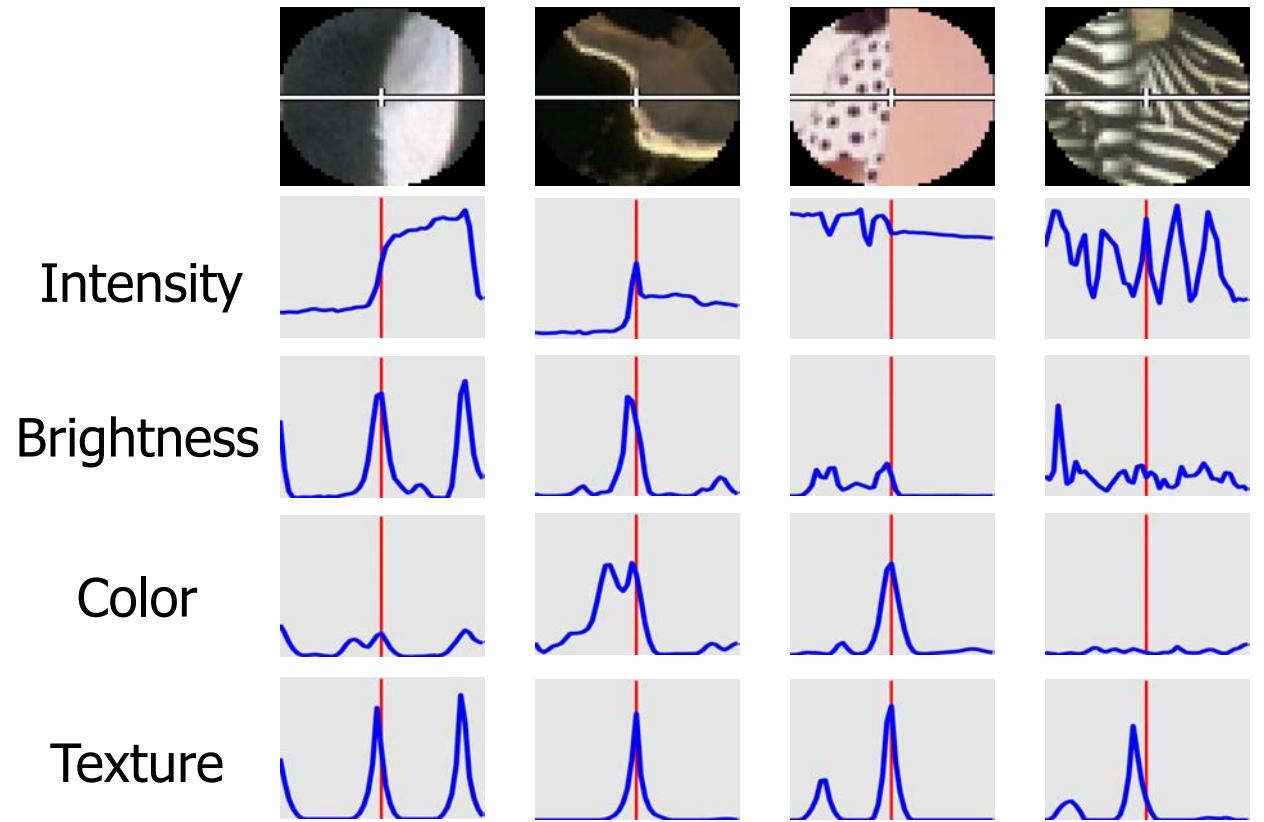
DIFFERENT BOUNDARY TYPES



Non-boundaries



Boundaries

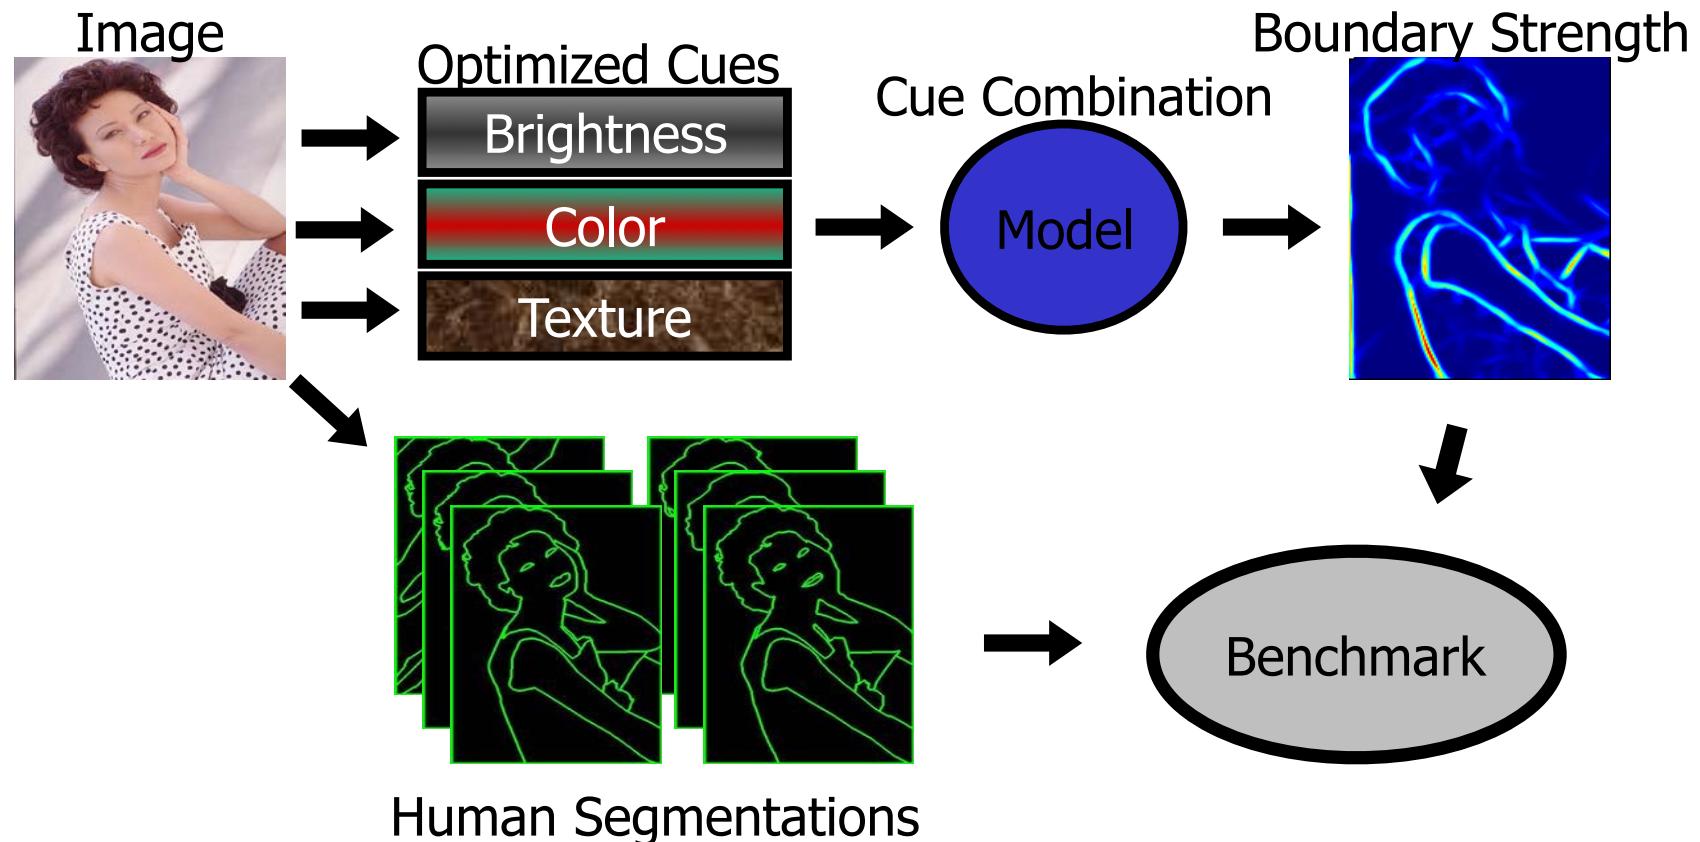


TRAINING DATABASE



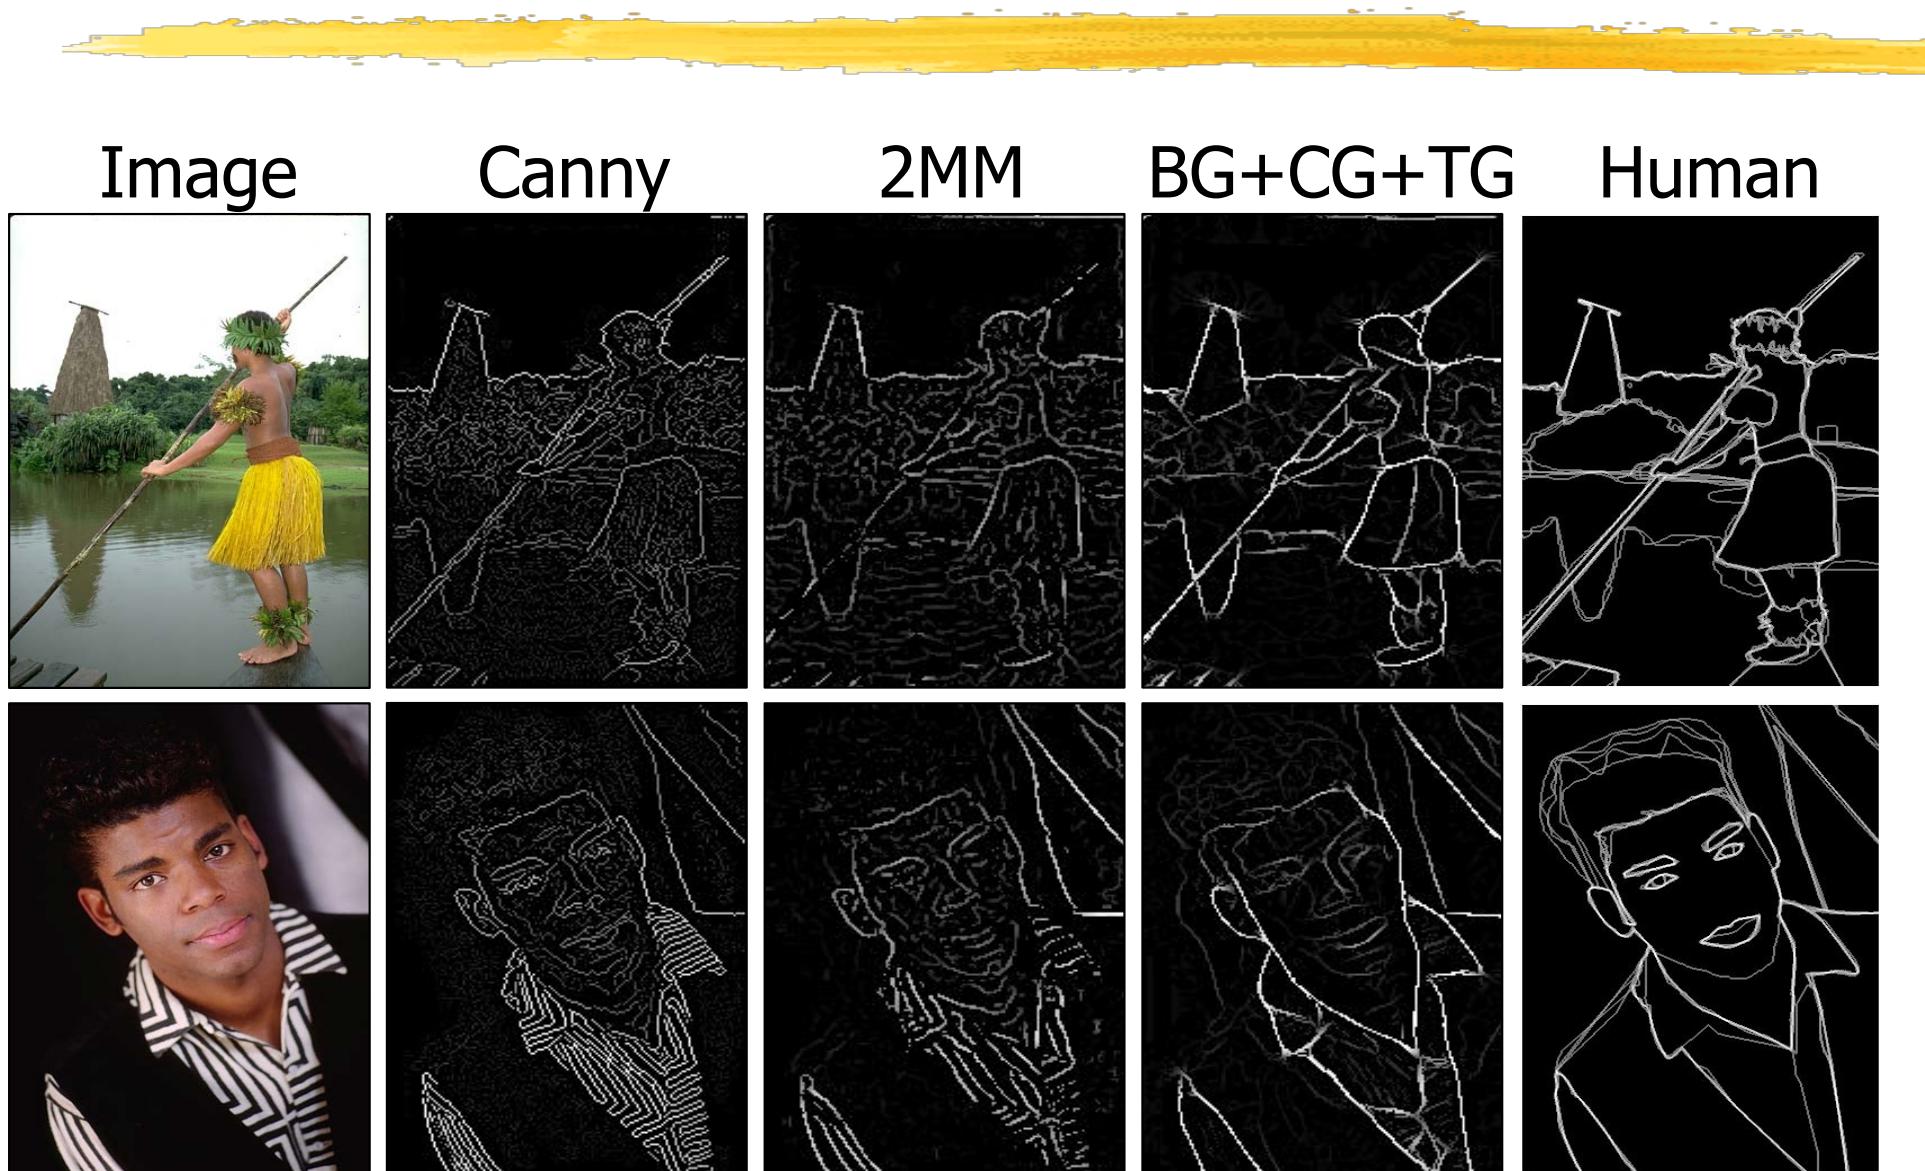
1000 images with 5 to 10 segmentations each.

MACHINE LEARNING

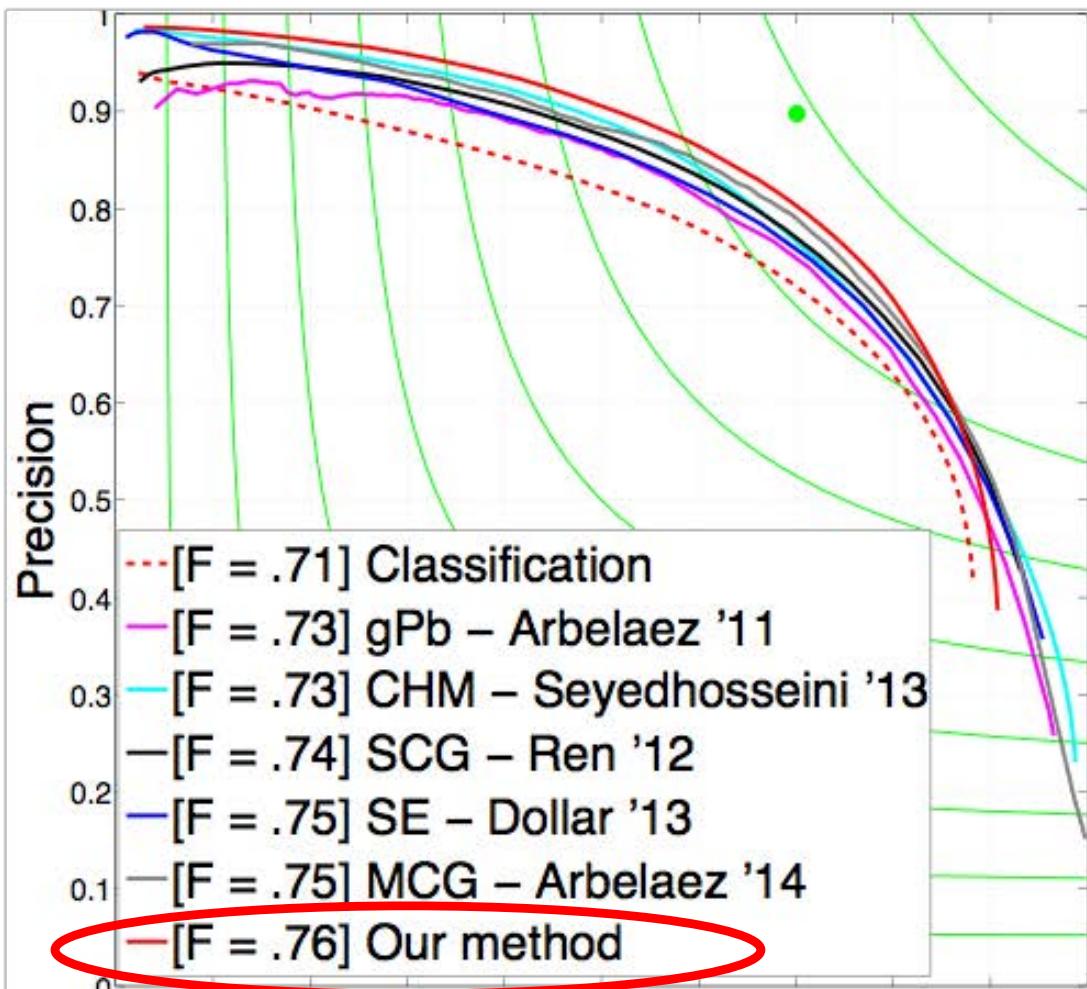


Learn the probability of being a boundary pixel on the basis of a set of features.

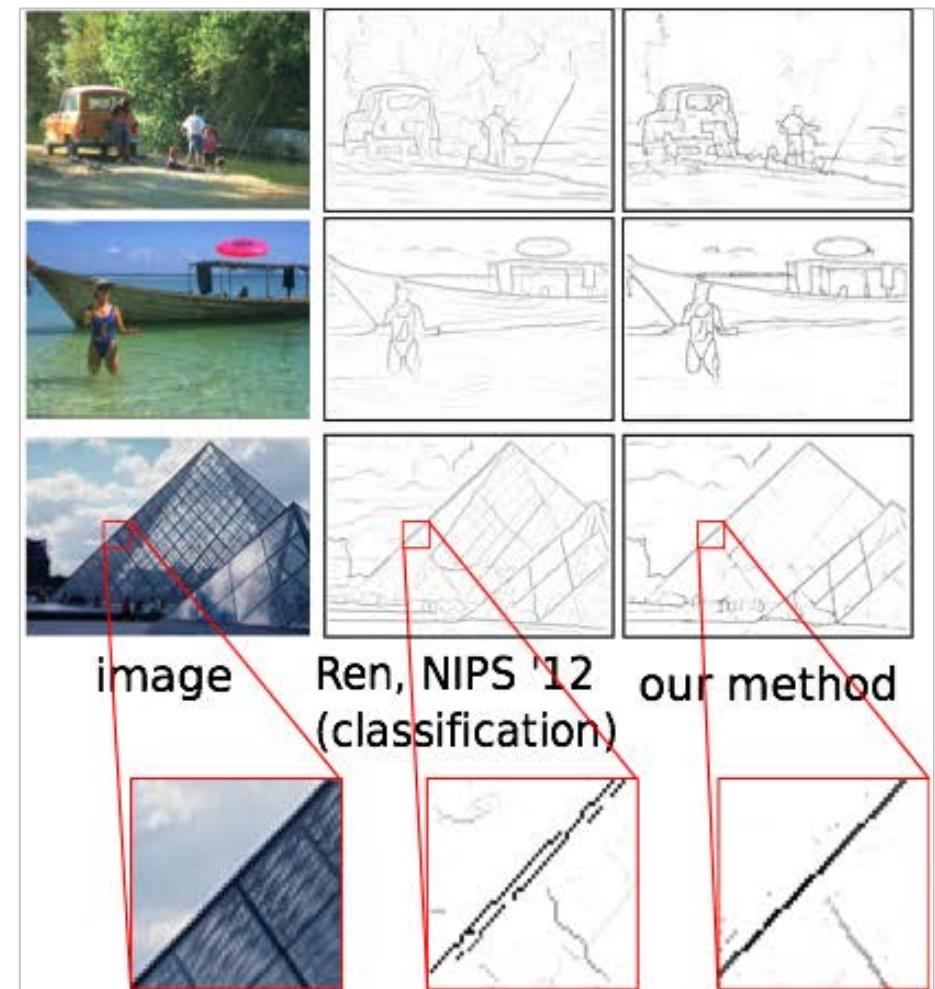
RESULTS



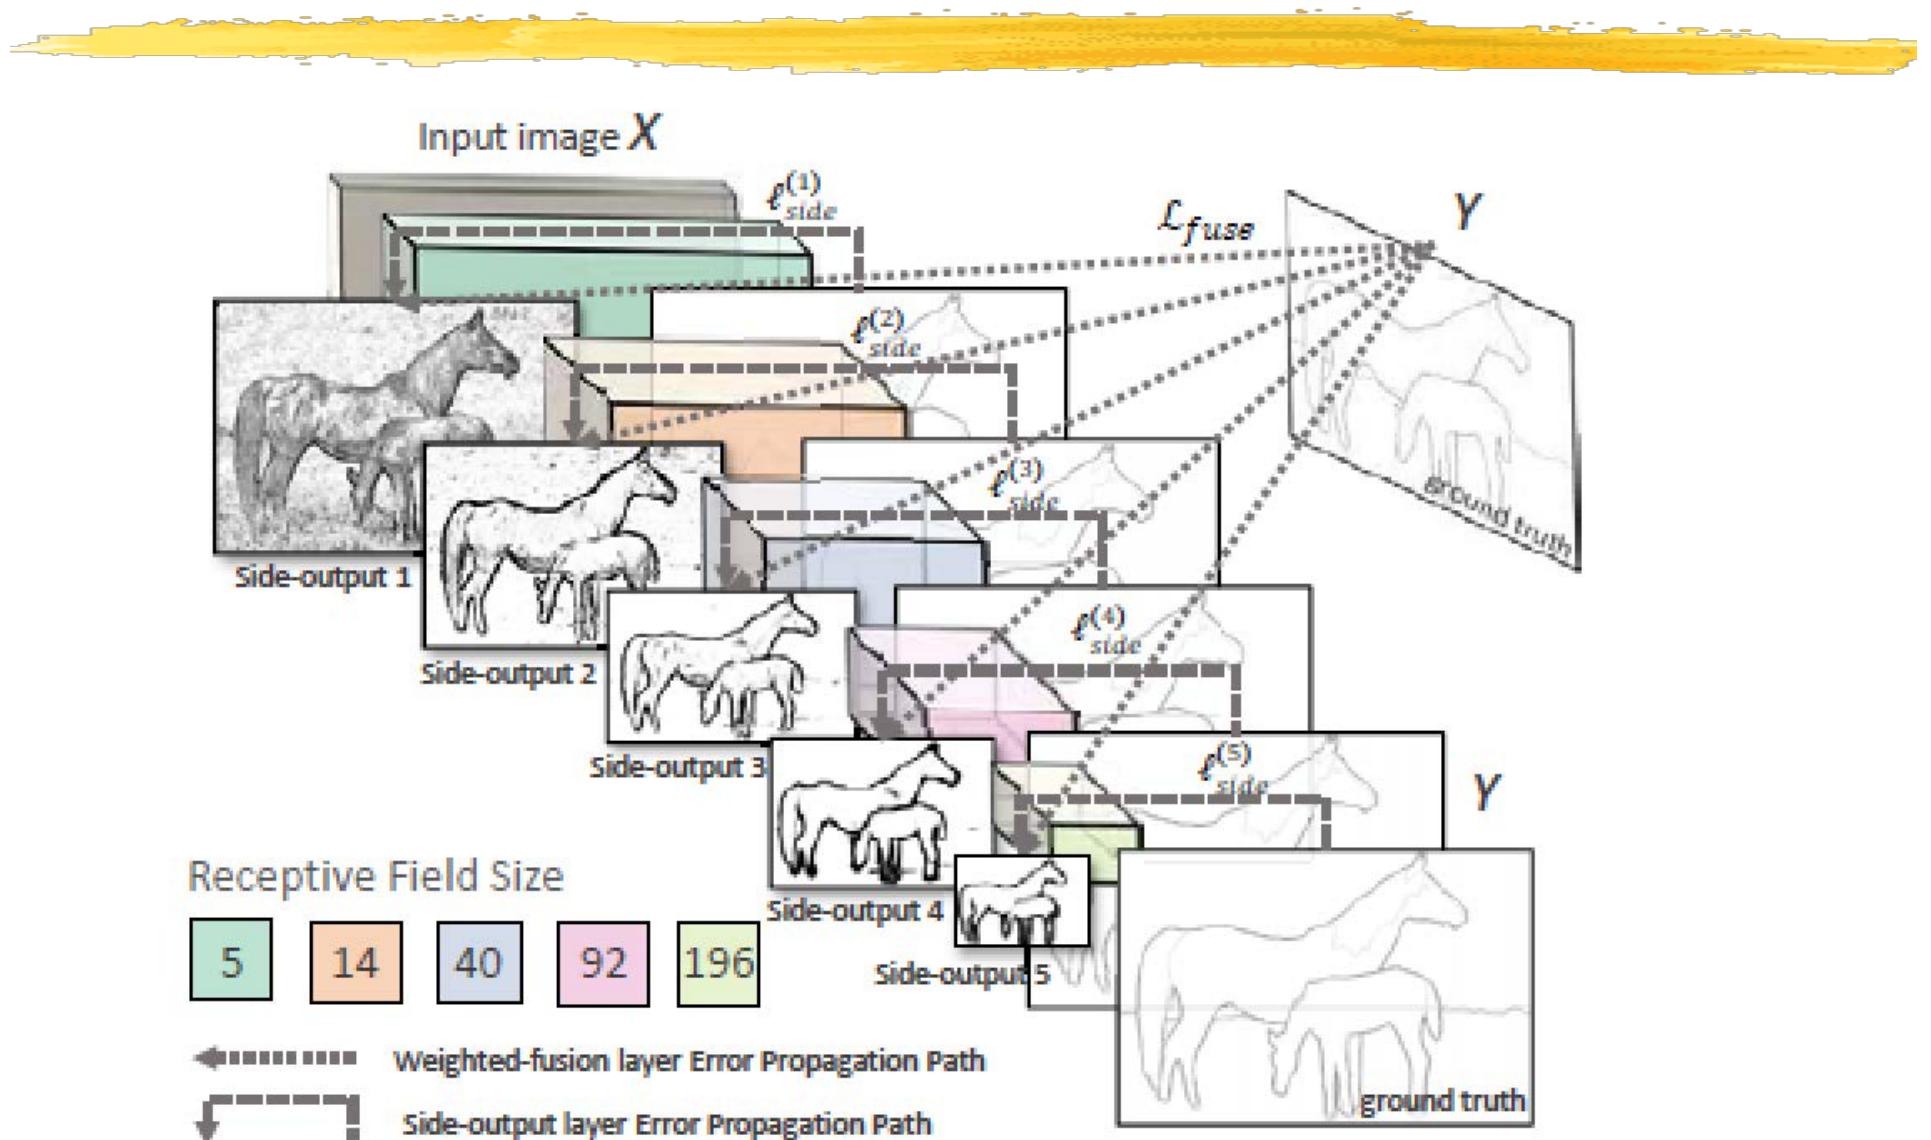
CLASSIFICATION vs REGRESSION



Yes!



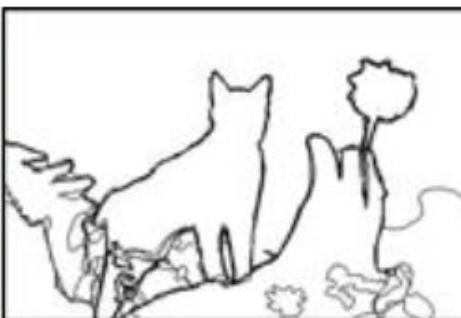
DEEP LEARNING



DEEP LEARNING VS CANNY



(a) original image



(b) ground truth



(c) HED: output



(d) HED: side output 2



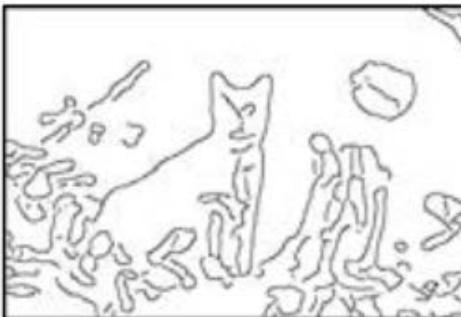
(e) HED: side output 3



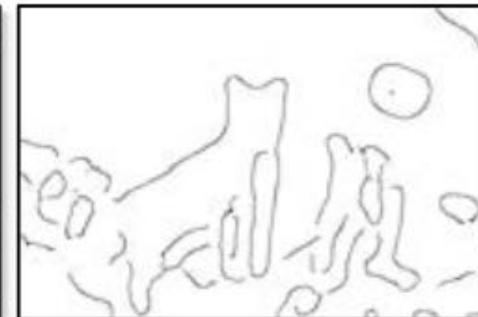
(f) HED: side output 4



(g) Canny: $\sigma = 2$



(h) Canny: $\sigma = 4$

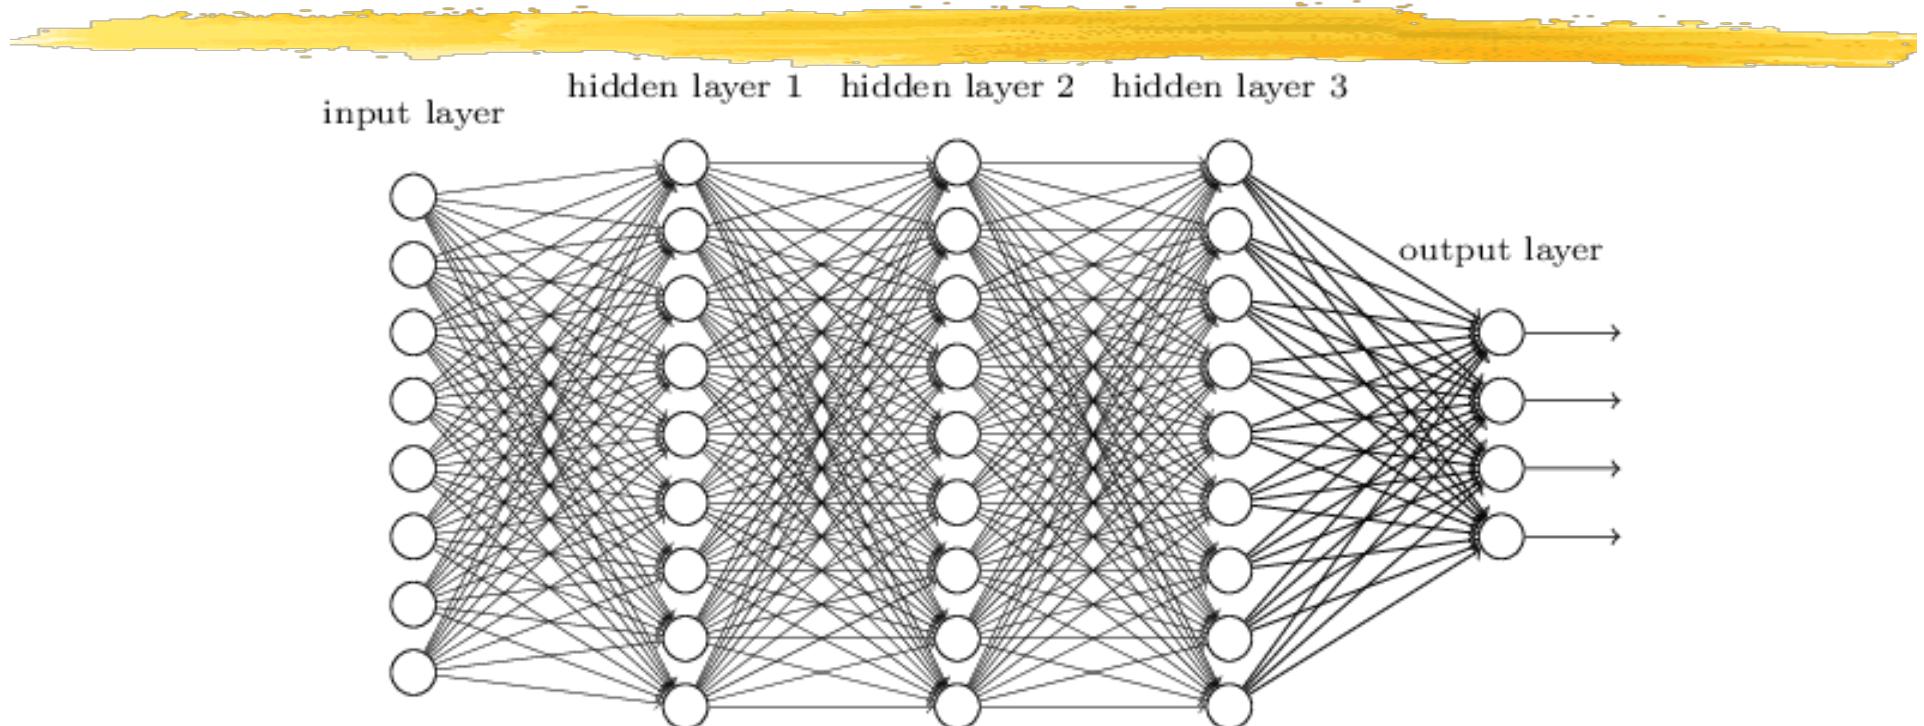


(i) Canny: $\sigma = 8$

DEEPER LEARNING

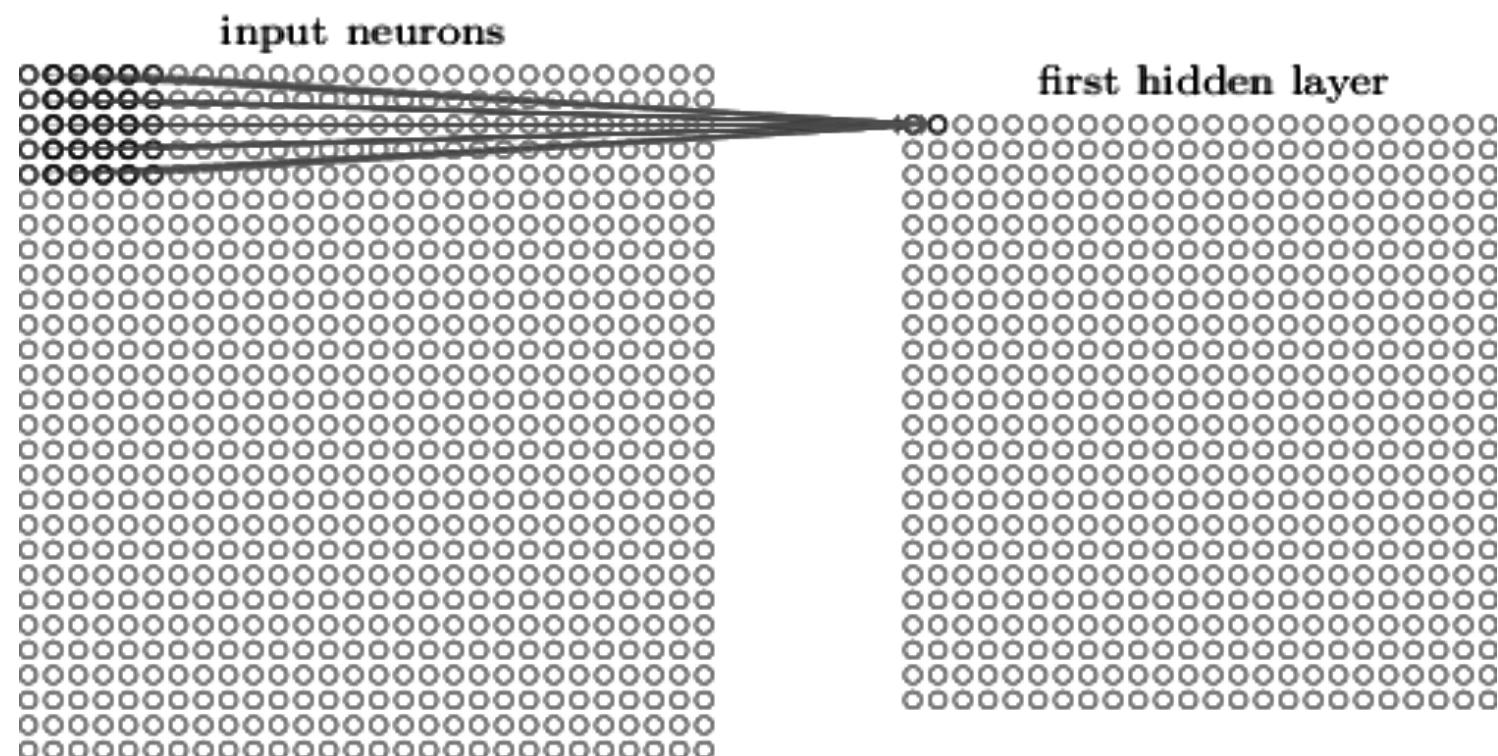


DEEP LEARNING



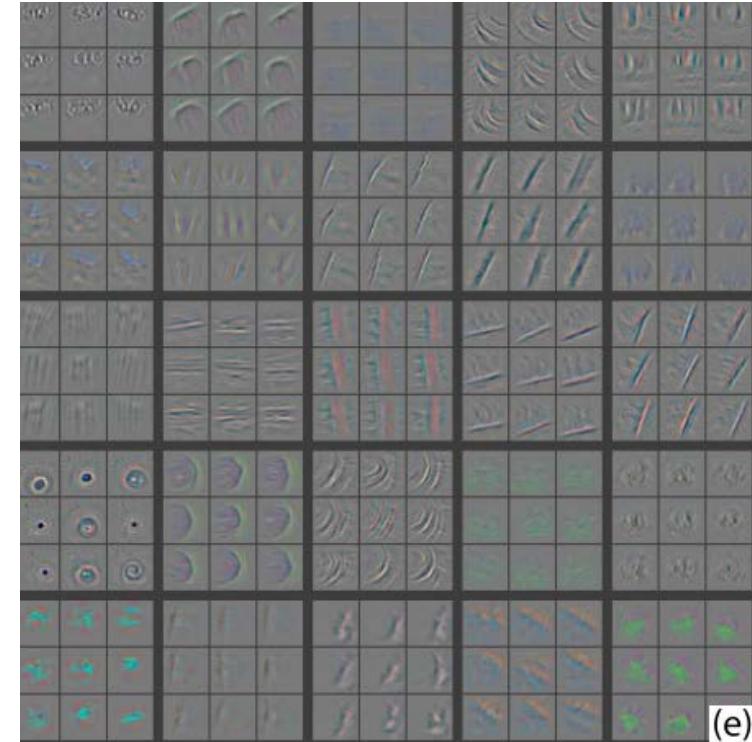
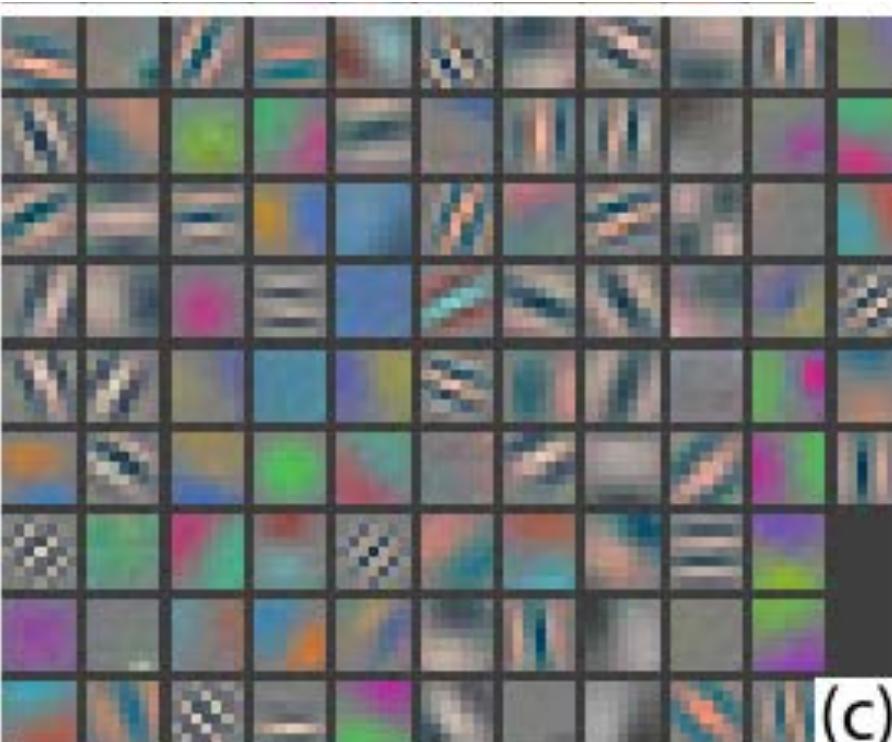
- Store in each node a function of the linear combination of the activations of all nodes in the previous layers.
- The network can be trained to produce a desired output given a specific input by learning the linear combination weights.

CONVOLUTIONAL LAYER



$$\sigma \left(b + \sum_{x=0}^{n_x} \sum_{y=0}^{n_y} w_{i,j} a_{i+x, j+y} \right)$$

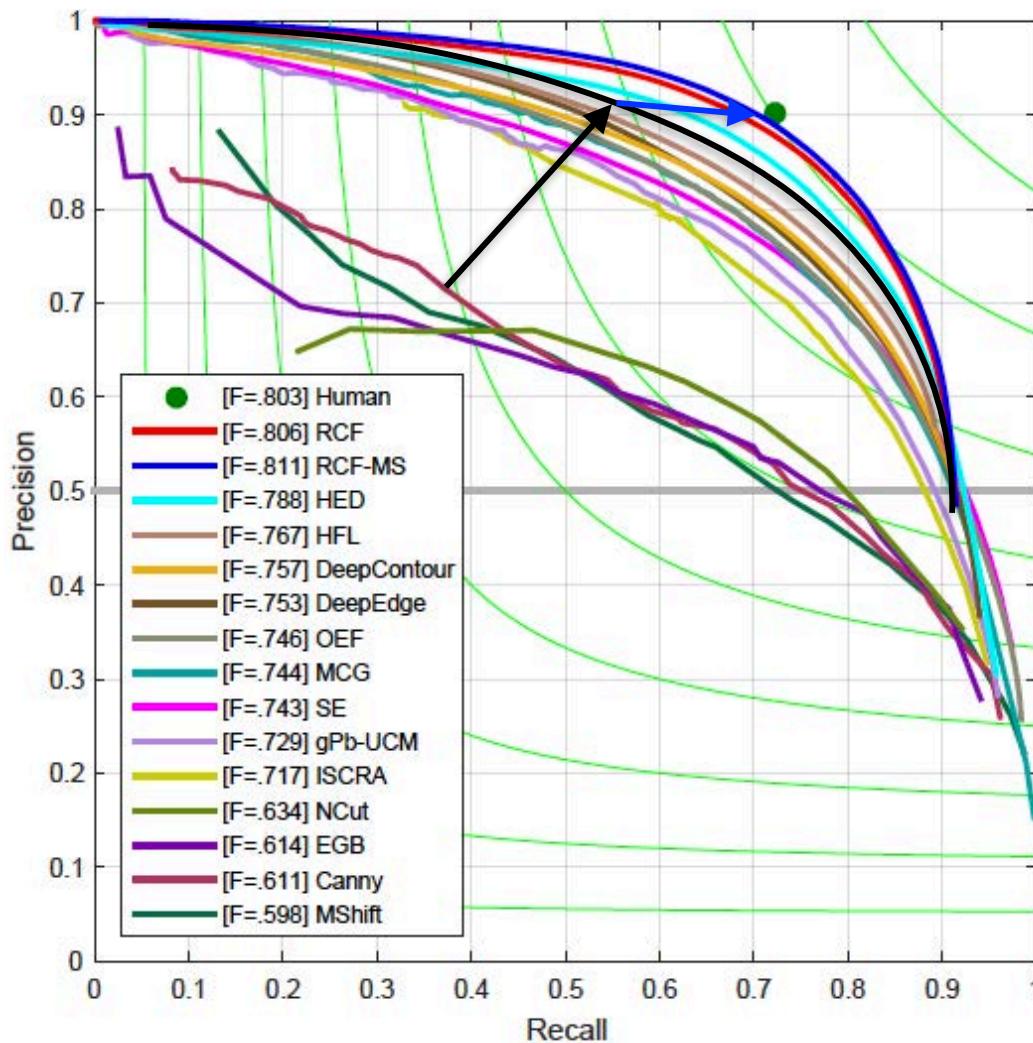
A PARTIAL EXPLANATION?



First and second layer features of a Convolutional Neural Net:

- They can be understood as performing multiscale filtering.
- The weights and thresholds are chosen by the optimization procedure.

50 YEARS OF EDGE DETECTION



- Convolution operators respond to steep smooth shading.
- Parametric matchers tend to reject non ideal edges.
- Arbitrary thresholds and scale sizes are required.
- Learning-based methods need exhaustive databases.
- There still is work to go from contours to objects.

Canny, PAMI'86 → Sironi et al. PAMI'15

Sironi et al. PAMI'15 → Liu et al., CVPR'17